



PERCONA

TECH DAYS

Elephant by the pool

PostgreSQL connection poolers overview

Sergey Kuzmichev & Jobin Augustine

Senor support Engineers - Percona

Percona Tech Days



Agenda

- Why do we need connection poolers at all
- How Connection Pooler address the problem
- External connection poolers, Features & Advantages
 - PgPool II
 - PgBouncer
 - Odyssey
- Limitations of External connection pooler one should be aware

Why we need new connection poolers

Cost of connection

- Every client connection need a dedicated new process
- Process forking
 - memory allocation, cancel key and other housekeeping
 - child process completes all authentication and handshakes
 - OS level overheads
- Overhead of connections are higher in process based databases.

Common DBA Observations

- New server processes are created when lot are “idle” in state.
- Too many PostgreSQL processes in the host OS
- Disconnection, deallocations / clean up

Cost of idling connections

- **Memory overhead**
- **Scheduling overhead**
- **Locking overhead**
 - Semaphores

SEMMNI	Maximum number of semaphore identifiers (i.e., sets)	at least $\text{ceil}((\text{max_connections} + \text{autovacuum_max_workers} + \text{max_wal_senders} + \text{max_worker_processes} + 5) / 16)$ plus room for other applications
SEMMNS	Maximum number of semaphores system-wide	$\text{ceil}((\text{max_connections} + \text{autovacuum_max_workers} + \text{max_wal_senders} + \text{max_worker_processes} + 5) / 16) * 17$ plus room for other applications
SEMMSL	Maximum number of semaphores per set	at least 17

- **Requirement of higher limits for connections.**
 - Server Abuse and Outages

Impact of authentication

Query over Fresh Connection

Trusted Connection

tps = **129.153727** (including connections establishing)

tps = 179.175602 (excluding connections establishing)

Password authentication

tps = **95.778541** (including connections establishing)

tps = 147.387896 (excluding connections establishing)

Password over TLS

tps = **72.597286** (including connections establishing)

tps = 117.009713 (excluding connections establishing)



Query over Persistent connection

Password over TLS

tps = **658.220014** (including connections establishing)

tps = 658.519444 (excluding connections establishing)

The solution : Connection poolers

- **Application side Connection pooler**
 - Support from ORMs
 - Right-sizing challenges
- **External connection pooler**
 - When there is no support by Application frameworks / language
 - Third party applications

What external poolers do

- Initialize a pool of connection at the startup or at the first connection.
- A user and database combination defines a pool.
- Connections are released back to pool. But not disconnected from the database side.
- Limited number of connections serving larger number of application connections.
- Connection **request queueing** at pool level.
- **Routing** of the connection

Popular External Poolers

pgpool-II

- **First version in 2003**
- **Connection Pooler**
- **High Availability Solution for PostgreSQL**
 - New quorum and consensus based backend failover
- **SPOF of pgpool can be avoided using multiple pgpool and Watchdog, Virtual IP failover**
- **Replication by managing multiple backend PostgreSQL**
- **Connection queuing.**
- **query cache**

pgpool-II

- **High complexity**
- **Overhead**

pgbouncer

- Lightweight and Performant
- SSL support
- Popular
- Monitoring
 - https://github.com/prometheus-community/pgbouncer_exporter
- Available in almost all repositories.

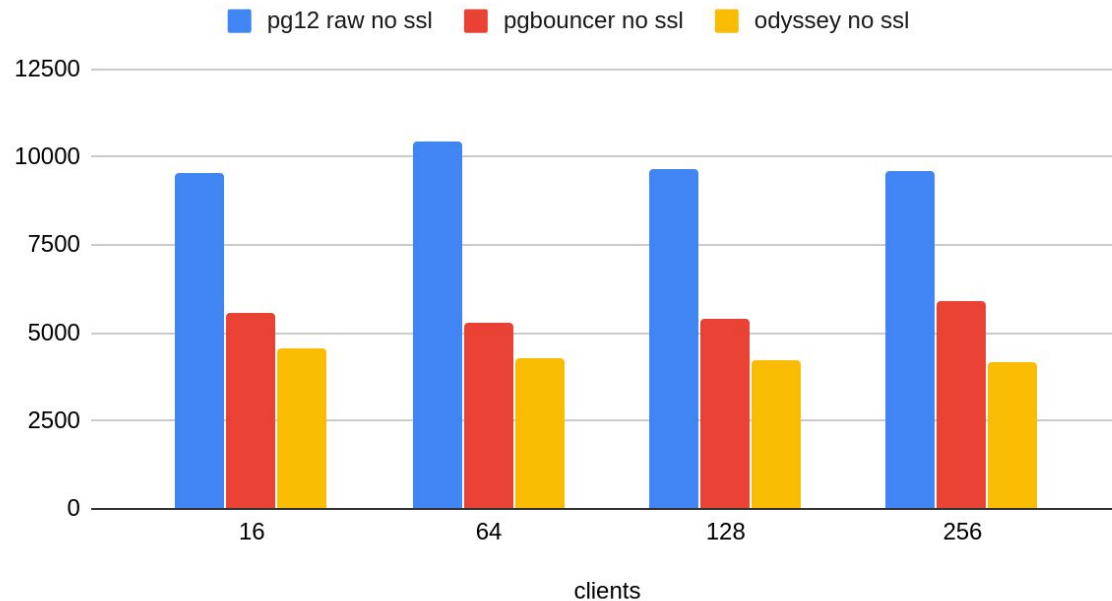
One Instance of pgbouncer uses only one CPU core.
Refer `so_reuseport` more details with latest kernel

Odyssey


- Multi-threaded using worker threads
- Feature parity with pgbouncer
- Very new
- Requires OpenSSL 1.1 (CentOS 8 or custom build)
- Monitoring (same as pgbouncer)
 - https://github.com/prometheus-community/pgbouncer_exporter
- Improves on pgbouncer pitfalls, like performance with SSL

Persistence at application

pgbench read-only no ssl

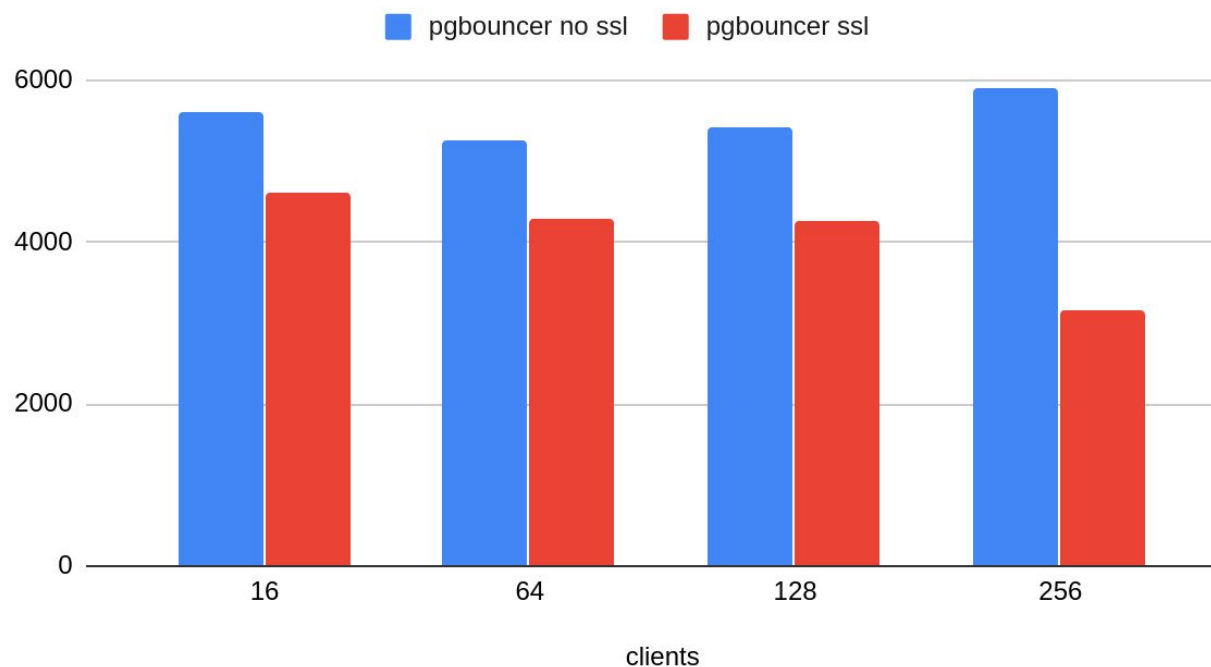


pgbench read-only SSL

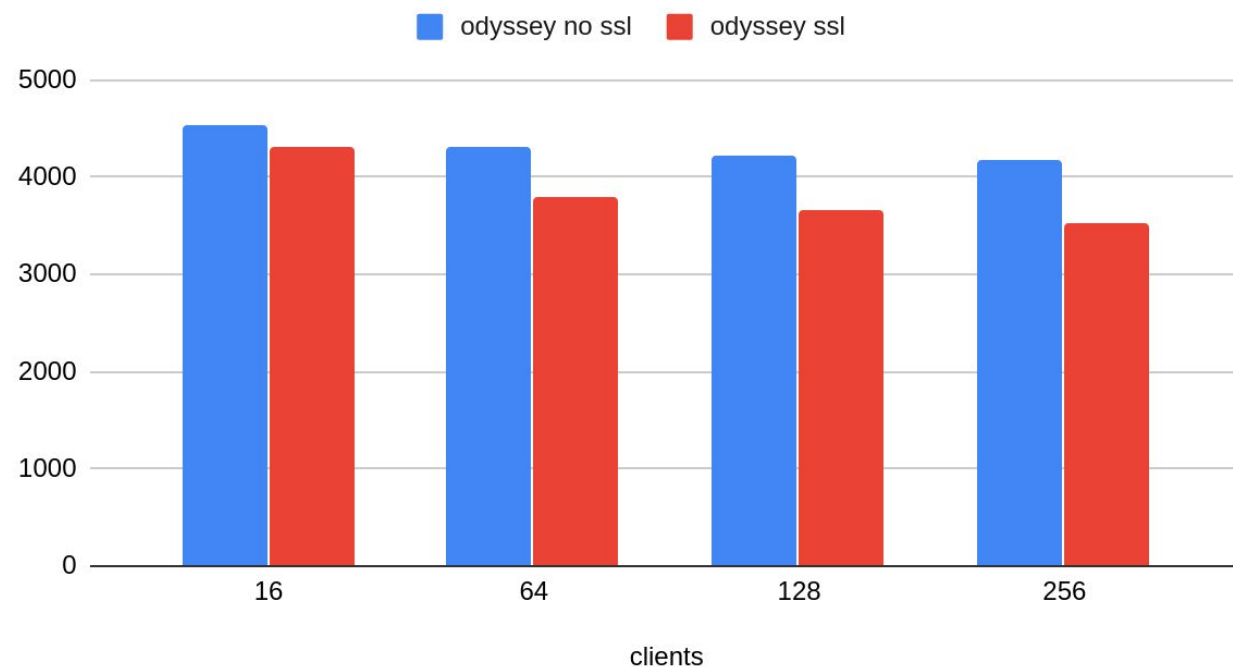


pgbouncer and odyssey

pgbouncer no ssl and pgbouncer ssl

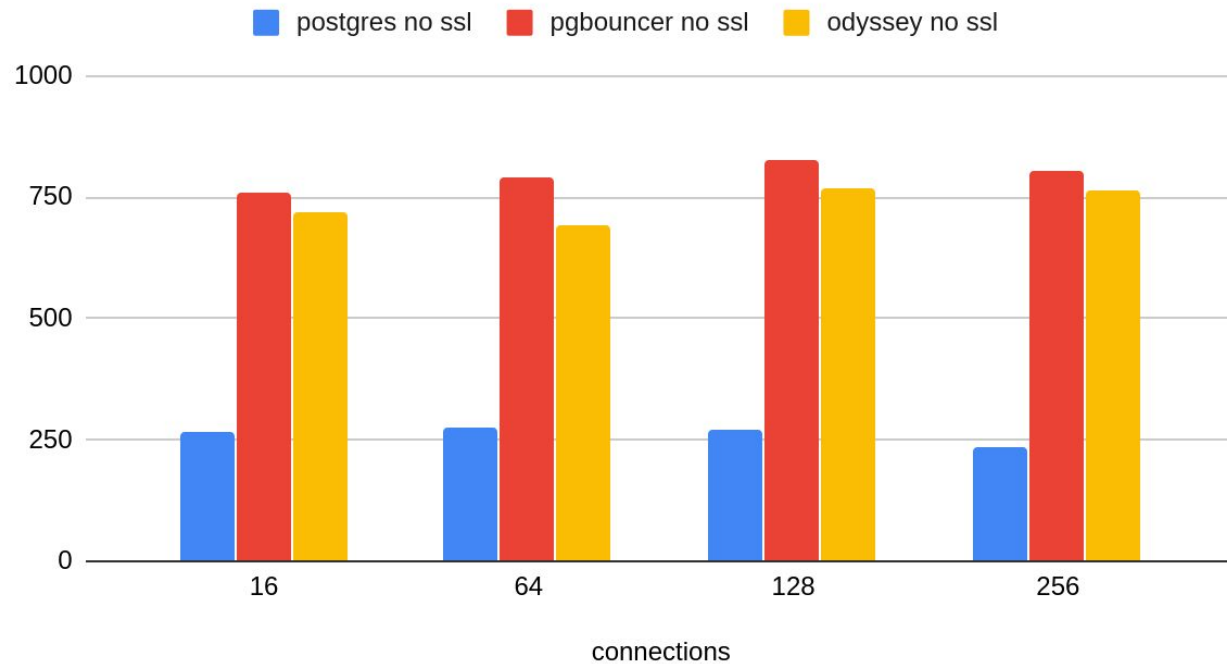


odyssey no ssl and odyssey ssl

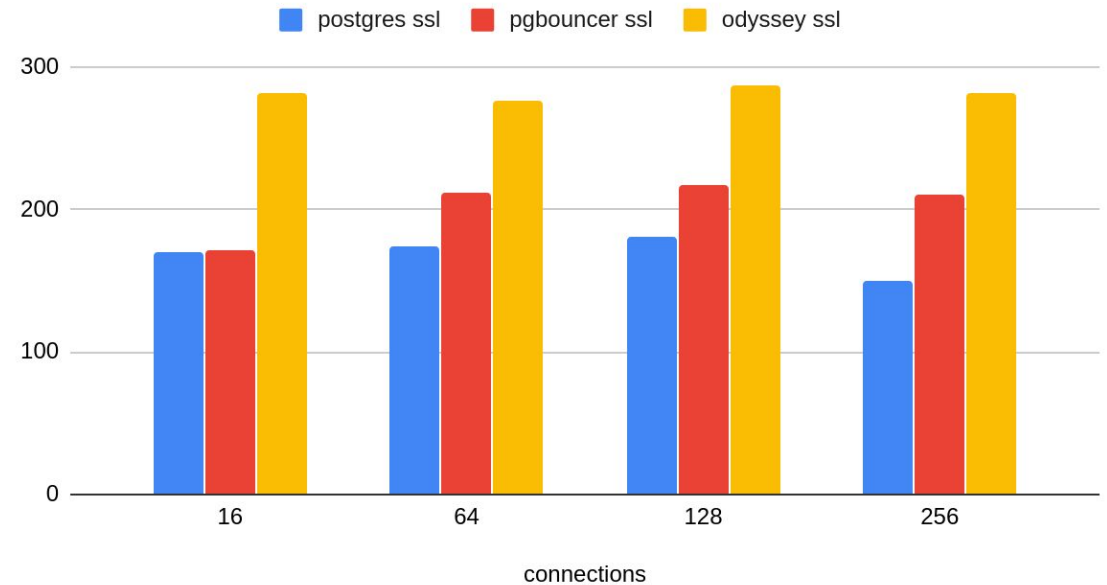


New Connections


No SSL, non-persistent client connections




SSL, non-persistent client connections



Overhead of SSL - No Pooler

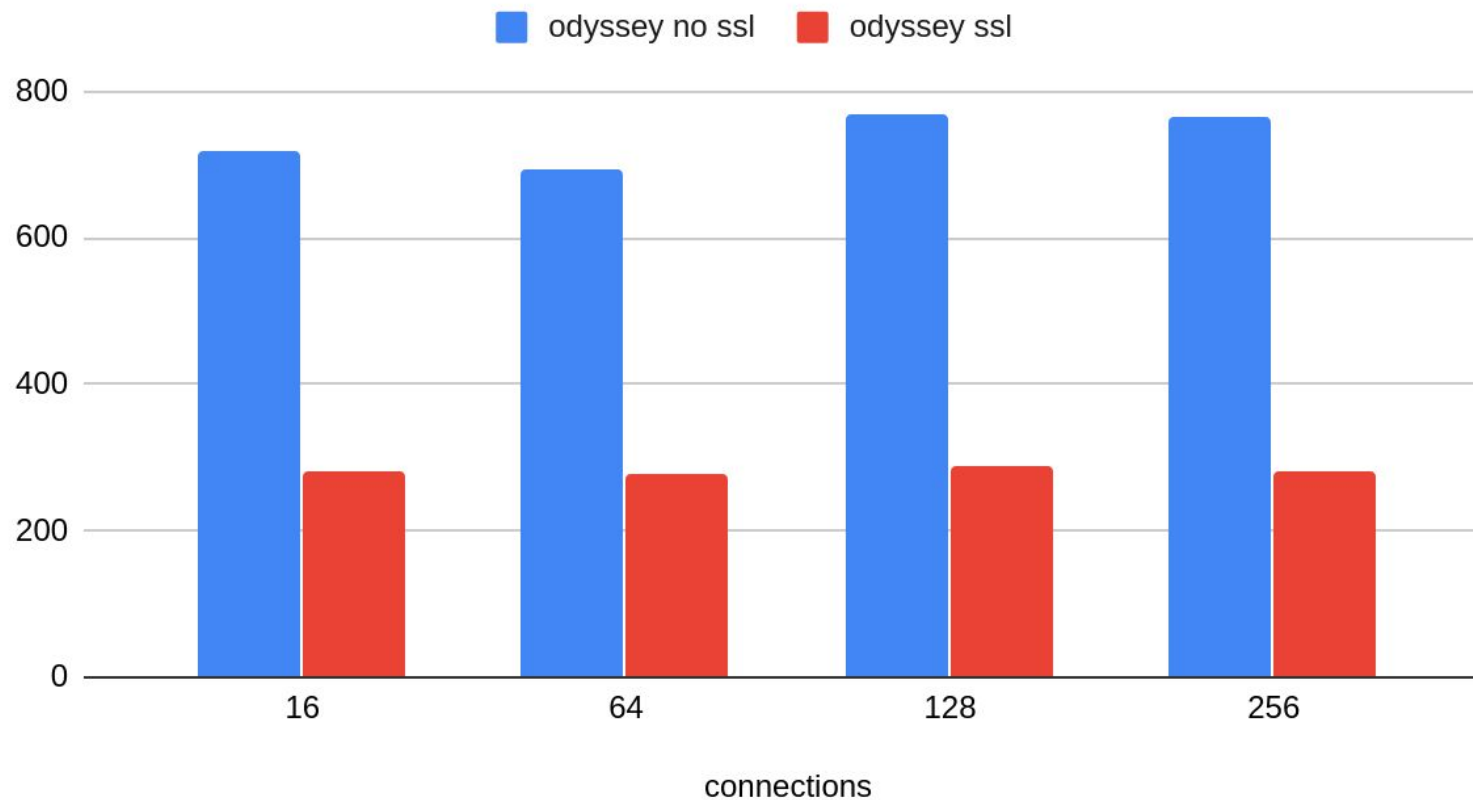


Overhead of SSL - PgBouncer



Overhead of SSL - odyssey

odyssey no ssl and odyssey ssl



Limitations of Connection Poolers

Limitations of External connection Poolers

Security, Auditing, Monitoring and troubleshooting Challenges

- Client address in the server logs points to pooler
- Server side audits won't have end client information
- tcpdump of the connection stream won't be possible
 - client may use different connection each time
- Sessions are shared

Its a man in middle

Limitations of External connection Pools

Limitations On functionality

- prepared statements
- temporary tables

Limitations of External connection Pools

Performance and Stability

- Applications which needs a persistent connection won't see any performance benefit.
 - Releasing to pool and getting it back
 - Extra hop
- I could be performance overhead
 - Eg : `serve_reset_query = DISCARD ALL`; or at least to `DEALLOCATE ALL`
- They won't remove SPOF (they become SPOF!)

References

prepared statements in Odyssey

<https://github.com/yandex/odyssey/issues/16>

<https://blog.bullgare.com/2019/06/pgbouncer-and-prepared-statements/>

https://pgconf.ru/media/2017/04/03/20170316H1_V.Borodin.pdf

potential prepared statements workaround

<https://github.com/dimitri/preprepare>

pg_hba in odyssey

<https://github.com/yandex/odyssey/issues/75>

cursor :

<https://github.com/yandex/odyssey/issues/45>

leader failover

<https://github.com/yandex/odyssey/issues/44>

prepared statements in pgbouncer

<https://github.com/pgbouncer/pgbouncer/issues/499>

application name reset in pgbouncer

<https://github.com/pgbouncer/pgbouncer/issues/457>

<https://www.pgpool.net/docs/latest/en/html/runtime-config-connection-pooling.html>

pgpool only has session-level pooling

<https://github.com/yandex/odyssey/issues/3>

<https://github.com/kwent/pgbouncerhero>

Pecrona Blog

<https://www.percona.com/blog/2018/06/27/scaling-postgresql-with-pgbouncer-you-may-need-a-connection-pooler-sooner-than-you-expect/>




PERCONA
LIVEONLINE
20-21 OCTOBER
2020

Two-Column Slide

This symmetrical layout is optimized for approximately square charts or graphics.


- Use only level-one bullets when possible
 - *Use level two bullets and below judiciously to maximize legibility*



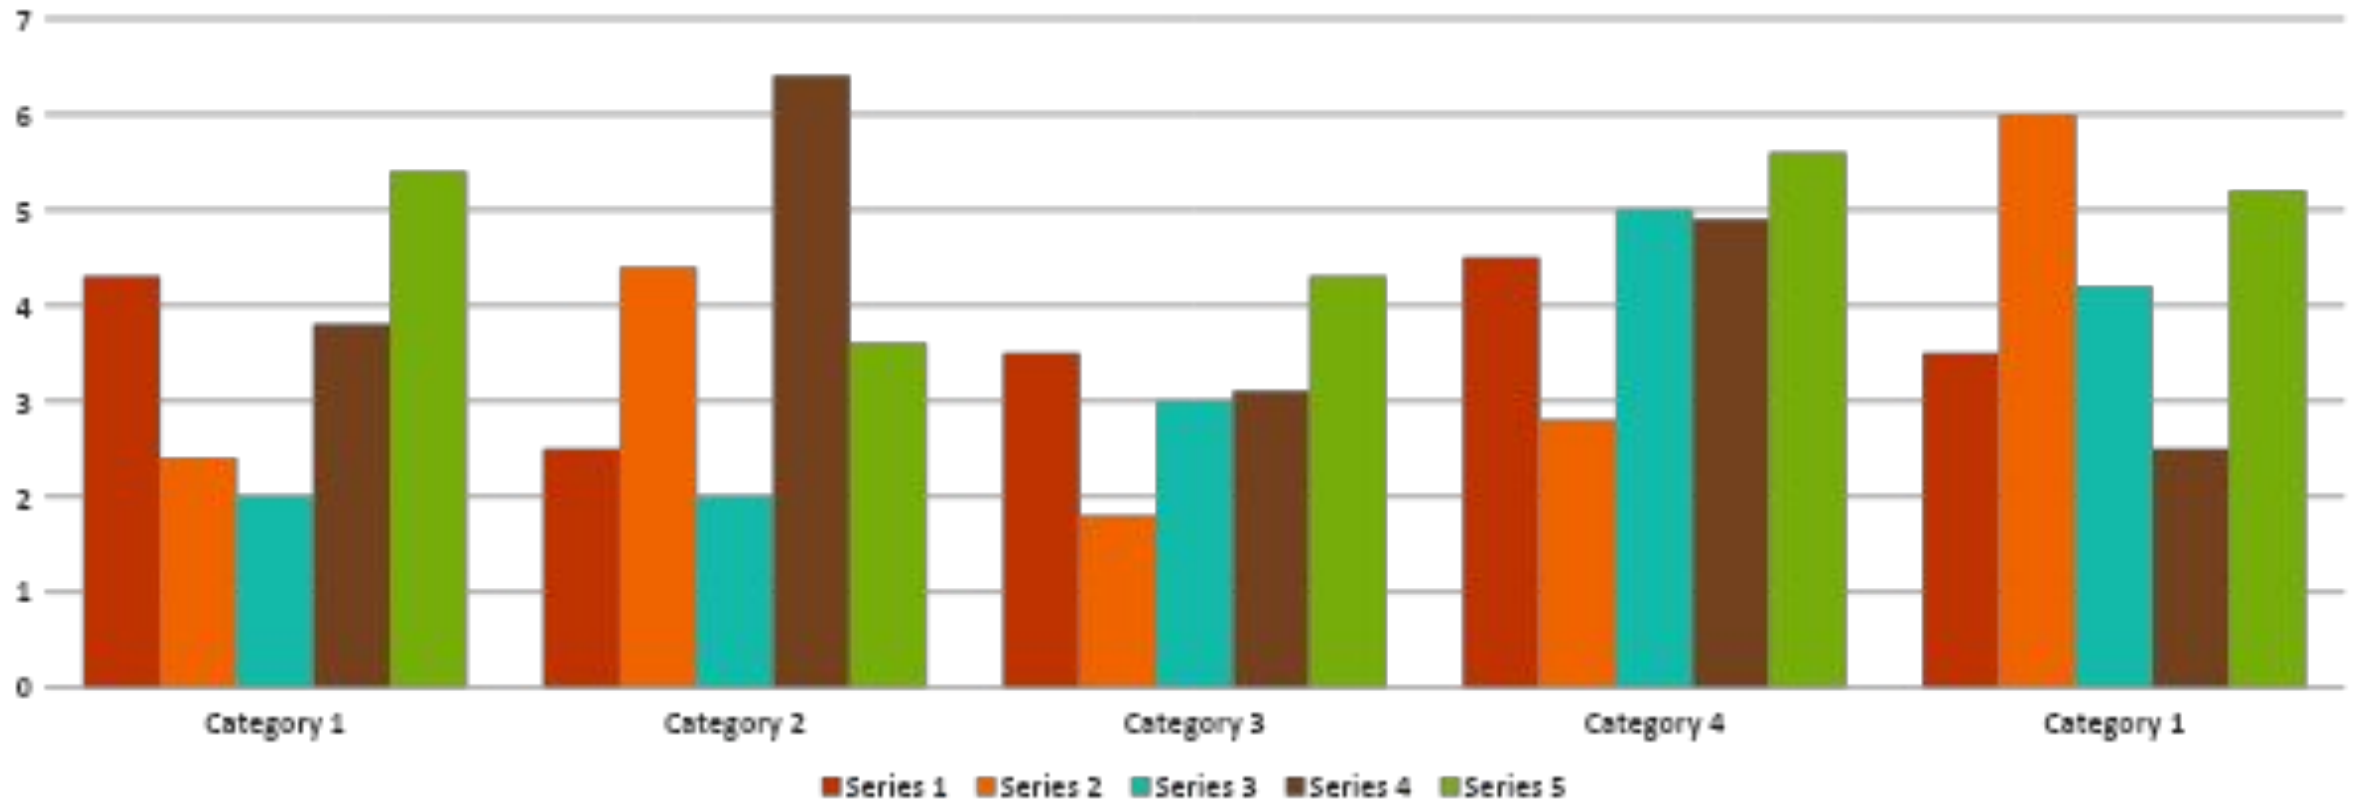
Two-Column Offset

**Narrowed text box
accommodates wider
charts and graphics**

- Bullets should be short and concise
- Use of level one bullets increases legibility and maintains a clean look



Single-Column Slide for Horizontal Content



About Percona

Solutions for your success with MySQL and MongoDB

Support, Managed Services, Software

Our Software is 100% Open Source

Support Broad Ecosystem – MySQL, MariaDB, Amazon RDS

In Business for 10 years

More than 3000 customers, including top Internet companies and enterprises

About This Presentation

Overview of
Existing Solutions
and History

Discuss what we
use at Percona

Show what
specific things to
look at

Smart Object Chevron List





**Champions of Unbiased
Open Source Database Solutions**