


PostgreSQL connections at scale

About me



<https://bhuvane.sh/>

Bhuvanesh (aka The DataGuy)

Data Architect + DBA



<https://linkedin.com/in/rbhuvanesh/>



BhuviTheDataGuy



BhuviTheDataGuy



BhuviTheDataGuy (Tech Blog)



<https://thedataguy.in> (Tech Blog)

Disclaimer



This presentation is based on my personal experience & research about PostgreSQL and pgbouncer.


No copyright !! Nothing reserved !!!

Icons used - <https://undraw.co/>

GIFs used - <https://giphy.com/>

Agenda

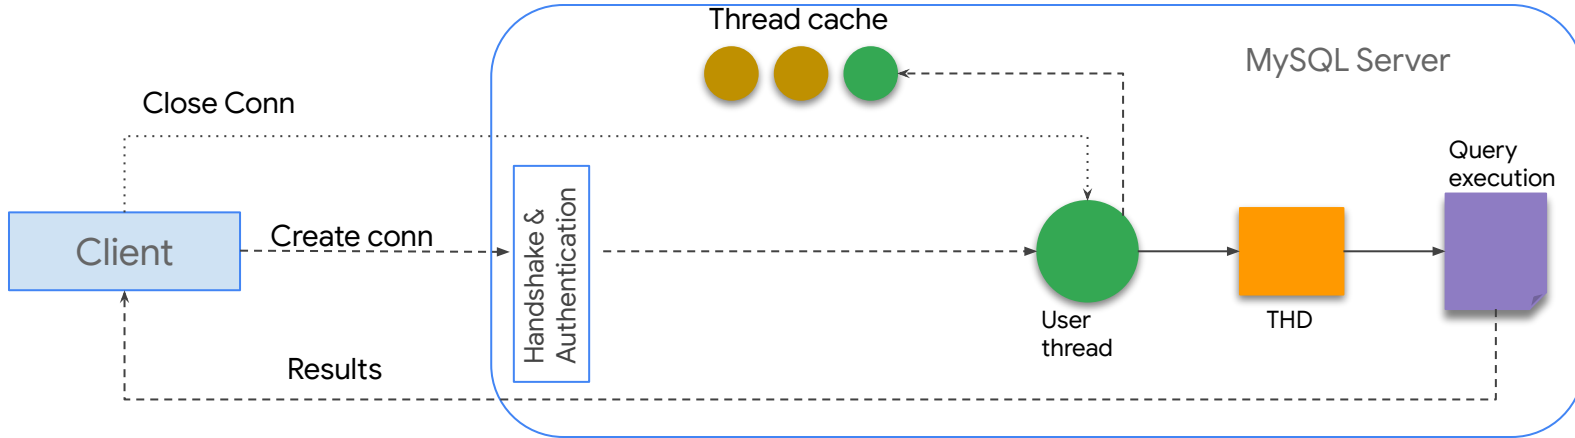
1. A database connection and its **Cost**
2. Need for a connection pooler
3. Pgouncer - Introduction
4. How pgouncer works and its features
5. Monitoring the pgouncer
6. Pgouncer deployment best practices





**Why PostgreSQL doesn't
scale to a large number of
connections?**

Life cycle of a connection in MySQL



Reference: <https://mysqlserverteam.com/mysql-connection-handling-and-scaling/>

Cost of a single PostgreSQL connection

- Each connection is a process fork with roughly consume **10 MB** memory (*a research from heroku*)
- Each connection will simultaneously open up to **1000 files** (*default configuration*)

Example:

Let's assume your DB is consuming 400 connections, then

- $10\text{MB} * 400 \text{ connections} = 4\text{GB Memory}$
- $1000 \text{ Files} * 400 \text{ connections} = 4,00,000 \text{ Files}$

What about Idle connections?

States of PostgreSQL connection:



Active	currently running
Idle	Its not doing anything
Idle in transaction	currently not doing anything and could be waiting for an input
Idle in transaction(aborted)	the connections that were idle in the transaction that have since been aborted

“Idle connections are not just idle, they eat the resources in all the ways.”

Cost of the idle connections

A benchmark from AWS:

- Open 100 connections.
- Leave the connections idle for 10 minutes.
- Close the connections.



Reference: <https://aws.amazon.com/blogs/database/resources-consumed-by-idle-postgresql-connections/>

More incoming connections?

Then your PostgreSQL server is in Danger!




We need a Superhero



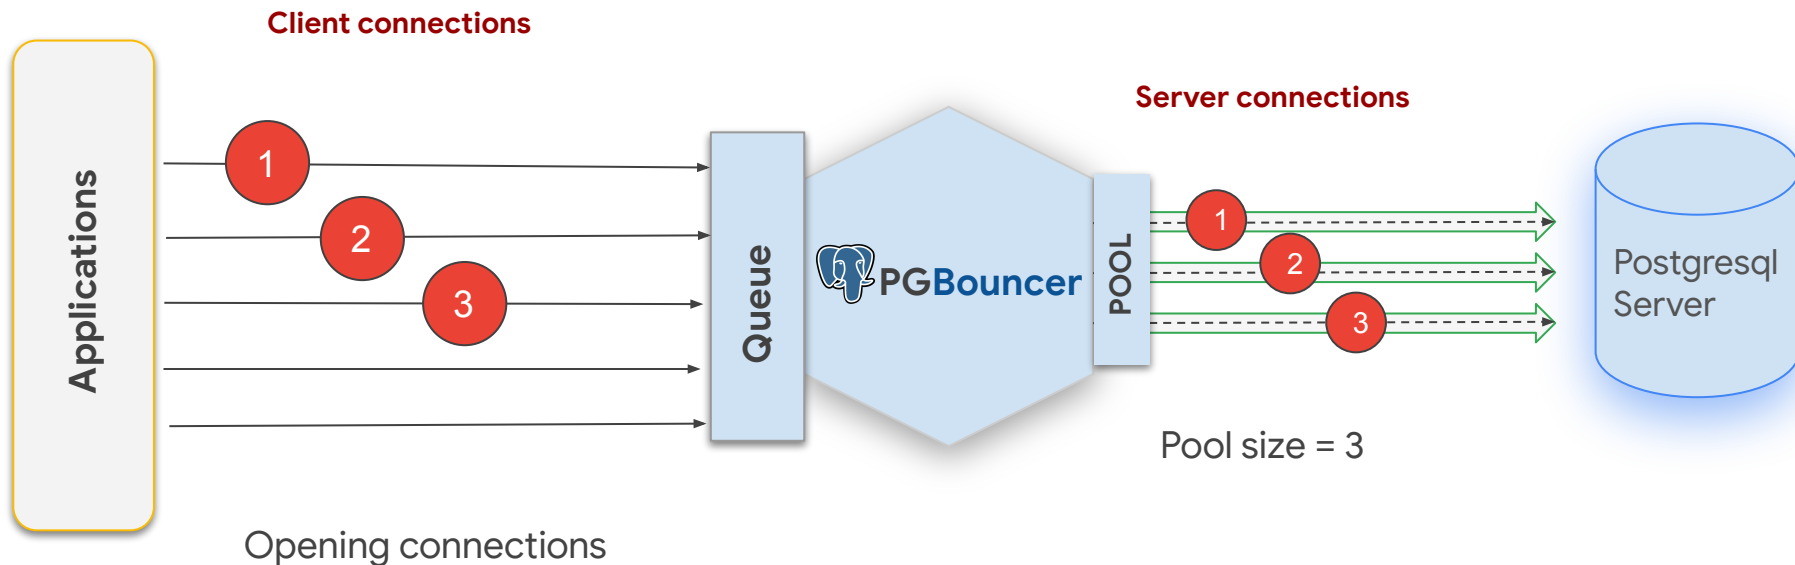
Not the Black Panther,
but
pgbouncer

pgbouncer


- Lightweight connection pooler
- Reduce the number of backend connections
- Connections economy (creating a connection is a fork and acquire a ProcArrayLock)
- No need for special authentication



What does a connection pooler do?




Lets deep dive into pgbouncer



Pgbouncer pooling mode

Session

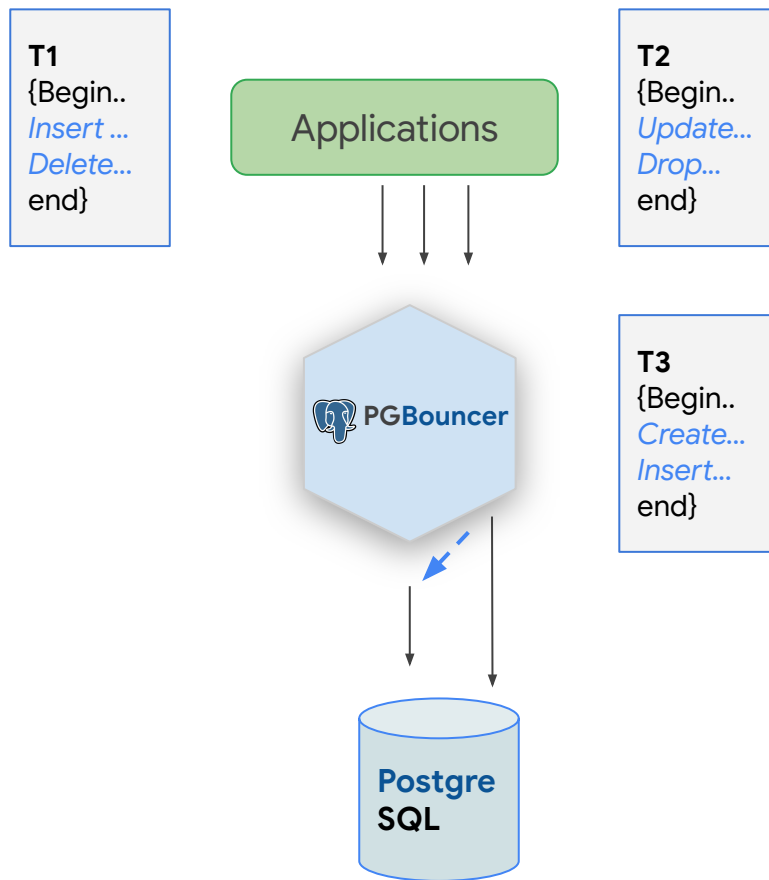
- Each client(app) connection will open backend connections and the connection will remain open until the client closes.
- Client and Server connections are mapped
- It is almost similar to using the database directly without any connection pooler.
- The connection is transparent.



Pgbouncer pooling mode

Transaction

- Client connect to the server connection only during the transaction, after the transaction the same server connection can be used to run another session's transaction.
- All the queries inside the {begin... end} will be executed in one server connection.
- Session variables and prepared statements will not work here.
- The connection is not transparent.




Pgbouncer pooling mode

Statement

- Very aggressive mode.
- The server connection back to the pool when the statement completes.
- Multi statement is not possible here.
- Auto commit is always on in this mode.

T1

*Insert ...
Update...
delete*



Adding the PostgreSQL Server

Like a connection string

:: Common patterns

```
db1 = host=primary dbname=db1
db2 = host=standby dbname=db2
* = host=standby
```

:: Custom pool size

```
db2 = host=standby dbname=db2 pool_size=50
reserve_pool=10
```



Pool size value can be override based on users

- Pool_size =10
- No of databases = 10
- No of users = 2
- Pool size can be override to **(no of db * no of users) = 20**

List of parameters:

- dbname
- host
- port
- user
- password
- auth_user
- client_encoding
- datestyle
- timezone
- pool_size
- reserve_pool
- max_db_connections
- pool_mode
- connect_query
- application_name

Authentication

Auth File with auth type

- A **txt file that contains username and the password**
- Password can be plain text of MD5 hash(recommended)
- Supported auth types:
 - any
 - trust
 - plain
 - md5
 - cert
 - hba
 - pam

pg_hba.conf

- Its very similar to postgresql's hba method.
- Same postgresql's syntax will work here.
- But LDAP, pam and a few other methods will not work.

Auth_user with query

- Automatically loads the username and password from the target database.
- Just give an user and password(mention it in the **auth_file**) for the authentication(like a dedicated user), then it'll fetch the user, password from the database.
- **example:**

```
SELECT username, passwd FROM  
pg_shadow WHERE username=$1
```

Authentication cont...

:: auth file

```
auth_type = trust  
auth_file = /etc/pgbouncer/userlist.txt
```

:: HBA-style

```
;auth_hba_file =
```

:: Auth user with query

```
;auth_query = SELECT username, passwd FROM pg_shadow WHERE username=$1
```

Connections and Pool

Parameter	Description
<code>max_client_conn</code>	Maximum number of connection allowed in pgbouncer
<code>default_pool_size</code>	How many server connections to allow per user/database pair.
<code>min_pool_size</code>	Minimum number of server connections
<code>reserve_pool_size</code> <code>reserve_pool_timeout</code>	If the pool is full, and a connection is waiting more than the <code>reserve_pool_timeout</code> then it'll use extra connection from this reserve pool. It should be less in size .

Example 1: How the connection pool makes the connection

Connection:

```
db1 = host=localhost dbname=db1
```

Pool:

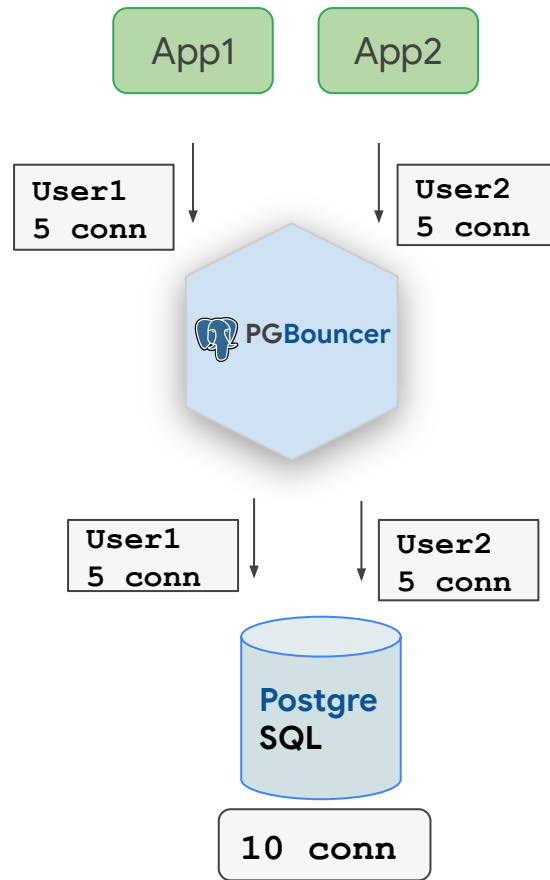
```
default_pool_size = 20  
pool_mode = session
```

Auth file:

```
user1:xxxxxx  
user2:yyyyy
```

Connection scenario:

- App1 is connecting to the DB via user1 with **5 connections**
- App2 is connecting to the DB via user2 with **5 connections**



Example 2: where the pool values override

Connection:

```
db1 = host=localhost dbname=db1
```

Pool:

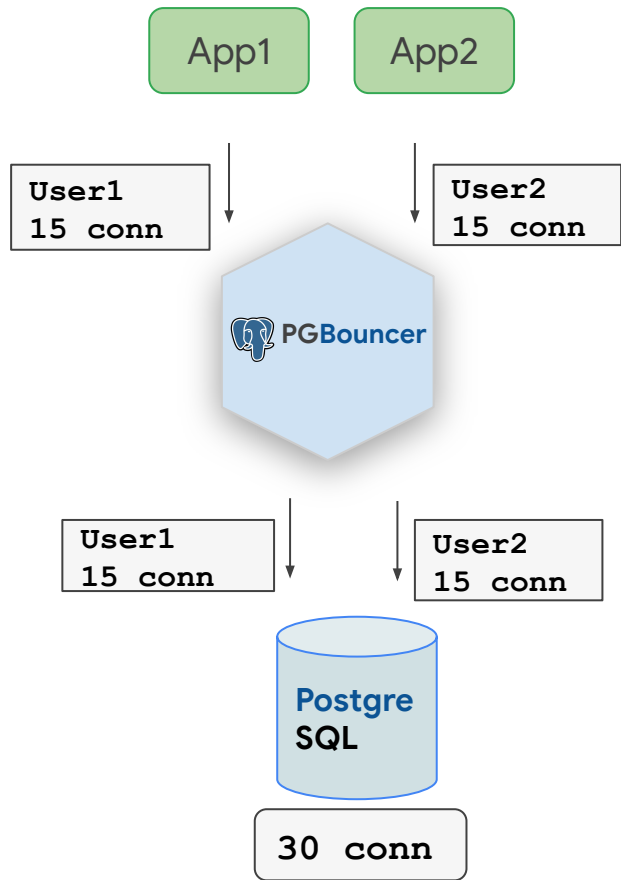
```
default_pool_size = 20  
pool_mode = session
```

Auth file:

```
user1:xxxxxx  
user2:yyyyy
```

Connection scenario:

- App1 is connecting to the DB via user1 with **15 connections**
- App2 is connecting to the DB via user2 with **15 connections**



Set the Hard limit for the pool size

Option 1: Set the limit based on DB or user level

Parameter	Description
<code>max_db_connections</code>	Pgbouncer will not allow more than this value for a database
<code>max_user_connections</code>	Maximum pool for a user


Option 2: Define separate connections for each DB and set the pool size via connection

```
db1 = host=34.72.164.39 dbname=db1
pool_size=10
db2 = host=34.72.164.39 dbname=db2
pool_size=10
```

[OR]

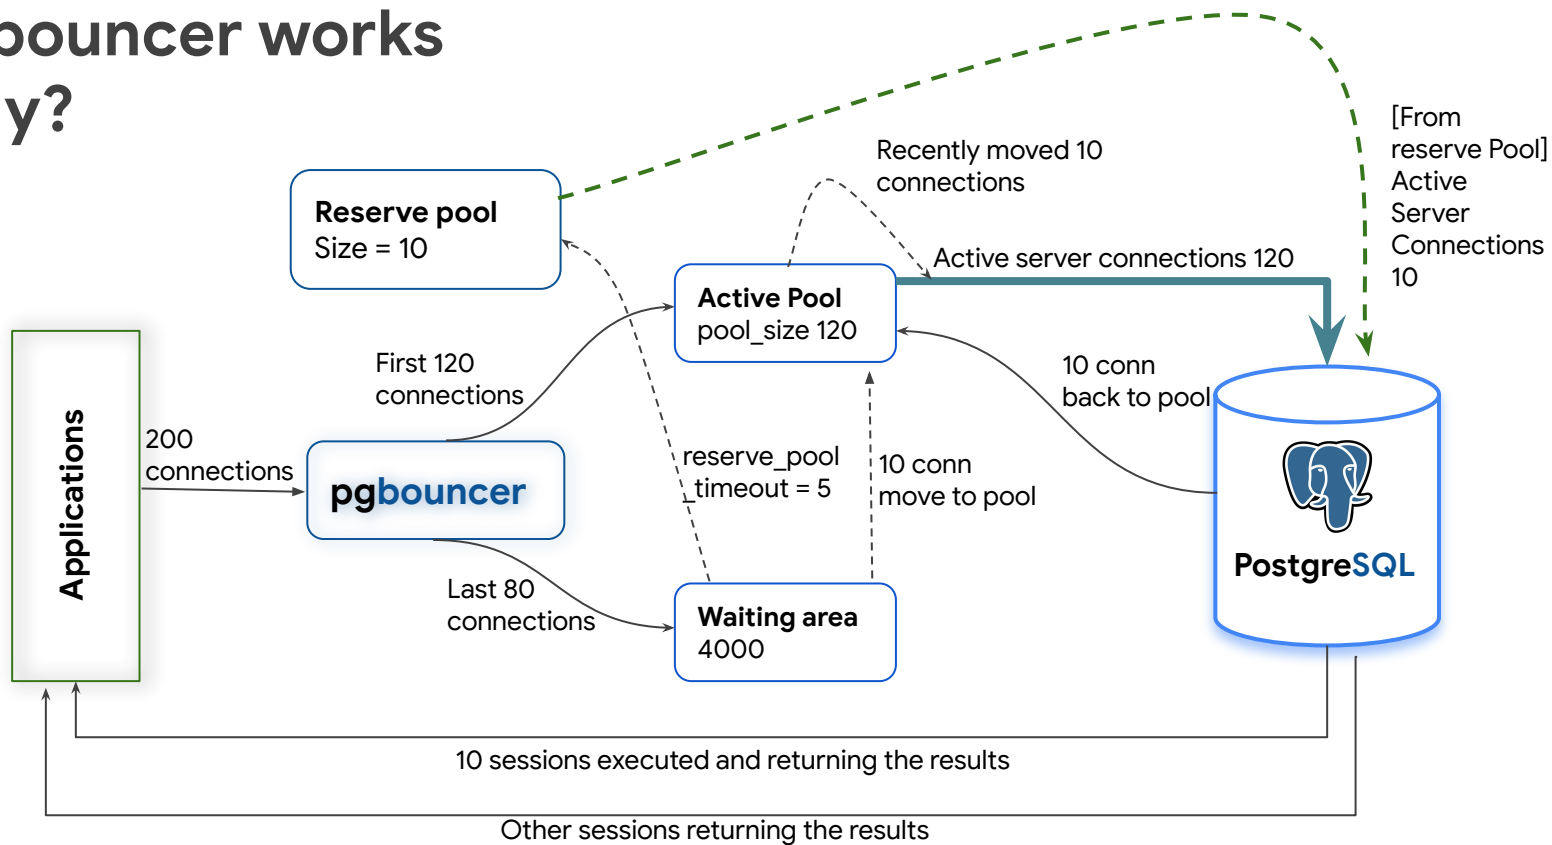
```
db1 = host=34.72.164.39 dbname=db1
pool_size=10 max_db_connections=10
db2 = host=34.72.164.39 dbname=db2
pool_size=10 max_user_connections=10
```

Lets see the real world example




- PostgreSQL max connections = 300
- No of databases = 1
- No of users = 1
- Pool size = 120
- Pool mode = session
- Max client connections = 4k
- Reserve pool size = 10
- Reserve pool timeout = 10s

How pgbouncer works internally?



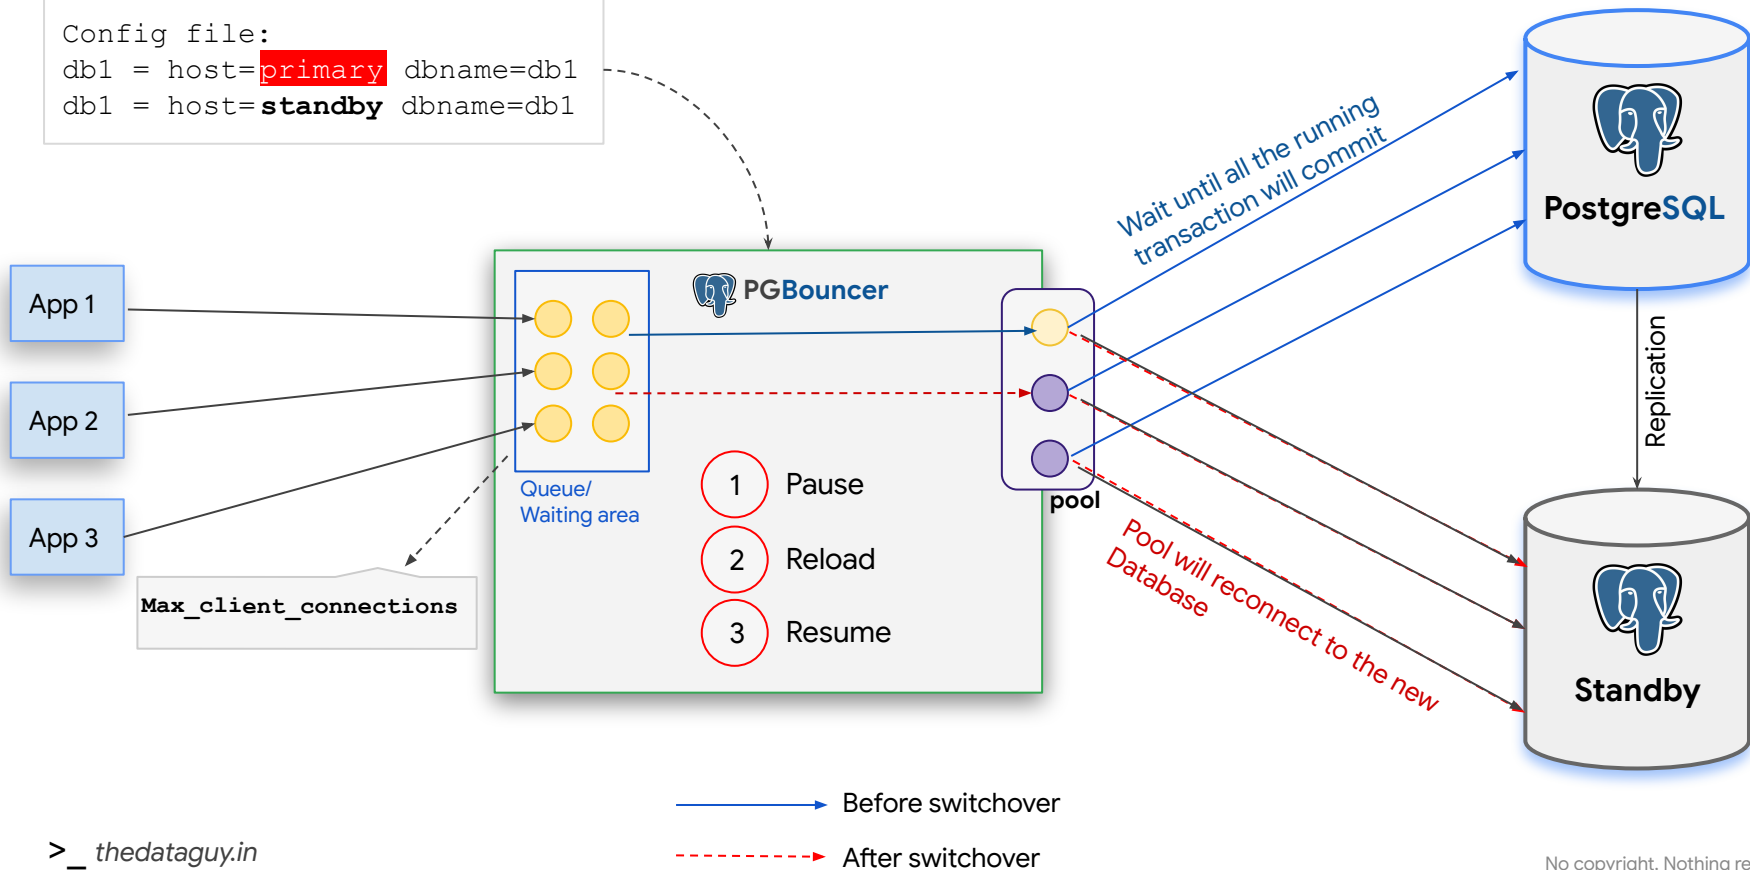
Zero downtime Maintenance



Seamless switchover to a standby

Config file:

```
db1 = host=primary dbname=db1  
db1 = host=standby dbname=db1
```



Restart without an actual restart

-R, --reboot

Do an online restart. That means connecting to the running process, loading the open sockets from it, and then using them. If there is no active process, boot normally.


Use this in your **systemctl** or **systemd** file:

```
pgbouncer -R /etc/pgbouncer/pgbouncer.ini
```

Online Restart Flow

1. Create new pgbouncer instance, make server connections
2. Send SUSPEND command to the old pgbouncer instance
3. Transfer the pool from old instance to new instance
4. Shutdown the old instance
5. Resume the work on the new instance

```
[LOG] takeover_init: launching connection
[LOG] S-0x55bbf85104b0: pgbouncer/pgbouncer@unix:6432 new connection to server
[LOG] S-0x55bbf85104b0: pgbouncer/pgbouncer@unix:6432 login OK, sending SUSPEND
[LOG] SUSPEND finished, sending SHOW FDS
[LOG] got pooler socket: 127.0.0.1:6432
[LOG] got pooler socket: unix:6432
[LOG] C-0x55bbf8517230: postgres/user1@127.0.0.1:59152 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8517460: postgres/user1@127.0.0.1:59162 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8517690: postgres/user1@127.0.0.1:59166 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf85178c0: postgres/user1@127.0.0.1:59168 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8517af0: postgres/user1@127.0.0.1:59170 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8517d20: postgres/user1@127.0.0.1:59172 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8517f50: postgres/user1@127.0.0.1:59174 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf8518180: postgres/user1@127.0.0.1:59176 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf85183b0: postgres/user1@127.0.0.1:59178 login attempt: db=postgres user=user1 tls=no
[LOG] C-0x55bbf85185e0: postgres/user1@127.0.0.1:59180 login attempt: db=postgres user=user1 tls=no
[LOG] SHOW FDS finished
[LOG] diskio over, going background
[LOG] kernel file descriptor limit: 1024 (hard: 1048576); max_client_conn: 1000, max expected fd use: 1072
[LOG] sending SHUTDOWN;
[LOG] S-0x55bbf85104b0: pgbouncer/pgbouncer@unix:6432 closing because: diskio over (age=10s)
[LOG] waiting for old pidfile to go away
[LOG] old process killed, resuming work
```





Monitoring

CLI and Prometheus exporter

Monitor via CLI

Pgbouncer config:

- Max pool = 300
- Max client = 800

Connect to pgbouncer db:

```
psql -h 127.0.0.1 -p 6432 -U  
user1 pgbouncer
```

Run a pgbench test:

```
pgbench -h 127.0.0.1 \  
-U user1 -C -c 800 \  
-j 2 \  
-t 10000 \  
-p 6432 \  
postgres
```

> thdataguy.in

```
pgbouncer=# show pools;  
  
-[ RECORD 1 ]-----  
database      | postgres  
user          | user1  
cl_active     | 216  
cl_waiting    | 582  
sv_active     | 216  
sv_idle       | 0  
sv_used       | 0  
sv_tested     | 0  
sv_login      | 0  
maxwait       | 12  
maxwait_us    | 880893  
pool_mode     | session
```

Monitoring with the prometheus exporter and Grafana

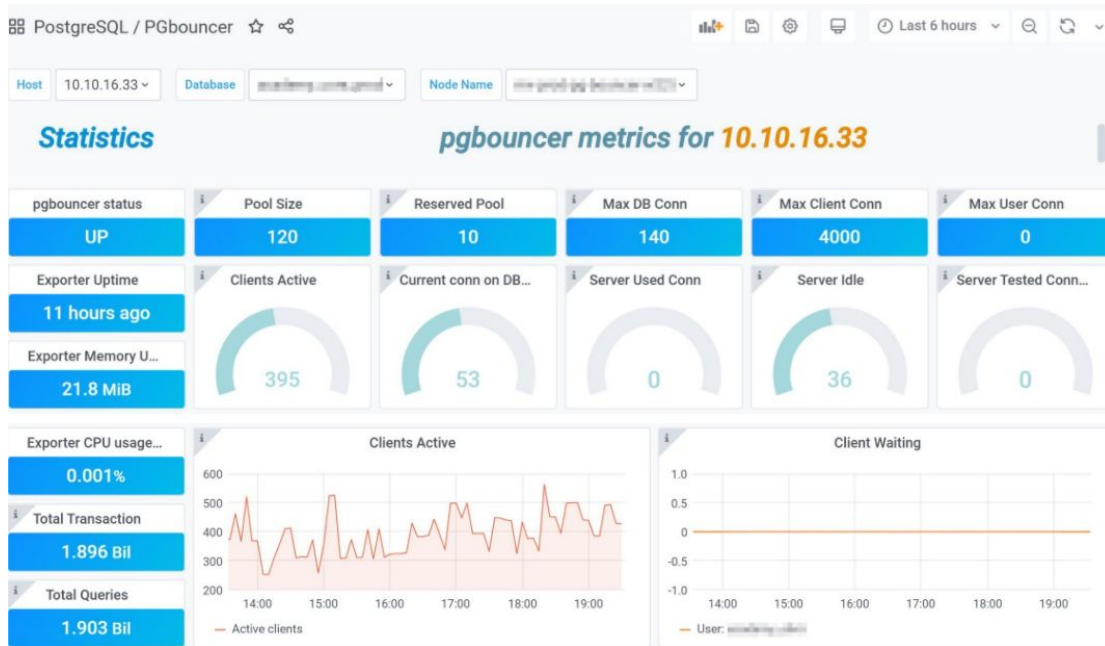
Prometheus exporter:

```
pip3 install  
prometheus-pgbouncer-exporter
```

Grafana Dashboard:

The dashboard is open source(we
developed this dashboard)


Dashboard ID: 13353




Reference: <https://medium.com/search/grafana-dashboard-for-pgbouncer-and-monitor-with-percona-pmm-3170d3eb4d14>

Pgbouncer deployment pattern


Pattern 1




Pattern 2




Pattern 3



Pattern 4



Pattern 5



What I didn't cover?

1. Timeout parameters
2. High availability of pgbouncer
3. TLS
4. `server_reset_query` - an important parameter
5. And more

Thank you !!!

And any Questions?

Get this deck here:
bit.ly/pgbouncer-deck

