

PostgreSQL Distributed & Secure Database Ecosystem Building

Trista Pan, panjuan@apache.org

»» About Me



Trista Pan

SphereEx Co-Founder & CTO

Apache Member

AWS Data Hero

Tencent Cloud TVP

Apache ShardingSphere PMC

Apache brpc (Incubating) & Apache AGE (Incubating) mentor

China Mulan Community Mentor

Bio: <https://tristazero.github.io>

LinkedIn: <https://www.linkedin.com/in/panjuan>

GitHub: <https://github.com/tristaZero>

Twitter: @tristaZero

Project Twitter: @ShardingSphere

» Contents

- Why You're Expecting More From PostgreSQL
- What is Apache ShardingSphere?
- What is the Value of this Project?
- Distributed SQL's Kernel Processing
- Distributed PostgreSQL Database Solution
- Data Encryption with PostgreSQL

PostgreSQL: The World's Most Advanced Open Source Relational Database

Highlights

Proven
Architecture
Reliability
Robust Features
Set
Extensibility
Open Source


Distributed
Solution
Cluster Management
Observability
Data Encryption

New Needs

Apache ShardingSphere Community Overview

Why Open Source Software

- Build new software easily and fast.
- Prevent vendor lock effectively.
- Proved to be secure and reliable.
- Fast Upgrade & Low Switching Cost
- Sound community ecosystem & quality software.



- The World's Largest Open Source Foundation
- 227M+ Lines of Code in Stewardship
- 300+ Top-Level Projects



• **15000+** Stars

• **5000+** Forks

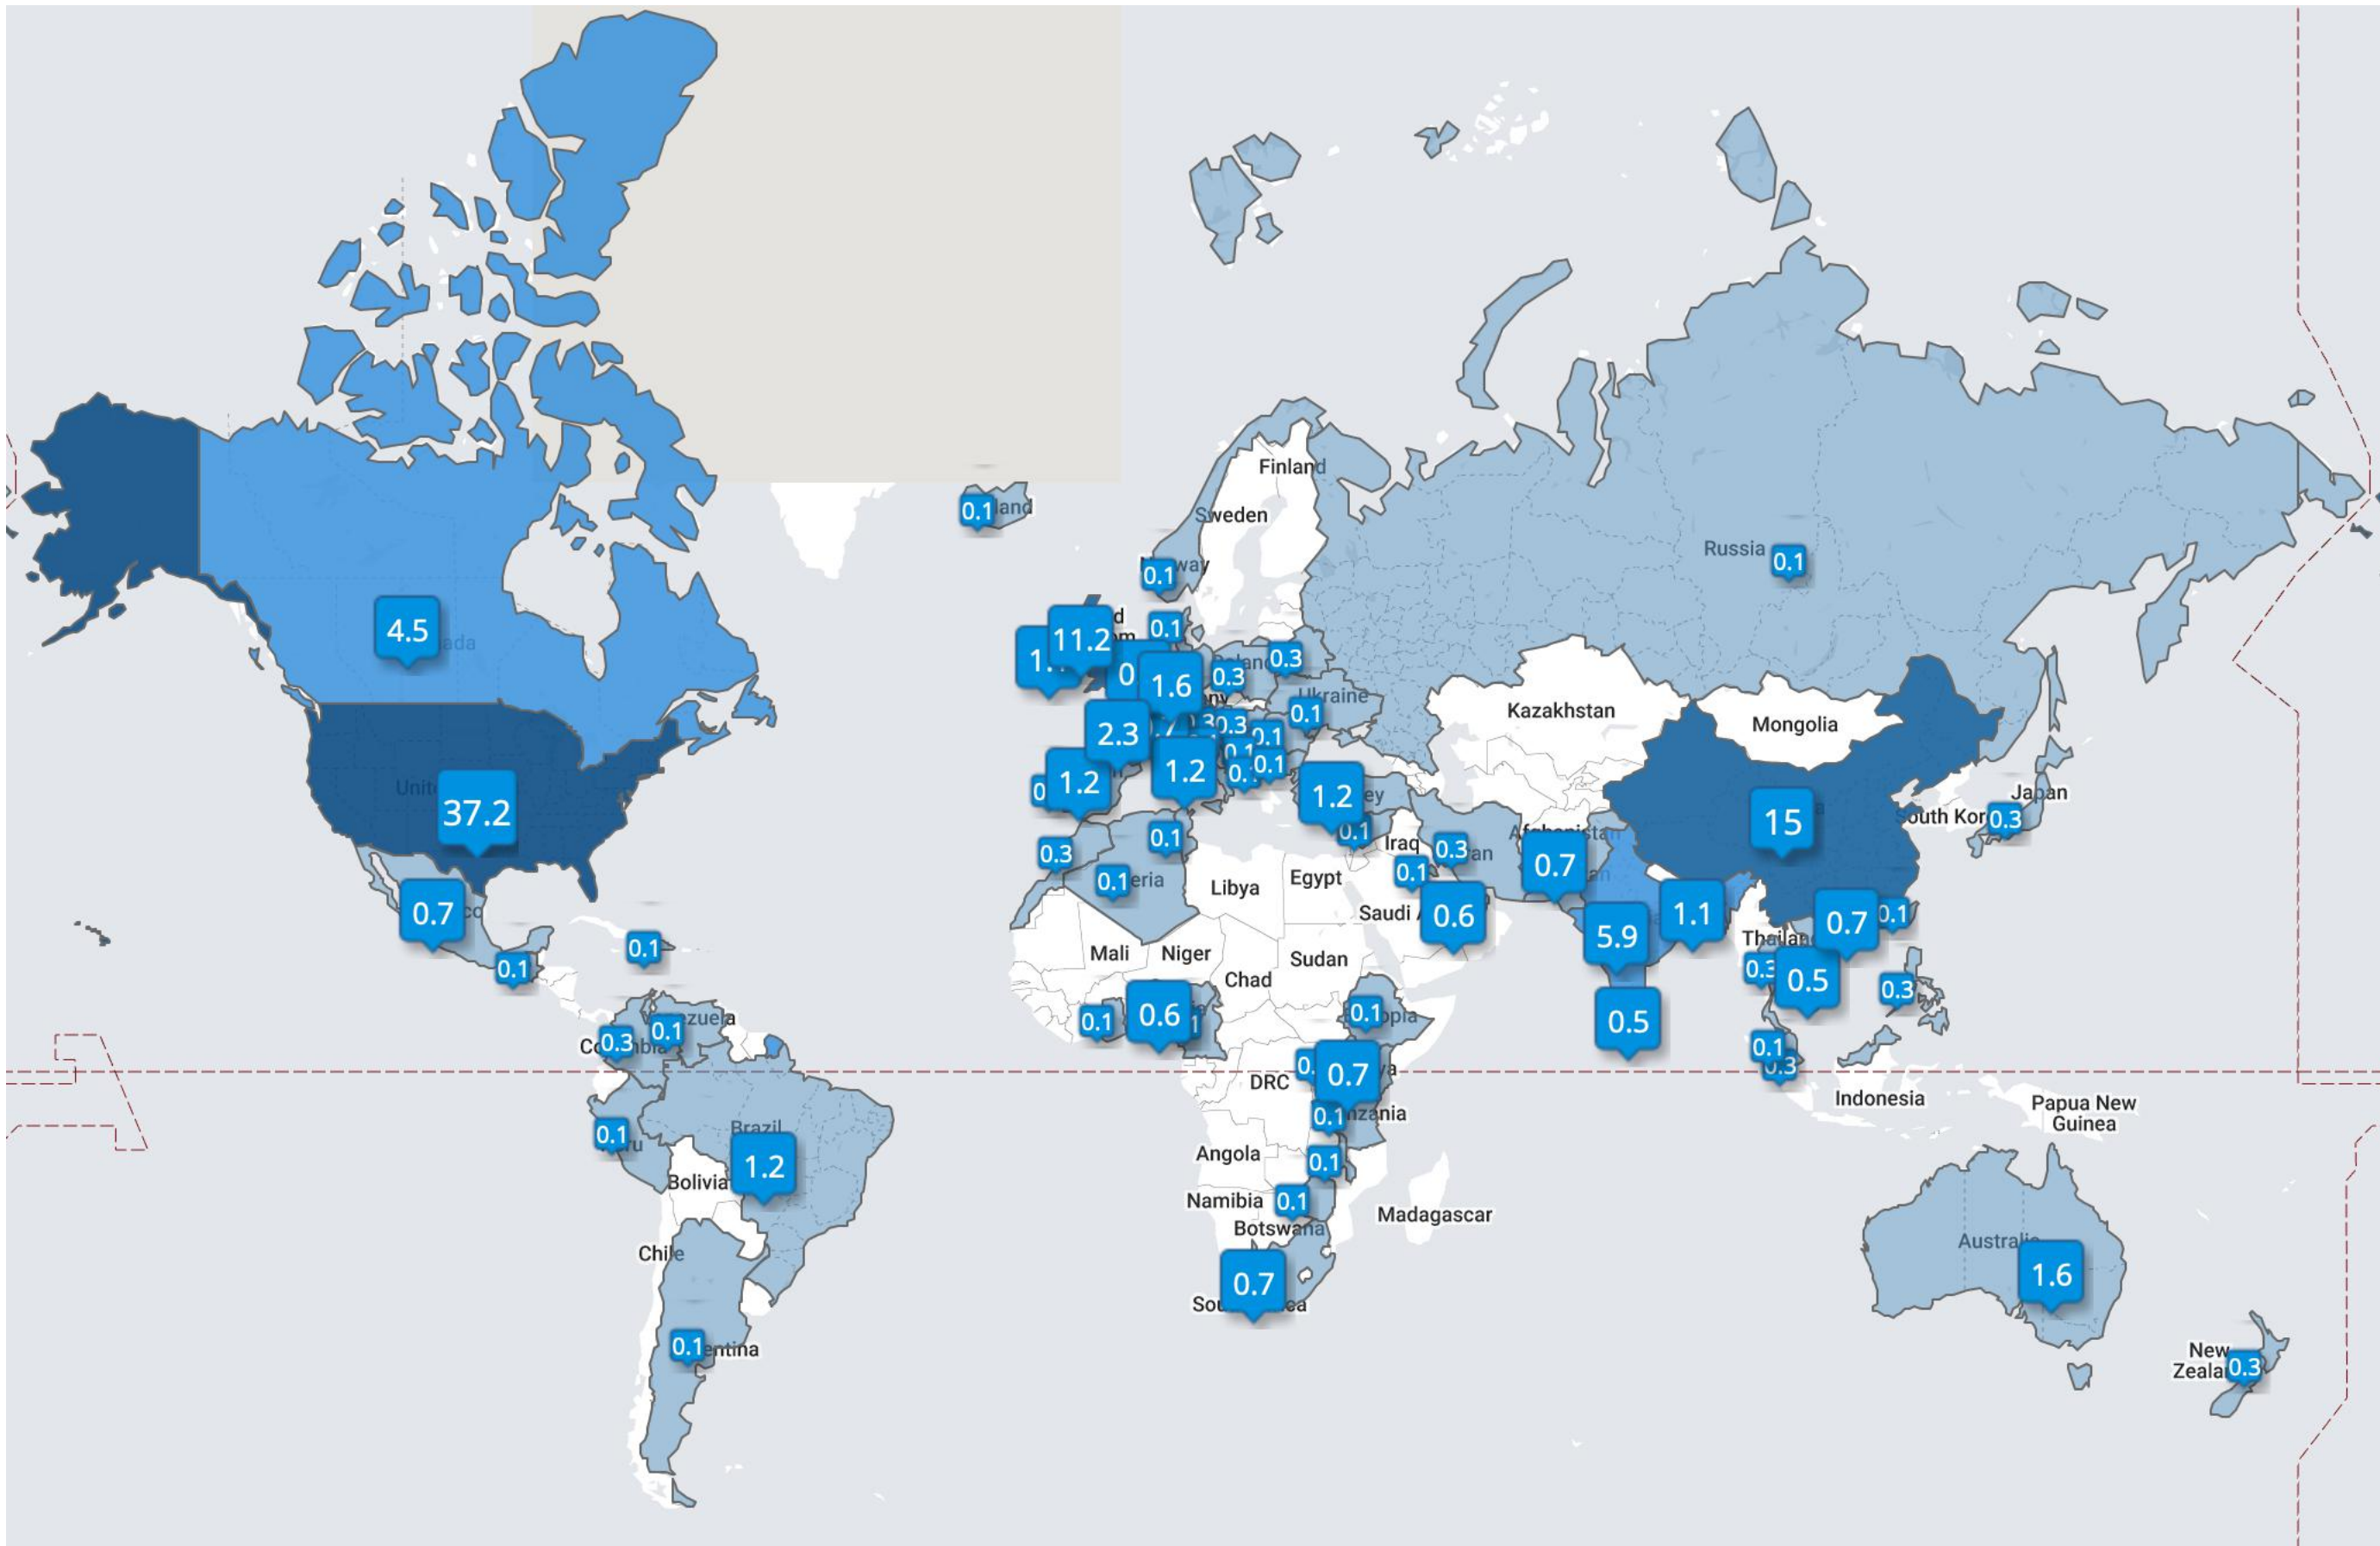
• **300+** Contributors

• **8000+** Pull Requests


- ShardingSphere is one of the Top-Level Apache Projects.
- Top Chinese Led Apache Project, 2020
- Top 10 Apache Project by Number of Commits, 2021

Apache ShardingSphere's Global Development


Twitter Social Followers



Stargazers Over Time



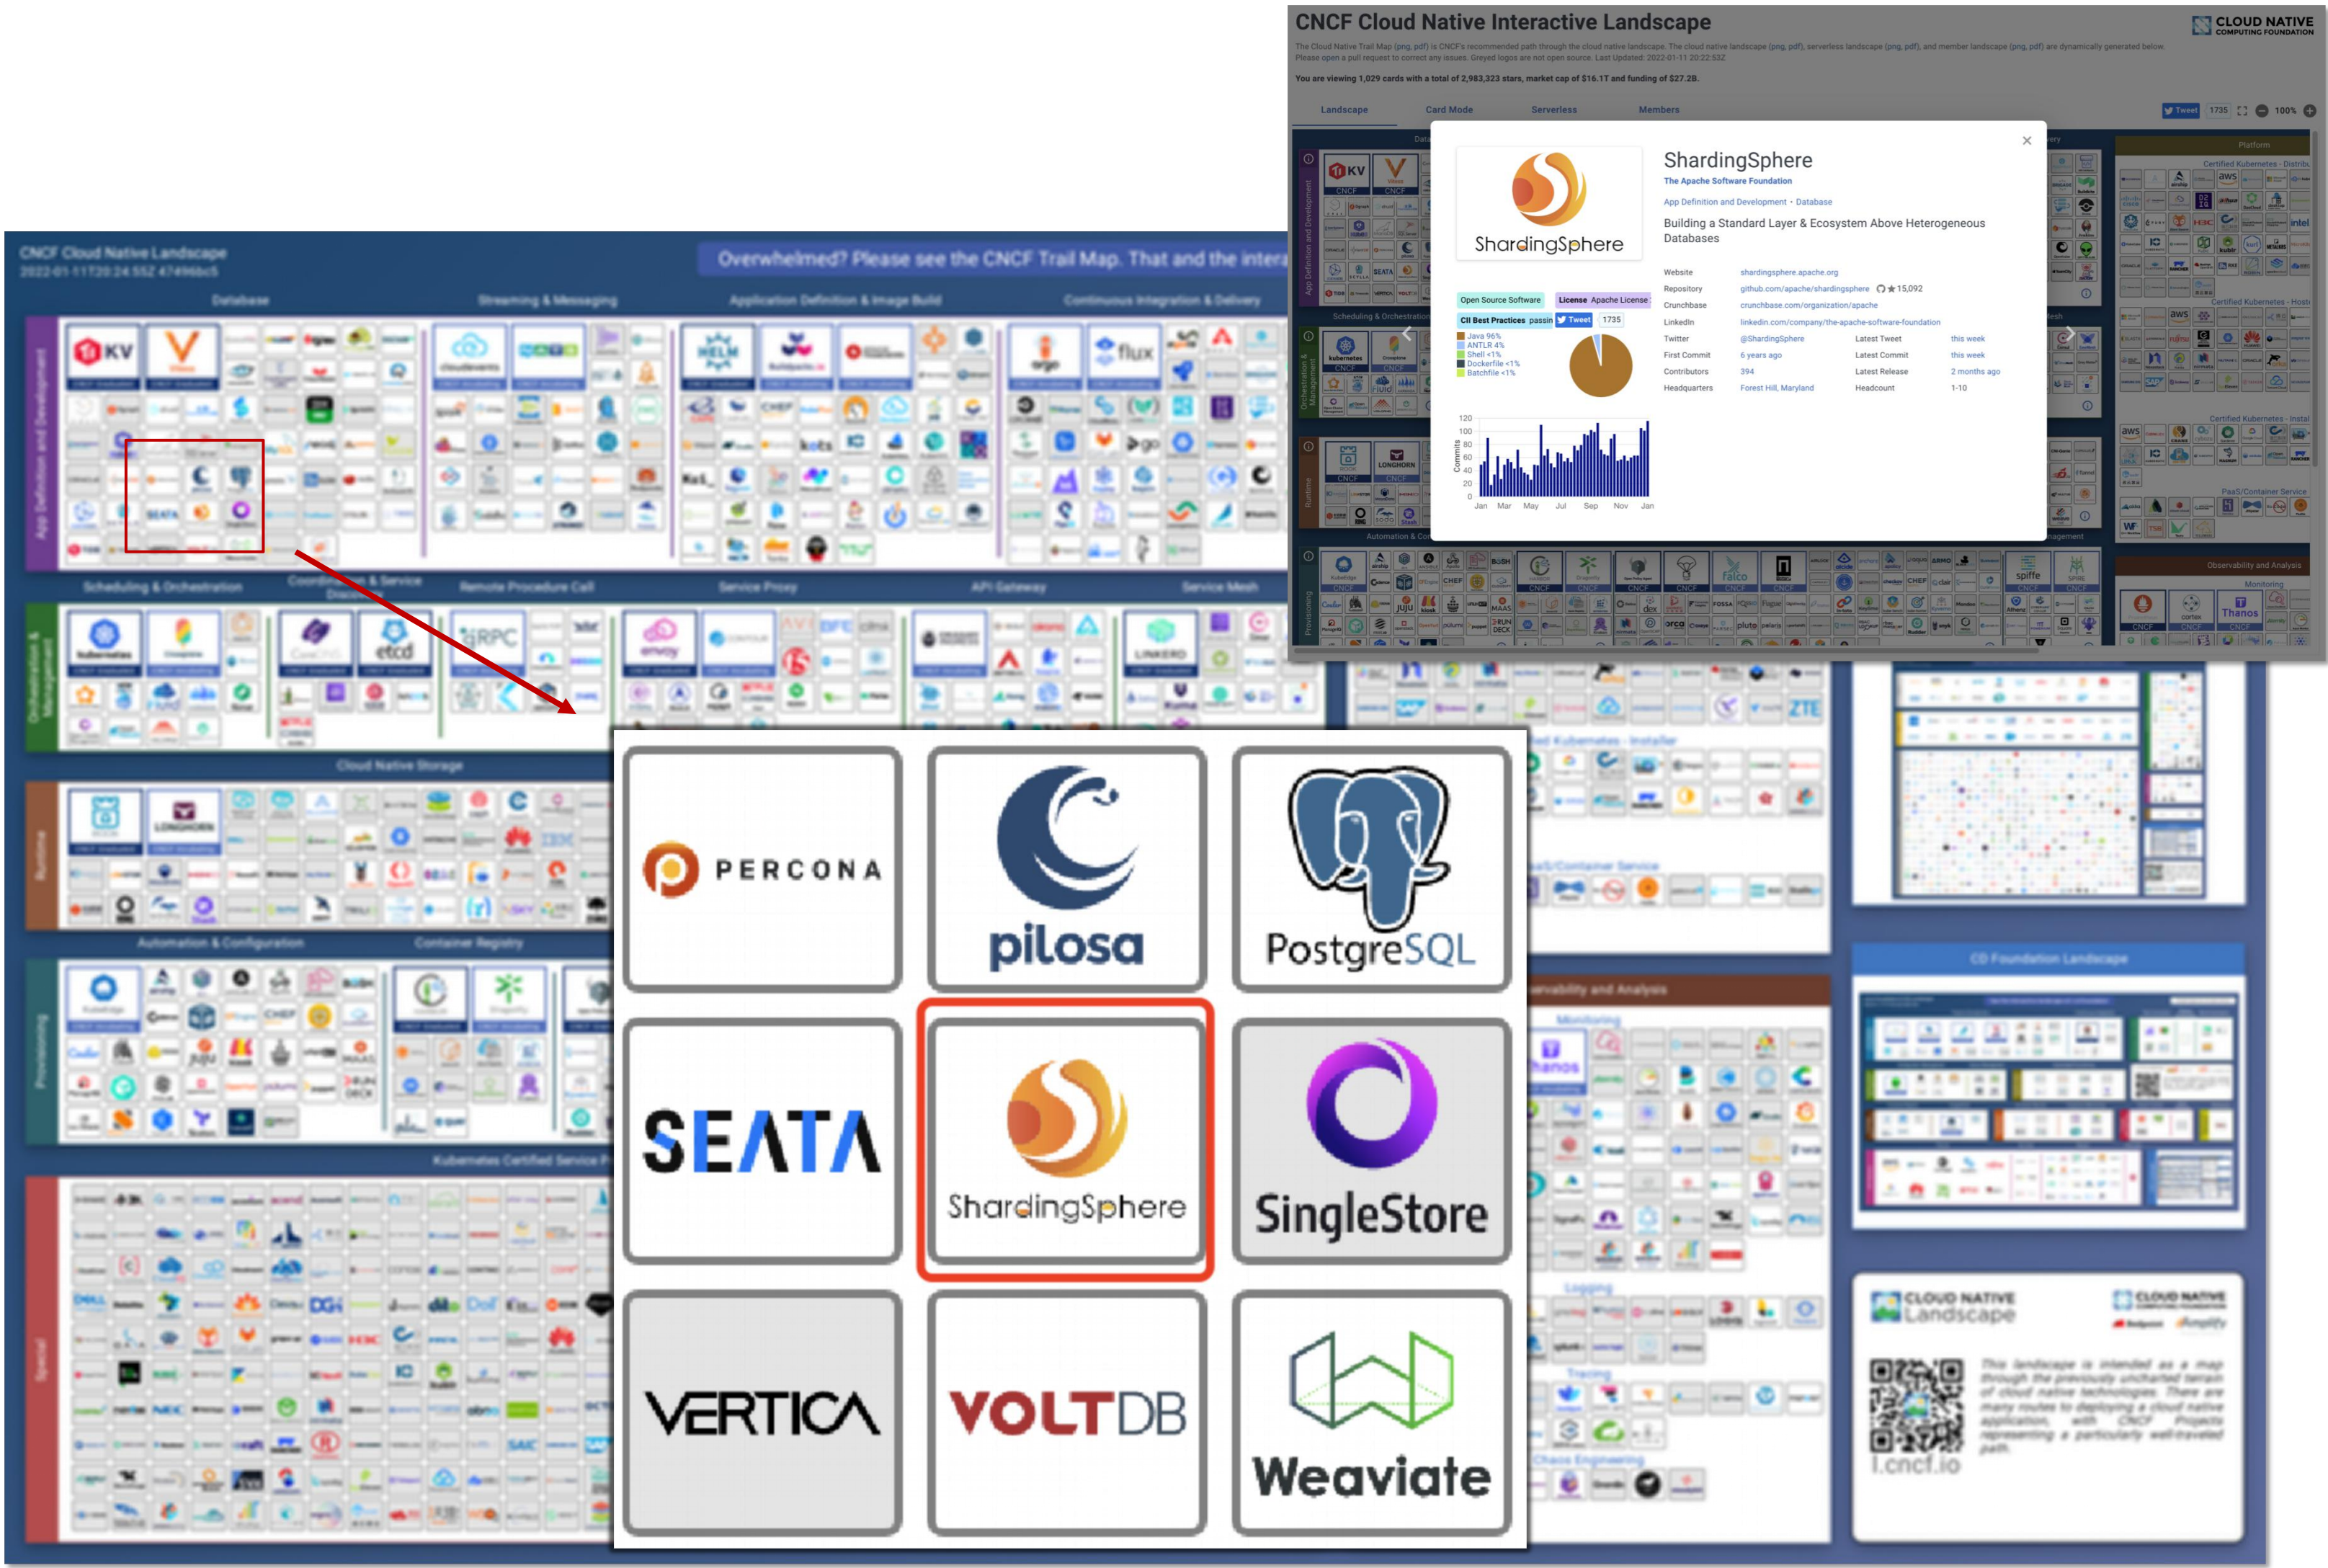
Contributors Over Time



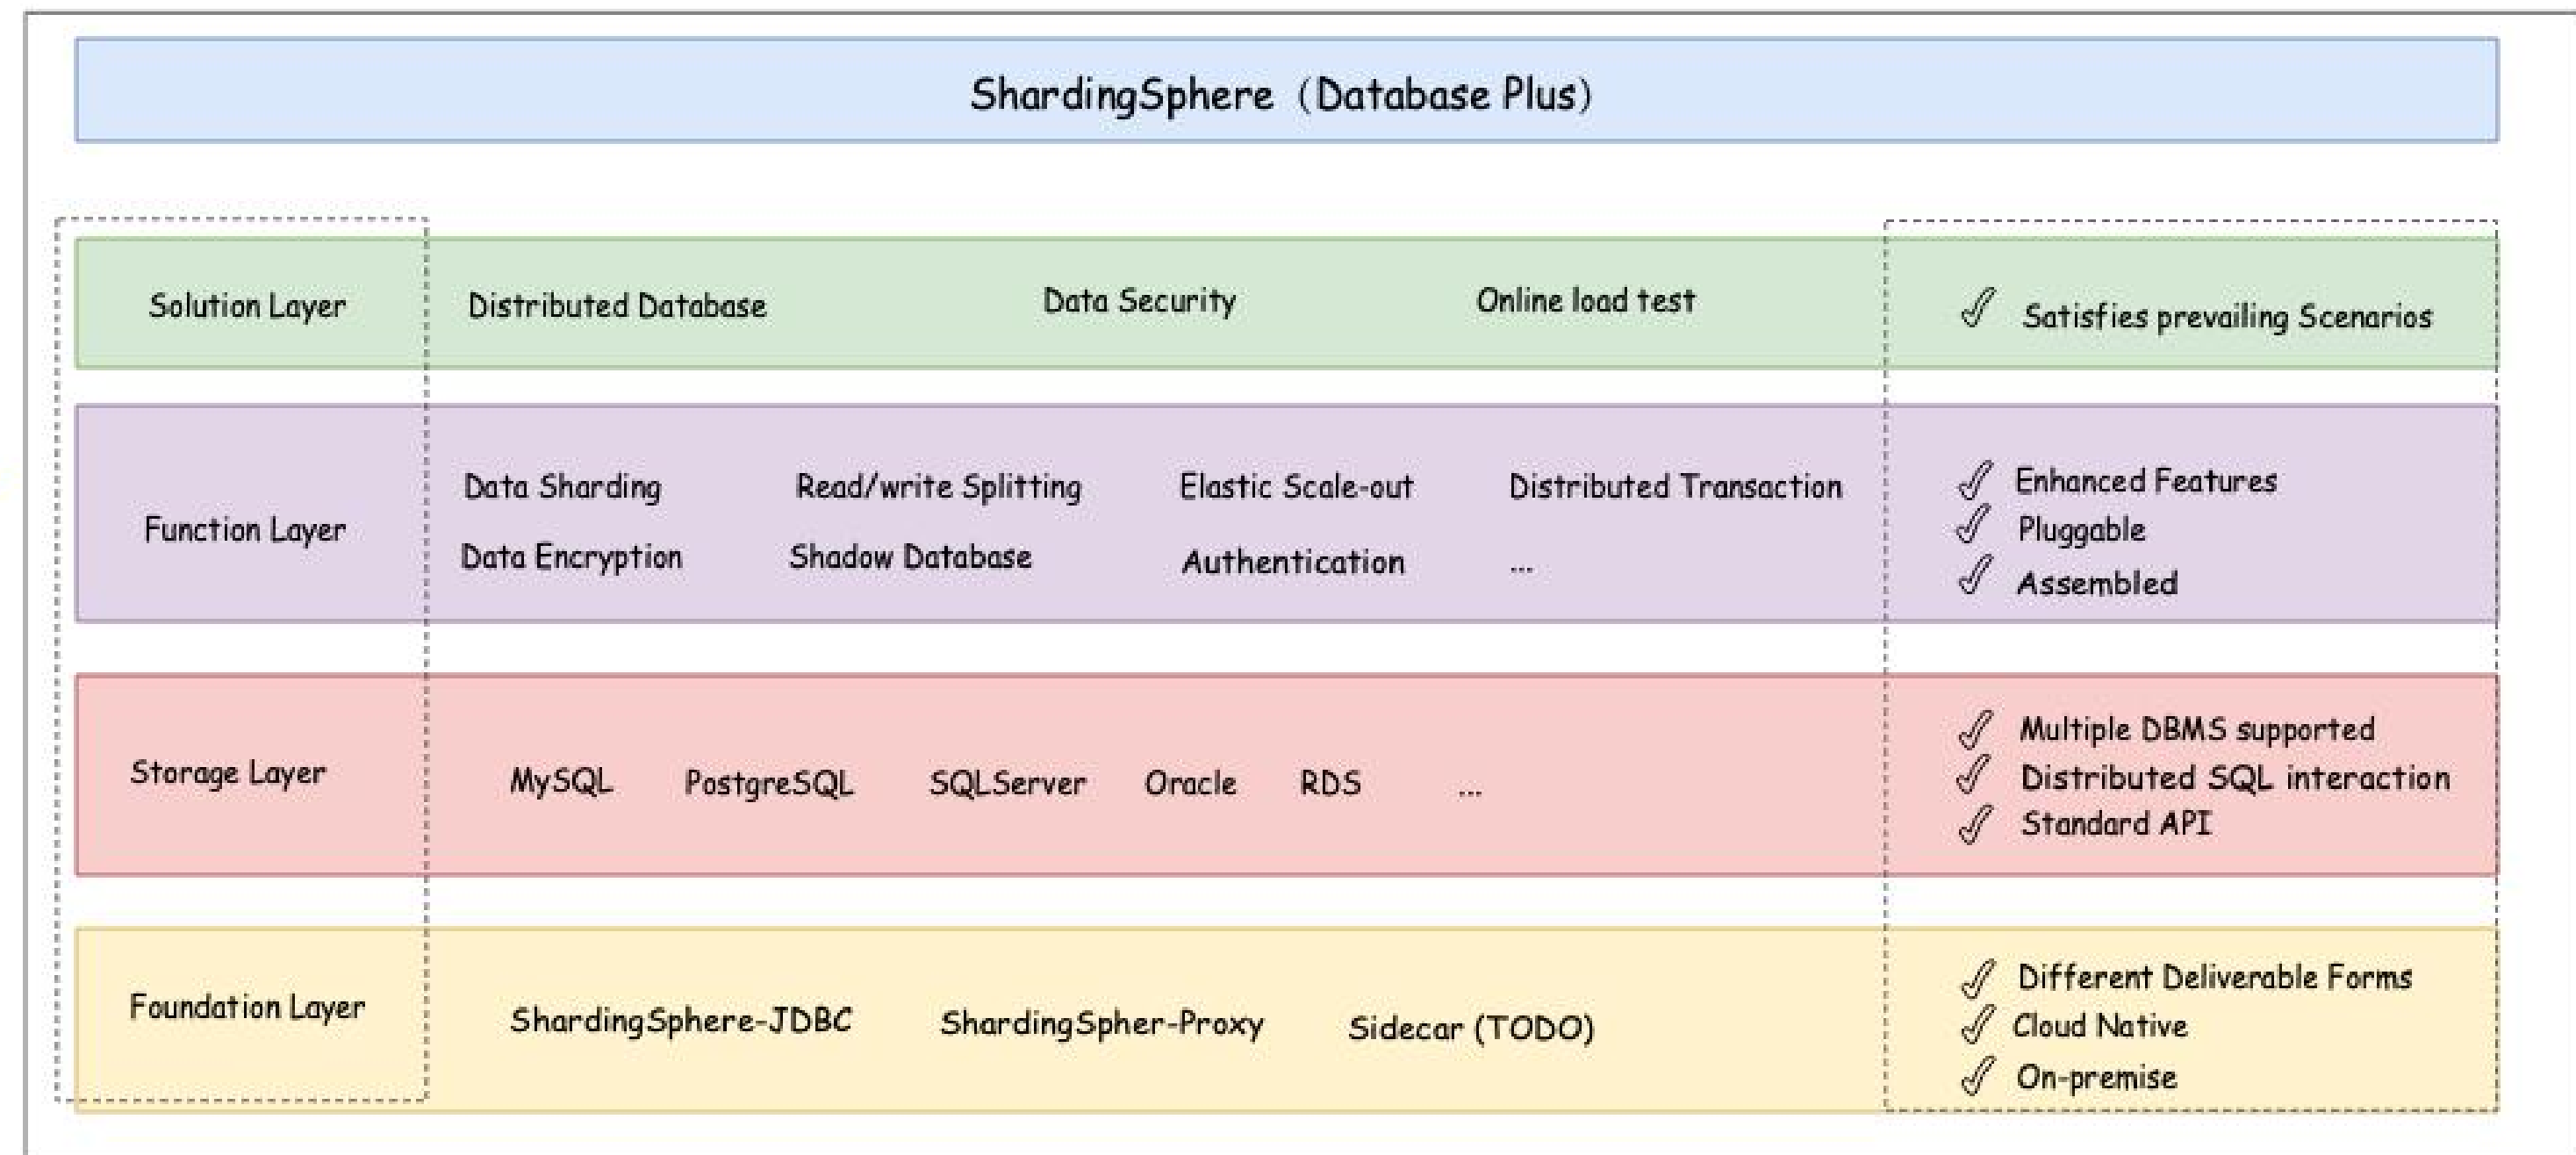
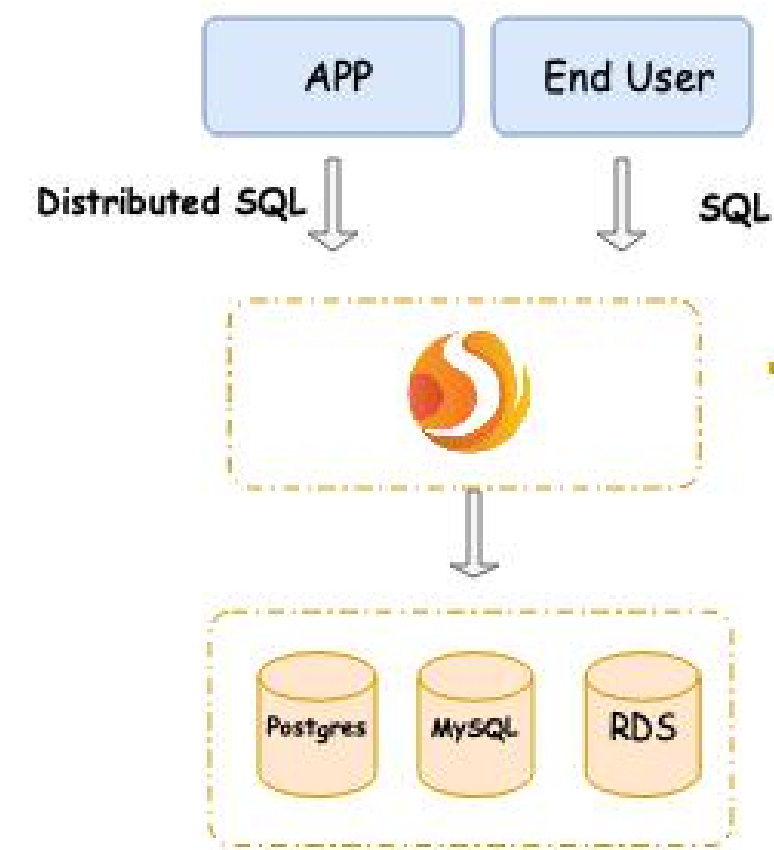
Apache ShardingSphere GitHub Stargazers & Contributors Over Time

CNCF : Cloud Native Computing Foundation

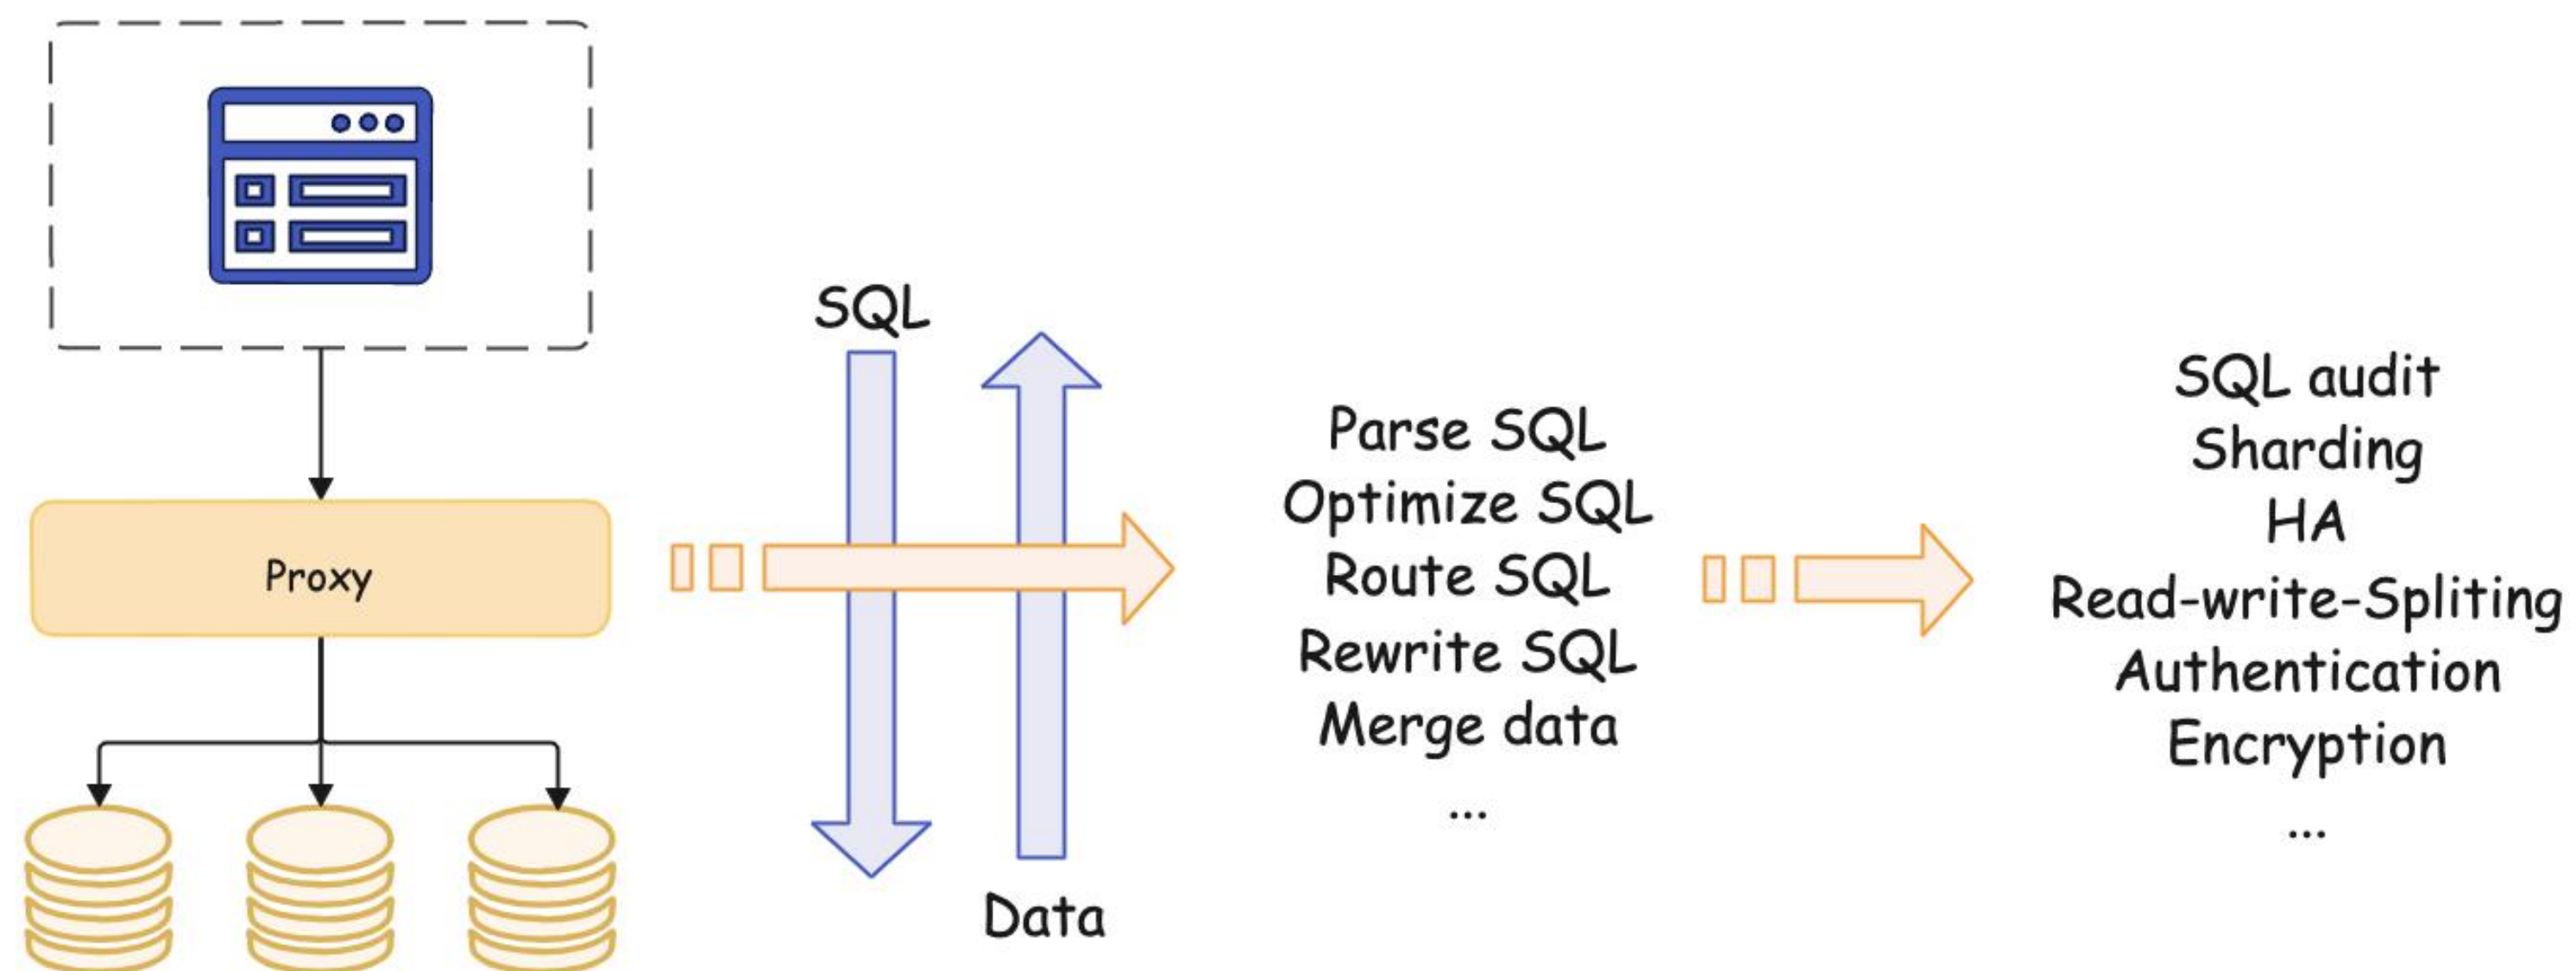
- As a part of the Linux Foundation, CNCF's mission is to make cloud native computing ubiquitous.
- Currently listed are 1,029 cards with a total of 2,984,816 Stars, a market cap of \$16.3T and funding of \$27.2B.



» Apache ShardingSphere


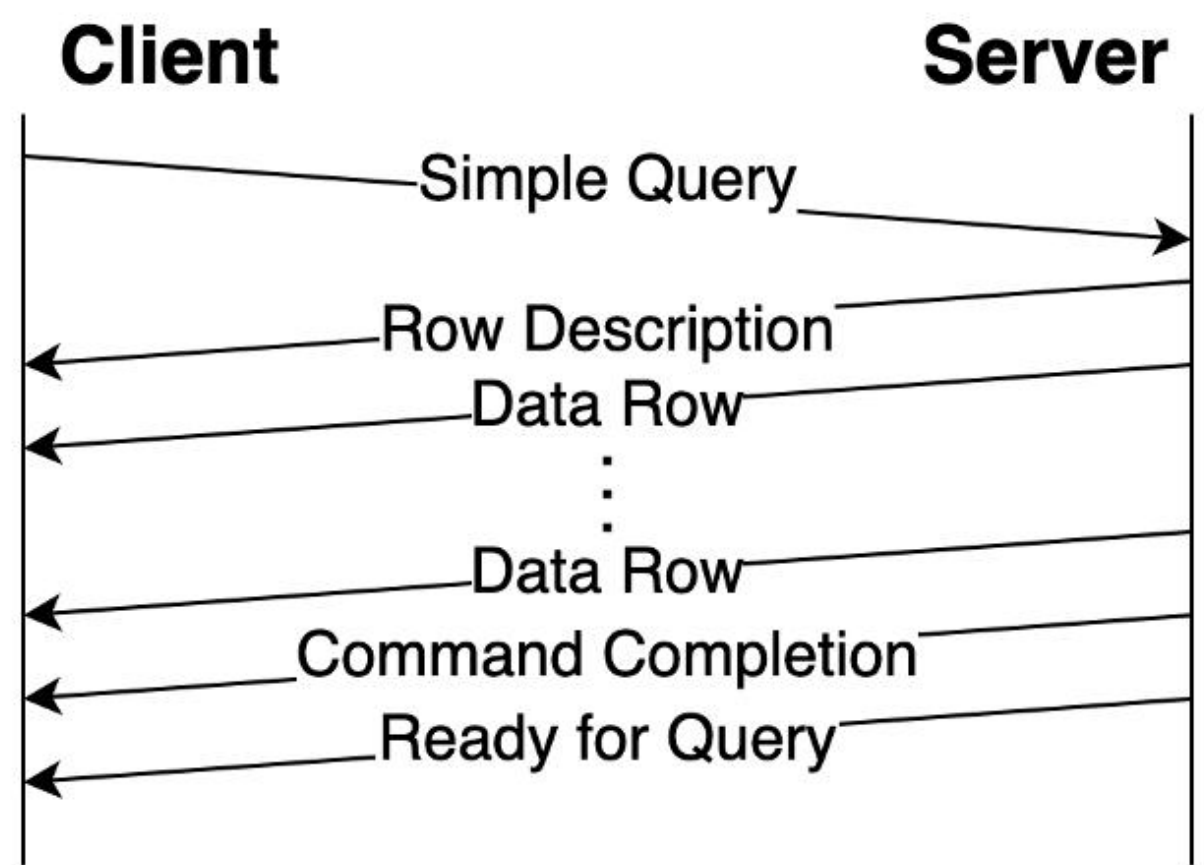


Kernel Process



» Kernel Process

```
biz_order=# select order_id, create_time from t_order limit 3;
 order_id | create_time
-----+-----
 614158771245527040 | 2021-09-01 10:32:44.893
 614127204259311616 | 2021-09-01 07:55:18.471
 608625285797490688 | 2021-08-31 18:17:42.173
(3 rows)
```

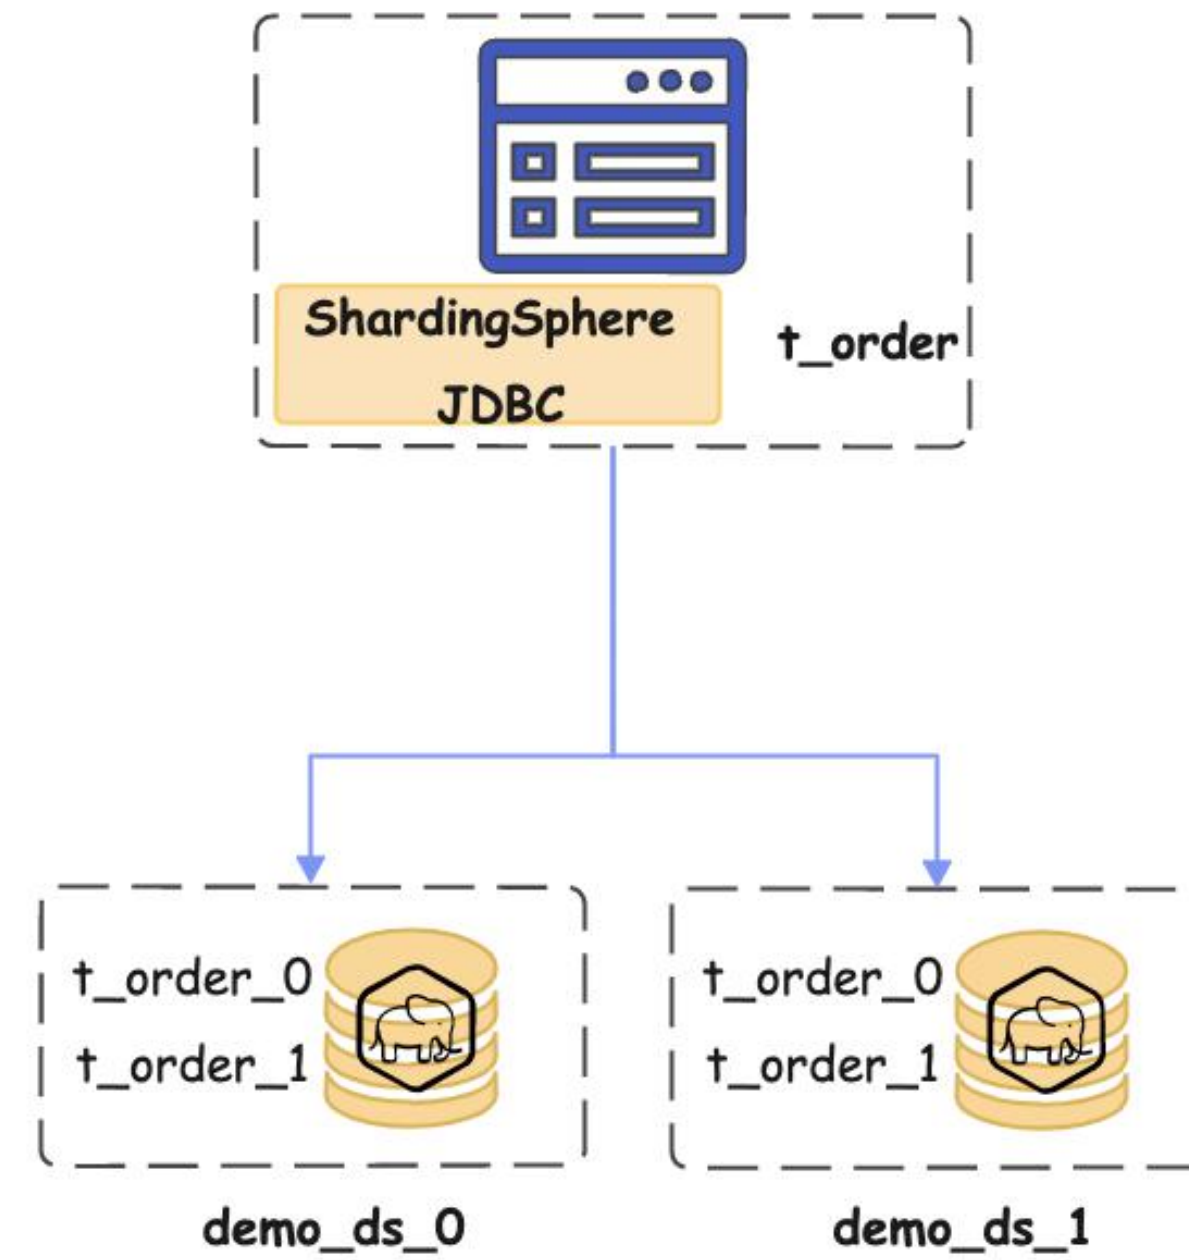
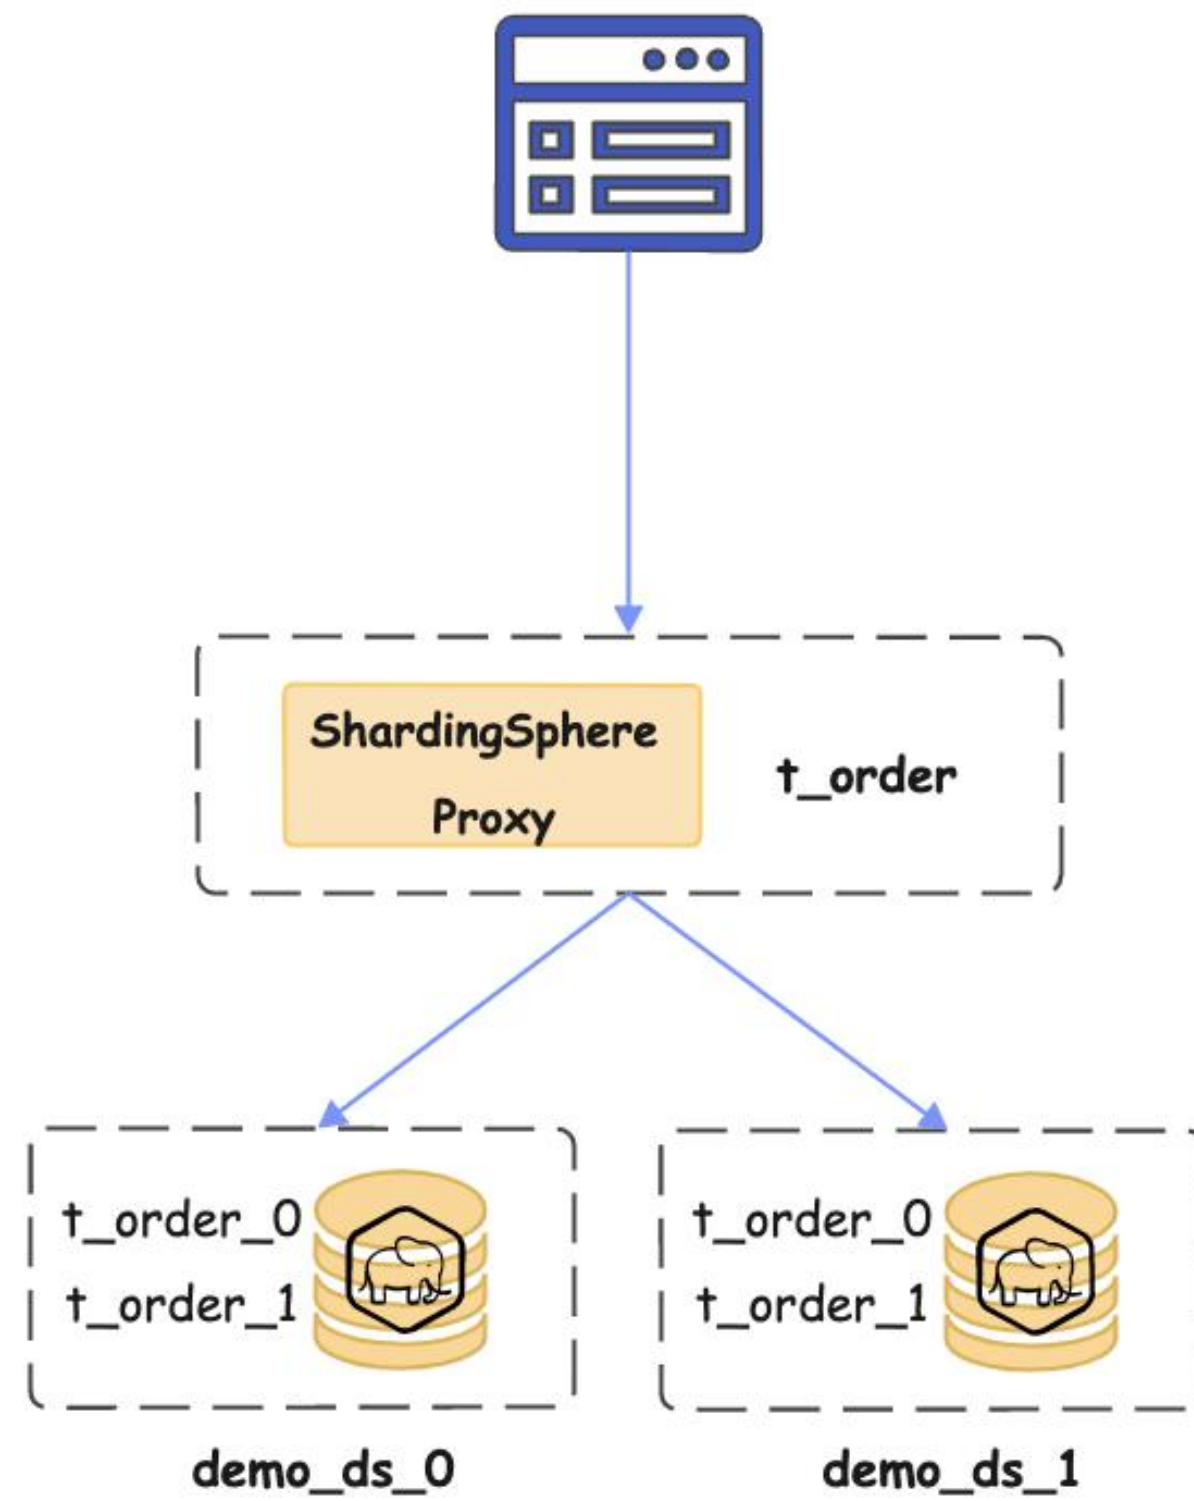


▼ PostgreSQL

Type: Simple query
Length: 14
Query: select 1;

0000	02 00 00 00 45 00 00 43	00 00 40 00 40 06 00 00E..C ..@..@...
0010	7f 00 00 01 7f 00 00 01	f7 34 15 38 d2 e2 a4 534.8...S
0020	ad cb 8e 78 80 18 18 dd	fe 37 00 00 01 01 08 0a	...x....7.....
0030	ed 8c c4 91 ad 89 67 39	51 00 00 00 0e 73 65 6cg9 Q....sel
0040	65 63 74 20 31 3b 00		ect 1;.

» Distributed Postgres Database Solution



» Distributed Postgres Database Solution

1. Create Database


```
$ psql -Uroot -h127.0.0.1 -ddemos -p3307
Password for user root:
psql (14.1, server 9.6.24-ShardingSphere-Proxy 5.0.1-SNAPSHOT-ddf0e2b)
Type "help" for help.

demos=> create database demo_ds;
CREATE DATABASE
demos=> exit
```

2. Add Data Source

```
demo_ds=> ADD RESOURCE demo_ds_0 (
demo_ds(>     HOST=127.0.0.1,
demo_ds(>     PORT=15432,
demo_ds(>     DB=demo_ds_0,
demo_ds(>     USER=postgres,
demo_ds(>     PASSWORD=postgres
demo_ds(> ), demo_ds_1 (
demo_ds(>     HOST=127.0.0.1,
demo_ds(>     PORT=25432,
demo_ds(>     DB=demo_ds_1,
demo_ds(>     USER=postgres,
demo_ds(>     PASSWORD=postgres
demo_ds(> );
CREATE
```


» Distributed Postgres Database Solution



3. Create Rules & Tables

```
demo_ds=> CREATE SHARDING ALGORITHM table_inline (  
demo_ds(> TYPE(NAME=inline,PROPERTIES("algorithm-expression"="t_order_${user_id % 2}"))  
demo_ds(> );  
  
demo_ds=> CREATE SHARDING TABLE RULE t_order (  
demo_ds(> DATANODES("demo_ds_${0..1}.t_order_${0..1}"),  
demo_ds(> DATABASE_STRATEGY(TYPE=standard,SHARDING_COLUMN=order_id,SHARDING_ALGORITHM(TYPE(NAME=inline,PROPERTIES("algorithm-expression"="demo_ds_${order_id % 2}")))),  
demo_ds(> TABLE_STRATEGY(TYPE=standard,SHARDING_COLUMN=user_id,SHARDING_ALGORITHM=table_inline)  
demo_ds(> );  
CREATE
```

```
demo_ds=> CREATE TABLE t_order (  
demo_ds(>   order_id bigint NOT NULL,  
demo_ds(>   user_id int NOT NULL,  
demo_ds(>   status varchar(50) DEFAULT NULL,  
demo_ds(>   PRIMARY KEY (order_id)  
demo_ds(> );  
CREATE TABLE
```

4. Insert Data

```
demo_ds=> INSERT INTO t_order values(1,1,'OK');  
INSERT 0 1  
demo_ds=> INSERT INTO t_order values(2,1,'OK');  
INSERT 0 1  
demo_ds=> INSERT INTO t_order values(3,2,'OK');  
INSERT 0 1  
demo_ds=> INSERT INTO t_order values(4,2,'OK');  
INSERT 0 1  
demo_ds=> SELECT * FROM t_order order by order_id;  
  order_id | user_id | status  
-----+-----+-----  
         1 |         1 | OK  
         2 |         1 | OK  
         3 |         2 | OK  
         4 |         2 | OK  
(4 rows)
```

» Distributed Postgres Database Solution

5. Preview SQL

```
demo_ds=> preview select * from t_order order by order_id;
data_source_name |          sql
-----+-----
demo_ds_0        | select * from t_order_0 order by order_id
demo_ds_0        | select * from t_order_1 order by order_id
demo_ds_1        | select * from t_order_0 order by order_id
demo_ds_1        | select * from t_order_1 order by order_id
(4 rows)
```

6. The Data in Actual PostgreSQL

demo_ds_0

```
demo_ds_0=# select * from t_order_0;
order_id | user_id | status
-----+-----+-----
4 | 2 | OK
(1 row)

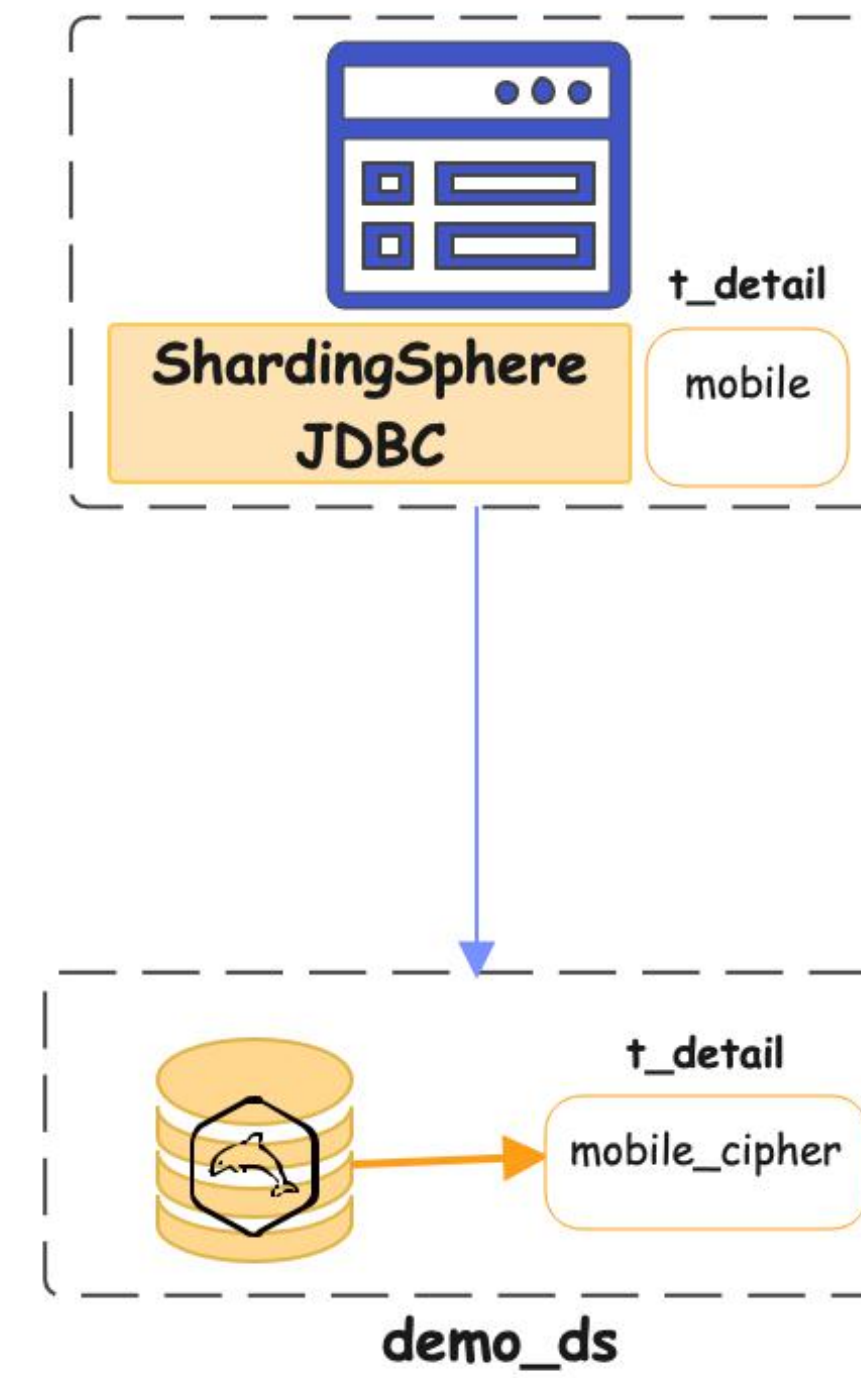
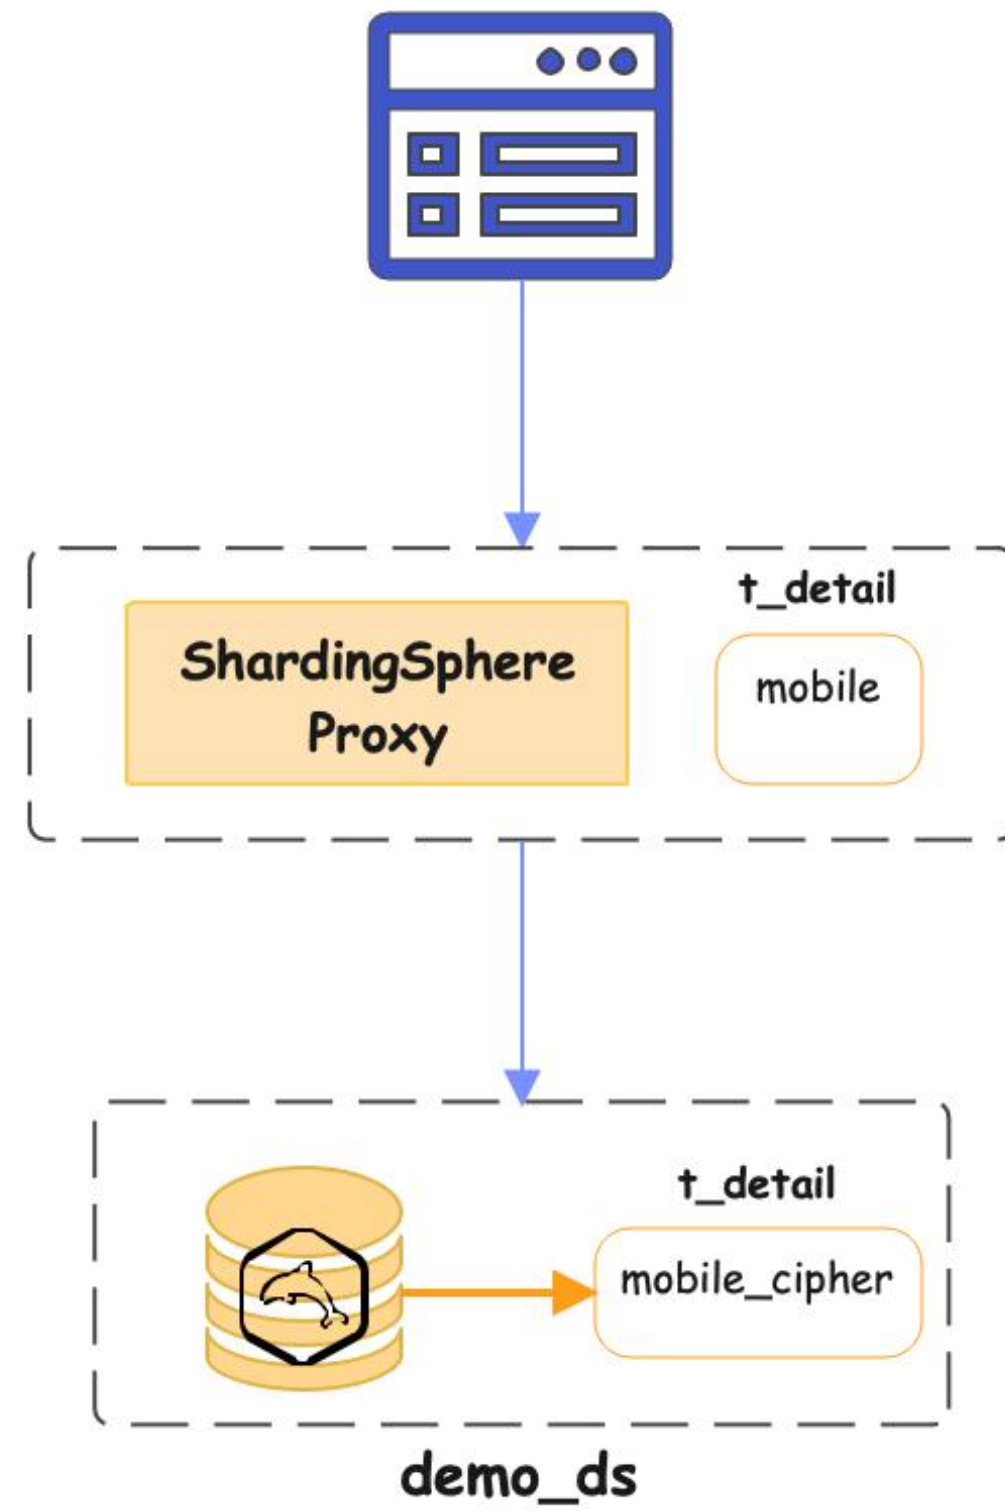
demo_ds_0=# select * from t_order_1;
order_id | user_id | status
-----+-----+-----
2 | 1 | OK
(1 row)
```

demo_ds_1

```
demo_ds_1=# select * from t_order_0;
order_id | user_id | status
-----+-----+-----
3 | 2 | OK
(1 row)

demo_ds_1=# select * from t_order_1;
order_id | user_id | status
-----+-----+-----
1 | 1 | OK
(1 row)
```


Secure Postgres Database Solution



Secure Postgres Database Solution

1. Create Database

```
$ psql -Uroot -h127.0.0.1 -ddemos -p3307
Password for user root:
psql (14.1, server 9.6.24-ShardingSphere-Proxy 5.0.1-SNAPSHOT-ddf0e2b)
Type "help" for help.

demos=> create database demo_ds;
CREATE DATABASE
demos=> exit
```

2. Add Data Source

```
demo_ds=> ADD RESOURCE ds_0 (
demo_ds(>   URL="jdbc:postgresql://127.0.0.1:15432/demo_ds"
demo_ds(>   USER=postgres,
demo_ds(>   PASSWORD=postgres
demo_ds(> );
CREATE
```

3. Create Rule

```
demo_ds=> CREATE ENCRYPT RULE t_detail (COLUMNS( (NAME=idCard,CIPHER=idcard_cipher,TYPE(NAME=AES,PROPERTIES('aes-key-value'='123456abc'))),
demo_ds(>   (NAME=mobile, CIPHER =mobile_cipher,TYPE(NAME=AES,PROPERTIES('aes-key-value'='123456abc')))));
```


Secure Postgres Database Solution

4. Create Table

```
demo_ds=> CREATE TABLE t_detail (id INT, mobile VARCHAR(50), idcard VARCHAR(50));
CREATE TABLE
demo_ds=> █
```

5. Insert Data

```
demo_ds=> INSERT INTO t_detail (id, mobile, idcard)
demo_ds-> VALUES (1, 18236483857, 220605194709308170),
demo_ds->          (2, 15686689114, 360222198806088804),
demo_ds->          (3, 14523360225, 411601198601098107),
demo_ds->          (4, 18143924353, 540228199804231247),
demo_ds->          (5, 15523349333, 360924195311103360),
demo_ds->          (6, 13261527931, 513229195302236086),
demo_ds->          (7, 13921892133, 500108194806107214),
demo_ds->          (8, 15993370854, 451322194405305441),
demo_ds->          (9, 18044280924, 411329199808285772),
demo_ds->          (10, 13983621809, 430204195612042092);
INSERT 0 10
demo_ds=>
demo_ds=> select * from t_detail;
 id |  mobile  |  idcard
-----+-----+-----
  1 | 18236483857 | 220605194709308170
  2 | 15686689114 | 360222198806088804
  3 | 14523360225 | 411601198601098107
  4 | 18143924353 | 540228199804231247
  5 | 15523349333 | 360924195311103360
  6 | 13261527931 | 513229195302236086
  7 | 13921892133 | 500108194806107214
  8 | 15993370854 | 451322194405305441
  9 | 18044280924 | 411329199808285772
 10 | 13983621809 | 430204195612042092
(10 rows)
```

Secure Postgres Database Solution

6. Check Data from Proxy

```
demo_ds=>
demo_ds=> SELECT * FROM t_detail WHERE mobile = '18236483857';
 id |   mobile   |   idcard
-----+-----+-----
  1 | 18236483857 | 220605194709308170
(1 row)
```

7. The Data in Actual PostgreSQL

```
demo_ds=# SELECT * FROM t_detail ;
 id |   mobile_cipher   |   idcard_cipher
-----+-----+-----
  1 | p31Pk19nIunYdH+AngyNUA== | pQv0JEkM94QzktJdM8UMg/uLrU71G6n6DALdPp9w6L0=
  2 | CV8+uYRaW0zcTQnQX3RcwA== | dCF7k4haK0aIV/d7dtwgzIb4lIFlJ913hrPim1+J278=
  3 | jnfu7o44KgN/PV1zhiu7jw== | 8iulp3+XTSv2XHGUUHKV0UsLuFx7yEpQVT+47EFfg94=
  4 | ZJDrTv/XIjdqdG1yp0t95w== | iqU6myMGfgI/XnxCtjhbMrwIauriWu8crxPS6BH2pMk=
  5 | FnQMYGnFJaiWmThENYzbFA== | KAPrCXoo1svMt5NWe0UaKYZIl1rSEVddHbBJ01jPIqw=
  6 | lv2ECfTCgQQksvdPp6k3Ug== | BBBPAuwU+iJluI9d9TA+H81BPnVXBaly1BE3Ep1N4e8=
  7 | z46vpnHCFTkIF2EtntxpHQ== | Bc39nPtyz1ji9Rc8k4f7G9CKfPew23mKFwp8guK7ybg=
  8 | p/IJdGcCikhpCu5gVZj4jg== | nnv/kS1i7uHXKncU0uLzE80WM0nGlcGkLokT2dltSaQ=
  9 | NvPcQv4w3EqD77+VAX0KCA== | +yeo5LWKNWcekFqYawCKjsctAZqe104DrI7AeZdR/Uk=
 10 | x0yg9E0X9lhy9mUx0QyL0A== | U7P1CMcxn6VPHYHPgTAtjHEbb6N6vhG0pdJtVjAdH1A=
(10 rows)
```


Thank You!

Any questions?

Bio: <https://tristazero.github.io>

LinkedIn: <https://www.linkedin.com/in/panjuan>

GitHub: <https://github.com/tristaZero>

Twitter: @tristaZero

Project Twitter: @ShardingSphere