



# Must-Know Postgres Extensions for DBAs and Developers During Migration

Deepak Mahto  
DataCloudGaze Consulting

# About Me

---

I am Deepak Mahto, and I like to call myself a Database Guy.

- Founder of DataCloudGaze Consulting.
- I have 15+ years of database experience, with more than 7 years in cloud and migrations.
- I have published 150+ technical blogs on databases.
- I live in Mumbai and have a 3-year-old child.
- Loves to explore street food.



# Agenda

---

- What are Postgres Extensions and Their Importance
- Some Fun Facts About Extensions
- List of Extensions to Aid in Migration from Heterogeneous Sources Like Oracle and MSSQL

# What is extension?

PostgreSQL's extensibility allows for seamless integration of extensions, making them function like built-in features, thus enhancing its capabilities and flexibility.



# Extensive versatility of Extension's

ExtensionCategory	ExtensionCount	Extension List
Analytics	14	hydra_columnar/pg_timeseries/timescaledb
AnalyticsConnectors	1	pg_tier
Auditing / Logging	7	auto_explain/pgaudit
Change Data Capture	6	pg_ivm/pglogical/wal2json
Connectors	27	dblink/oracle_fdw/mysql_fdw/postgres_fdw
Data / Transformations	52	citext/postgis/postgresql_anonymizer
Debugging	2	plpgsql_check
Index / Table Optimizations	15	hypopg/pg_repack/pg_squeeze/pgtt
Machine Learning	3	pg_embedding-pgvector-postgresml
Machine LearningOrchestration	1	vectorize
Metrics	17	pg_proctab/pg_stat_kcache/pg_stat_statements
MetricsTooling / Admin	1	pgtelemetry
Orchestration	8	pg_cron/pg_dbms_job/pg_partman
Procedural Languages	16	plpgsql/plproxy/plrust/plv8/pljava
Query Optimizations	6	pg_hint_plan/pg_stat_monitor
Search	11	fuzzystmatch/pg_trgm
Security	11	passwordcheck/pgcrypto
Tooling / Admin	13	pgtap/adminpack/citus

Reference - <https://cloud.tembo.io>

# Create Extension in Postgres

---

CREATE EXTENSION loads and manages new extensions, requiring appropriate privileges.

extensionrepo=# \help create extension

Command: CREATE EXTENSION

Description: install an extension

Syntax:

```
CREATE EXTENSION [ IF NOT EXISTS ] extension_name
    [ WITH ] [ SCHEMA schema_name ]
    [ VERSION version ]
    [ CASCADE ]
```

URL: <https://www.postgresql.org/docs/16/sql-createextension.html>

# Postgres fun fact - default extension.

---

PL/pgSQL is a loadable procedural language for the PostgreSQL database system that is default extension created with each new database.

```
postgres=# load 'plpgsql';
```

```
LOAD
```

```
postgres=# \dx plpgsql
```

List of installed extensions

Name	Version	Schema	Description
plpgsql	1.0	pg_catalog	PL/pgSQL procedural language

(1 row)

```
postgres=# do language plpgsql 'begin end';
```

```
DO
```

# More supported Procedural Language as extensions

---

- PL/R - *PostgreSQL support for R as a procedural language (PL)*
- PL/V8 - *A procedural language in JavaScript powered by V8*
- PL/Tcl - *Tcl procedural language for PostgreSQL.*
- PL/Perl - *The Perl procedural language for PostgreSQL.*
- PL/Rust - *Procedural language in the Rust programming.*
- PL/Python - *Untrusted procedural language for PostgreSQL.*

More..



# pg\_stat\_statements

---

## Track statistics of SQL planning and execution

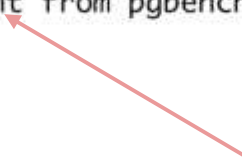
```
WITH statements AS (  
  SELECT * FROM pg_stat_statements pss  
  where pss.dbid in (SELECT oid from pg_database where datname=current_database())  
  and query like '%pgbench%'  
)  
SELECT calls,  
       mean_exec_time,  
       query  
FROM statements  
ORDER BY calls DESC  
LIMIT 3;
```

-[ RECORD 1 ]--+	
calls	225
mean_exec_time	<a href="#">0.34736720000000002</a>
query	SELECT abalance FROM pgbench_accounts WHERE aid = \$1
-[ RECORD 2 ]--+	
calls	1
mean_exec_time	0.0208
query	select count(*) from pgbench_branches

## pg\_stat\_statements - toplevel to track nested calls.

---

```
extension=# CREATE OR REPLACE FUNCTION public.sampleproceduralblock(id integer) RETURNS boolean
extension=#     LANGUAGE plpgsql
extension=#     AS $$
extension$# DECLARE
extension$# cnt bigint ;
extension$# BEGIN
extension$#     select count(1) into cnt from pgbench_accounts where bid = id;
extension$#     if cnt > 0 then
extension$#         return true;
extension$#     else
extension$#         return false;
extension$#     end if;
extension$# END $$;
CREATE FUNCTION
Time: 214.587 ms
extension=# explain analyze select public.sampleproceduralblock(col1) from generate_series(1,25) col1;
```



Problematic SQL!

## pg\_stat\_statements - toplevel to track nested calls.

---

*\*pg\_stat\_statements.track = 'all' (only for non prod or specific session)*

```
select query , total_exec_time , toplevel  from pg_stat_statements where  
query like 'select%sampleproceduralblock%' and toplevel;
```

```
-[ RECORD 1 ]---+-----
```

```
-----
```

```
query          | select public.sampleproceduralblock(col1) from generate_series($1,  
$2) col1
```

```
total_exec_time | 11203.602380999999
```

```
toplevel       | t
```

Time: 235.243 ms

```
select query , total_exec_time , toplevel  from pg_stat_statements where  
not toplevel;
```

```
-[ RECORD 1 ]---+-----
```

```
query          | select count($2)                                from pgbench_accounts where bid = id
```

```
total_exec_time | 11200.898884
```

```
toplevel       | f
```

Identify problematic sql within procedural block.



# pg\_hint\_plan - Influence sql performance

Makes it possible to tweak PostgreSQL execution plans using so-called "hints" in SQL comments

```
/*+SeqScan(t2)*/  
EXPLAIN (COSTS false) SELECT * FROM t1, t2 WHERE t1.id = t2.id;  
QUERY PLAN
```

-----  
Merge Join

Merge Cond: (t1.id = t2.id)

-> Index Scan using t1\_pkey on t1

-> Sort

Sort Key: t2.id

-> Seq Scan on t2

(6 rows)

Influence access  
path and  
Join method

# hypog - Invisible Indexes

---

PostgreSQL extension adding support for hypothetical indexes

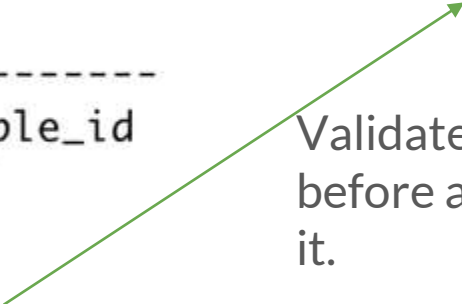
```
SELECT * FROM hypog_create_index('CREATE INDEX ON hypo_sample (id)') ;  
indexrelid | indexname
```

```
-----+-----  
13543 | <13543>btree_hypo_sample_id  
(1 row)
```

Time: 242.272 ms

```
EXPLAIN (COSTS OFF) SELECT val FROM hypo_sample WHERE id = 1;  
QUERY PLAN
```

```
-----  
Index Scan using "<13543>btree_hypo_sample_id" on hypo_sample  
Index Cond: (id = 1)  
(2 rows)
```



Validate and test indexes  
before actually creating  
it.

# orafce - Oracle's compatibility functions and packages

Emulate a subset of functions and packages from the Oracle RDBMS.

```
\df add_months
```

List of functions				
Schema	Name	Result data type	Argument data types	Type
oracle	add_months	date	day date, value integer	func
oracle	add_months	timestamp without time zone	timestamp with time zone, integer	func

(2 rows)

```
SELECT NVL(1,2) AS NVL , INSTR('POSTGRESQL ON MYDBOPS', 'SQL', 1) AS INSTR , DECODE(3, 1, 100, 2, 200,0) AS DECODE ;
nvl | instr | decode
```

```
-----
1 | 8 | 0
(1 row)
```

```
SELECT ADD_MONTHS(CLOCK_TIMESTAMP(),1) AS ADD_MONTHS ;
add_months
```

```
-----
2024-07-07 15:00:38
(1 row)
```

```
SELECT DBMS_RANDOM.RANDOM() AS RANDOM ;
random
```

```
-----
1873518366
(1 row)
```

```
SELECT COUNT(1) AS OUTPUT FROM DBA_SEGMENTS;
output
```

```
-----
288
```

# oracle\_fdw - Foreign Data Wrapper for Oracle

---

Facilitate seamless data migration and integration between Oracle databases and PostgreSQL.

```
postgres=> create extension oracle_fdw;
CREATE EXTENSION
postgres=> CREATE SERVER oradb1 FOREIGN DATA WRAPPER oracle_fdw
            OPTIONS (dbserver '//          :1521/ pdb1' );
CREATE SERVER
postgres=> CREATE USER MAPPING FOR postgres SERVER oradb1
            OPTIONS (user '          ', password '          ');
CREATE USER MAPPING
postgres=> IMPORT FOREIGN SCHEMA "SYS" LIMIT TO (DUAL)
            FROM SERVER oradb1 INTO public;
IMPORT FOREIGN SCHEMA
postgres=> select * from dual;
 dummy
-----
 X
(1 row)
```

# pgtt - Temporary Table across schema

---

Use Oracle-style Global Temporary Tables and the others RDBMS.

```
postgres=> \dx pgtt
                                List of installed extensions
  Name | Version | Schema | Description
-----+-----+-----+-----
 pgtt | 3.0.0   | pgtt_schema | Extension to add Global Temporary Tables feature to PostgreSQL
(1 row)

postgres=> load 'pgtt';
LOAD
postgres=> CREATE /*GLOBAL*/ TEMPORARY TABLE TMP1(COL1 INTEGER);
CREATE TABLE
postgres=> insert into TMP1 select 1;
INSERT 0 1
postgres=>

Session 1

postgres=> load 'pgtt';
LOAD
postgres=> table tmp1;
 col1
-----
(0 rows)

postgres=> insert into TMP1 select 2;
INSERT 0 1
postgres=> table tmp1;
 col1
-----
  2
(1 row)

Session 2
```



# pgaudit - Compliance and Regulatory

---

provides detailed session and/or object audit logging via the standard PostgreSQL logging facility.

```

                                List of installed extensions
  Name      | Version | Schema | Description
-----+-----+-----+-----
 pgaudit    | 16.0    | public | provides auditing functionality
(1 row)
```

```

extension=> \dconfig pgaudit.*
              List of configuration parameters
              Parameter              | Value
-----+-----
 pgaudit.log                        | write, ddl
 pgaudit.log_catalog                | on
 pgaudit.log_client                  | off
 pgaudit.log_level                   | log
 pgaudit.log_parameter               | off
 pgaudit.log_parameter_max_size     | 0
 pgaudit.log_relation                | off
 pgaudit.log_rows                    | off
 pgaudit.log_statement               | on
 pgaudit.log_statement_once          | off
 pgaudit.role                        |
```

# pgaudit

Configuring custom logging based on compliance requirement.

```
extension=> alter user test1 set pgaudit.log = 'ddl';
ALTER ROLE
extension=> alter user test2 set pgaudit.log = 'read,write';
ALTER ROLE
extension=> \drds test1|test2
```

List of settings

Role	Database	Settings
test1		pgaudit.log=ddl
test2		pgaudit.log=read,write

(2 rows)

```
DROP TABLE
CREATE TABLE
INSERT 0 1
```

ension:

ot logged>

ension:

id int,

>

ension:

o

user1',

# Table Partitioning Maintenance –

---

pg\_partm

partition sample=> █

I

# PL/pgSQL Conversion - Challenge

PL/pgSQL functions aren't syntax-checked until executed.

```
postgres=# CREATE OR REPLACE FUNCTION FUNC_DEMO1()  
postgres=# RETURNS void  
postgres=# LANGUAGE plpgsql  
postgres=# AS  
postgres=# $$  
postgres$$ DECLARE VAR1 INTEGER;  
postgres$$ BEGIN  
postgres$$ SELECT 1 INTO VAR1 WHERE COL1 = 1;  
postgres$$ END;  
postgres$$ $$;  
CREATE FUNCTION
```

postgres=# SELECT func\_demo1();  
**ERROR: column "col1" does not exist**  
LINE 1: SELECT 1 WHERE COL1 =  
1

QUERY: SELECT 1  
WHERE COL1 = 1  
CONTEXT: PL/pgSQL function  
func\_demo1() line 4 at SQL statement

# plpgsql\_check extension to rescue.

---

Extension serves as a comprehensive linter for plpgsql in Postgres

- Utilizes the internal PostgreSQL parser/evaluator to display runtime errors.
- Parses SQL inside routines to identify errors not typically found during "CREATE PROCEDURE/FUNCTION" commands.
- The plpgsql\_check extension detects issues in PL/pgSQL code: undefined/unused variables, type mismatches, control flow errors, incorrect function calls, trigger problems, and SQL statement errors.

## Some more essential extensions

---

- `pg_repack` - lets you remove bloat from tables and indexes as online.
- `pglogical` - PostgreSQL Logical Replication - Change Data Capture.
- `hydra_columnar` - Analytics- Columnar storage for Postgres
- `auto_explain` - logging execution plans of slow statements automatically.
- `postgresql_anonymizer` - Anonymization & Data Masking.
- `passwordcheck*` - Checks and rejects weak passwords.

# Thank you!

---

 <https://www.linkedin.com/in/mahtodeepak/>

 <https://x.com/mahtodeepak05>

 <https://databaserookies.wordpress.com/>



<https://www.datacloudgaze.com/>