



从查询计划入手优化数据库

郭峰

Pivotal 中国研发中心

Agenda

- Greenplum 查询处理过程
- Greenplum 查询计划解读
- Greenplum 分布式查询计划生成
- Greenplum 查询调优
- Q+A





Greenplum 查询处理过程

Greenplum 查询处理过程

Greenplum 对查询的处理过程可以分成四个步骤:

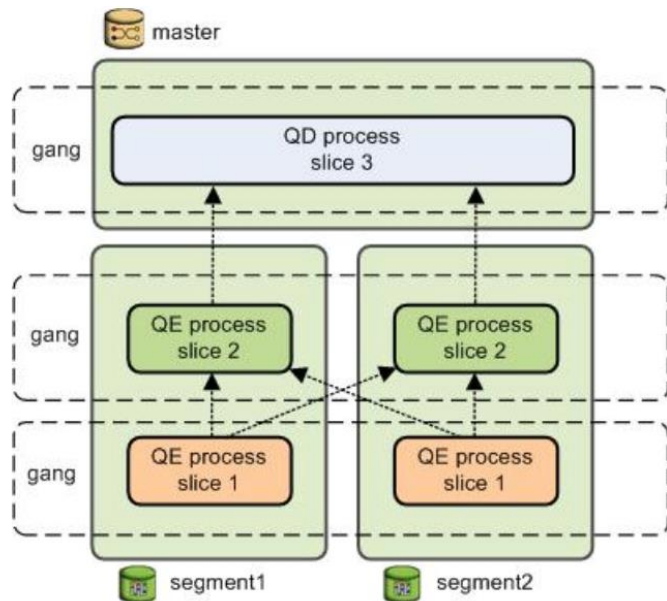
1. Master 接收查询语句并生成查询计划。
2. Master 把查询计划分发到Segments.



Greenplum 查询处理过程

Greenplum 对查询的处理过程可以分成四个步骤:

1. Master 接收查询语句并生成查询计划。
2. Master 把查询计划分发到Segments.
3. Segments 并发在各自本地的数据集上执行计划。
 - a. Slice: A portion of the plan that segments can work on independently.
 - a. Gang: Related processes that are working on the same slice of the query plan but on different segments.
1. Master 收集结果并返回给客户端。



Greenplum 查询计划解读

Greenplum 查询计划解读

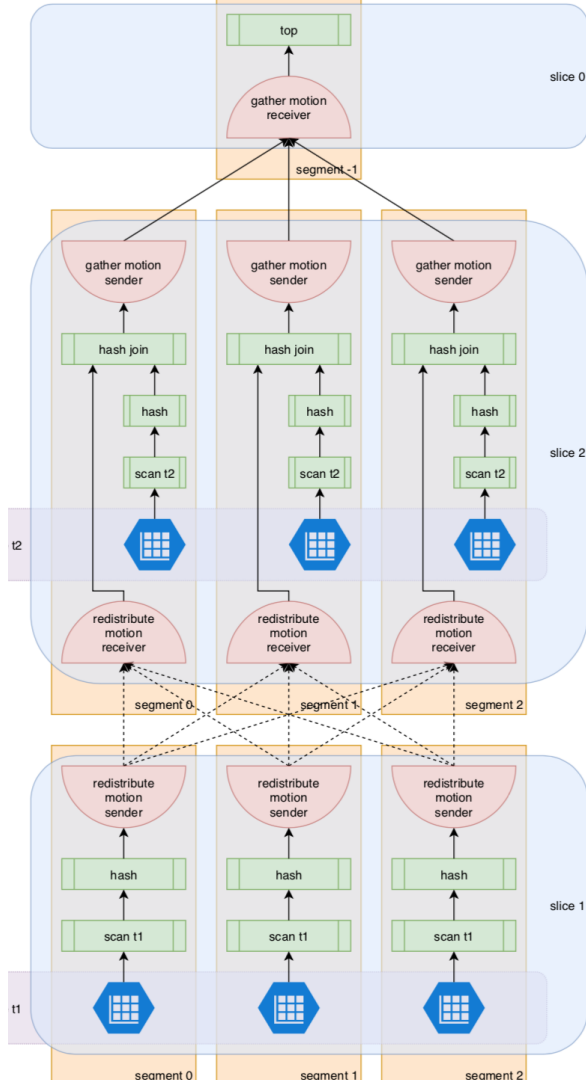
```
gpadmin=# EXPLAIN select * from t2 join t1 on t2.c1 = t1.c2;
```

QUERY PLAN

```
-----  
-----  
Gather Motion 3:1  (slice2; segments: 3)  (cost=2037.25..97600.62 rows=7413210  
width=16)  
  -> Hash Join  (cost=2037.25..97600.62 rows=2471070 width=16)  
      Hash Cond: t1.c2 = t2.c1  
        -> Redistribute Motion 3:3  (slice1; segments: 3)  (cost=0.00..2683.00  
rows=28700 width=8)  
            Hash Key: t1.c2  
              -> Seq Scan on t1  (cost=0.00..961.00 rows=28700 width=8)  
        -> Hash  (cost=961.00..961.00 rows=28700 width=8)  
            -> Seq Scan on t2  (cost=0.00..961.00 rows=28700 width=8)  
Optimizer status: legacy query optimizer  
(9 rows)
```

6个节点，其中2个Motion节点。

```
select * from t2 join  
t1 on t2.c1 = t1.c2
```



Motion: 在Segments之间传输数据

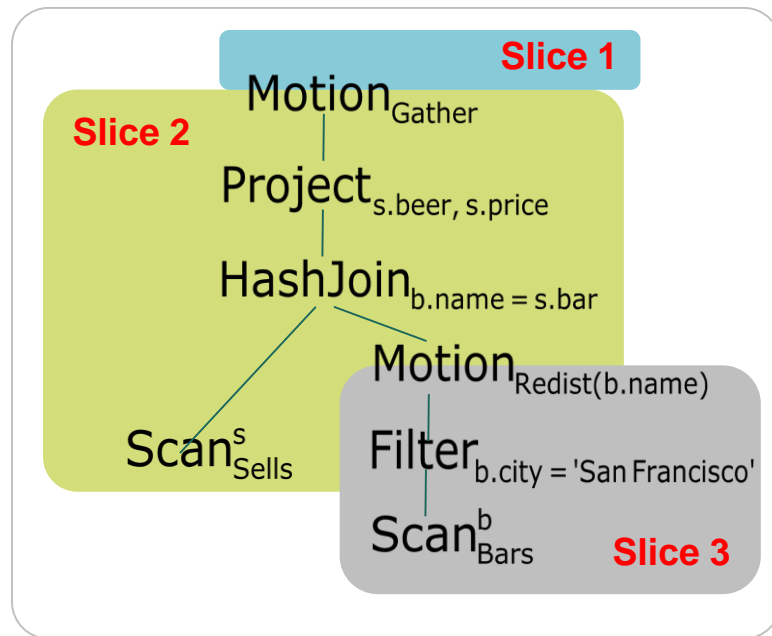
Slice: 根据Motion切割slice,
motion的两边各有一个slice.

Greenplum 查询计划解读

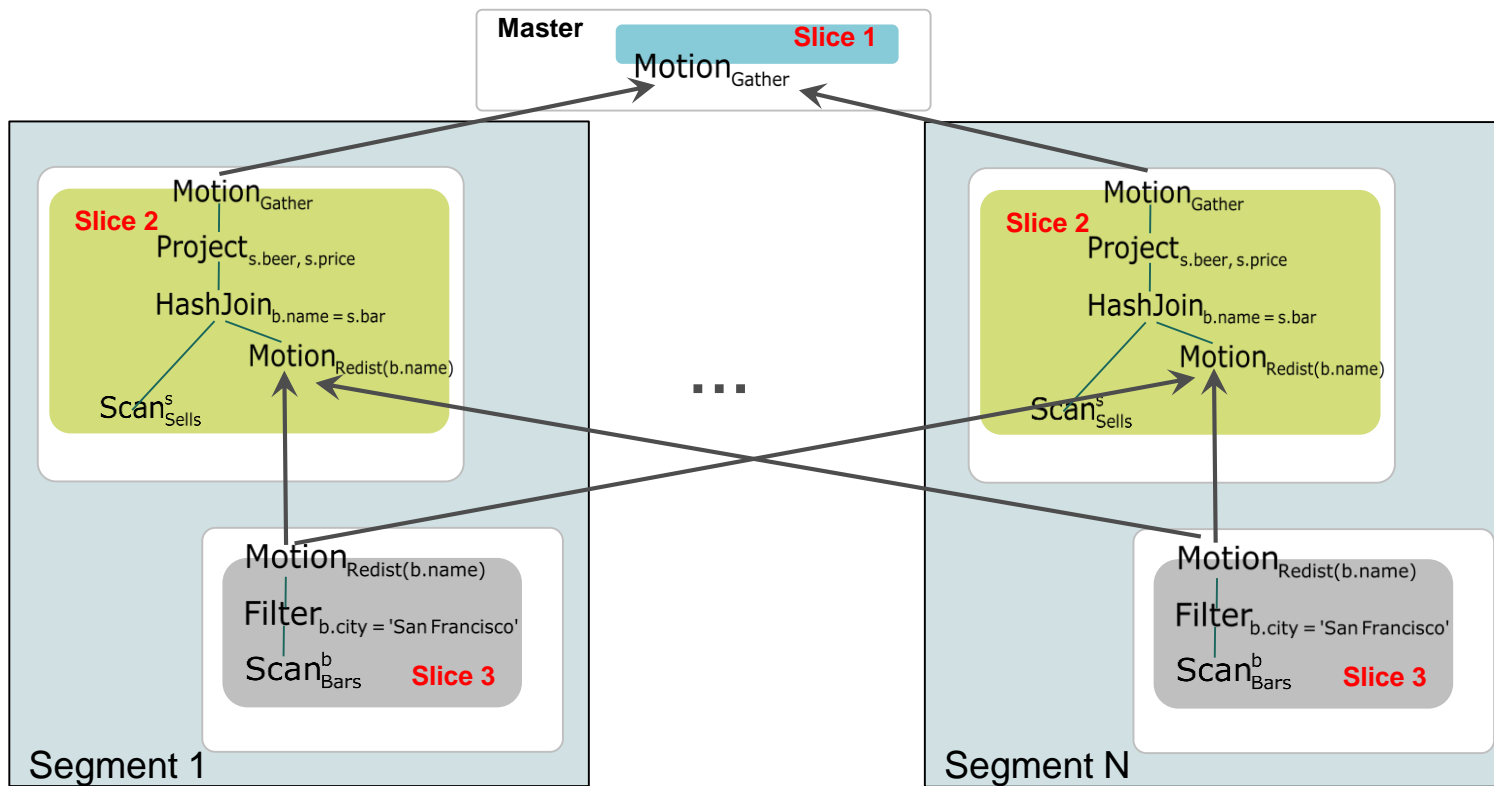
```
SELECT s.beer, s.price  
FROM Bars b, Sells s  
WHERE b.name = s.bar  
AND b.city = 'San Francisco'
```

Bars is distributed randomly

Sells is distributed by bar



Greenplum 查询计划解读



Greenplum 查询计划解读

每个节点含有三个估计值：**cost, rows, width.**

- Cost: Greenplum采用基于代价的优化器来评估执行查询的不同策略，并选择代价最低的方法。
 - 启动代价: 得到第一个tuple的代价。
 - 总代价：处理完所有tuple的代价。
- Rows: 查询节点输出的tuple的个数（估计值）。
- Width: 查询节点输出的tuple的长度（估计值），单位是字节。

Greenplum 查询计划解读

EXPLAIN ANALYZE

- 执行查询的总时间（以毫秒为单位）
- 内存使用量；
- 查询节点需要的工作进程个数
- 每个操作中处理最多行的Segment处理的最大行数以及Segment的序号
- 从处理最多行的Segment上获得第一行花费的时间（单位是毫秒）及从该Segment获得所有行花费的时间。

EXPLAIN ANALYZE 除了显示查询计划外还会执行查询，若需对DML语句使用EXPLAIN ANALYZE而不影响数据，可置EXPLAIN ANALYZE于事务中：

```
BEGIN;  
EXPLAIN ANALYZE ...;  
ROLLBACK;
```

Gather Motion 3:1 (slice2; segments: 3) (cost=2360.80..7000.90 rows=99992 width=16)

```
select * from t2 join t1 on  
t2.c1 = t1.c2
```

Rows out: 33333 rows at destination with 26 ms to first row, 93 ms to end, start offset by 0.530 ms.

-> Hash Join (cost=2360.80..7000.90 rows=33331 width=16)

Hash Cond: t1.c2 = t2.c1

Rows out: Avg 11111.0 rows x 3 workers. Max 11246 rows (seg1) with 24 ms to first row, 83 ms to end, start offset by 0.821 ms.

Executor memory: 1042K bytes avg, 1043K bytes max (seg0).

Work mem used: 1042K bytes avg, 1043K bytes max (seg0). Workfile: (0 spilling)

(seg1) Hash chain length 2.3 avg, 9 max, using 14237 of 16384 buckets.

-> **Redistribute Motion 3:3** (slice1; segments: 3) (cost=0.00..3137.97 rows=33633 width=8)

Hash Key: t1.c2

Rows out: Avg 33333.3 rows x 3 workers at destination. Max 33468 rows (seg0) with 0.433 ms to first row, 40 ms to end, start offset by 25 ms.

-> Seq Scan on t1 (cost=0.00..1119.99 rows=33633 width=8)

Rows out: Avg 33333.3 rows x 3 workers. Max 33348 rows (seg0) with 0.048 ms to first row, 4.813 ms to end, start offset by 0.991 ms.

-> Hash (cost=1110.91..1110.91 rows=33331 width=8)

Rows in: Avg 33333.3 rows x 3 workers. Max 33348 rows (seg0) with 24 ms to end, start offset by 0.904 ms.

-> Seq Scan on t2 (cost=0.00..1110.91 rows=33331 width=8)

Rows out: Avg 33333.3 rows x 3 workers. Max 33348 rows (seg0) with 0.045 ms to first row, 4.892 ms to end, start offset by 0.905 ms.

Slice statistics:

(slice0) Executor memory: 386K bytes.

(slice1) Executor memory: 212K bytes avg x 3 workers, 212K bytes max (seg0).

(slice2) Executor memory: 4767K bytes avg x 3 workers, 4767K bytes max (seg0). Work_mem: 1043K bytes max.

Statement statistics:

Memory used: 4096K bytes

Optimizer status: legacy query optimizer

Total time: 96.875 ms

(25 rows)

Greenplum 查询计划解读

with T1 ms to first row, T2 ms to end, start offset by T3 ms

- T1: 这个查询节点上得到第一个tuple的执行时间
- T2: 这个查询节点上得到所有tuple的总执行时间
- T3: 这个查询节点执行开始时间 减去 这个查询的开始时间
 - Interconnect的建立
 - 分发查询计划
 - 创建需要的gang

Greenplum 查询计划解读

Work_mem used

- 这个查询节点执行时用到的内存大小。
- 如果分配给这个查询节点的内存量不足够用来执行操作，部分数据将溢出的文件中。

Work_mem used: 68K bytes avg, 69K bytes max (seg0). Workfile: (3 spilling)

Work_mem wanted: 1042K bytes avg, 1043K bytes max (seg0) to lessen workfile I/O affecting 3 workers.

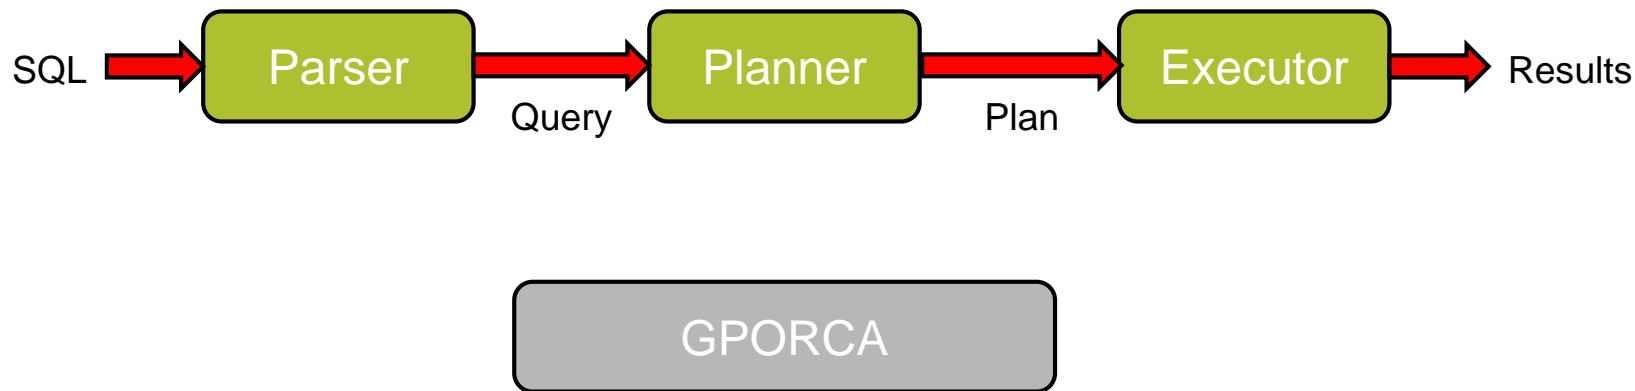
```
(seg0) Initial batch 0:
(seg0)   Wrote 480K bytes to inner workfile.
(seg0)   Wrote 480K bytes to outer workfile.
(seg0) Initial batches 1..15:
(seg0)   Read 611K bytes from inner workfile: 41K avg x 15 nonempty batches, 43K max.
(seg0)   Read 613K bytes from outer workfile: 41K avg x 15 nonempty batches, 43K max.
(seg0) Hash chain length 2.3 avg, 10 max, using 14298 of 16384 buckets.
```

Greenplum 分布式查询计划生成

Greenplum 分布式查询计划生成

Greenplum支持两种查询优化器：Planner 和 GPORCA

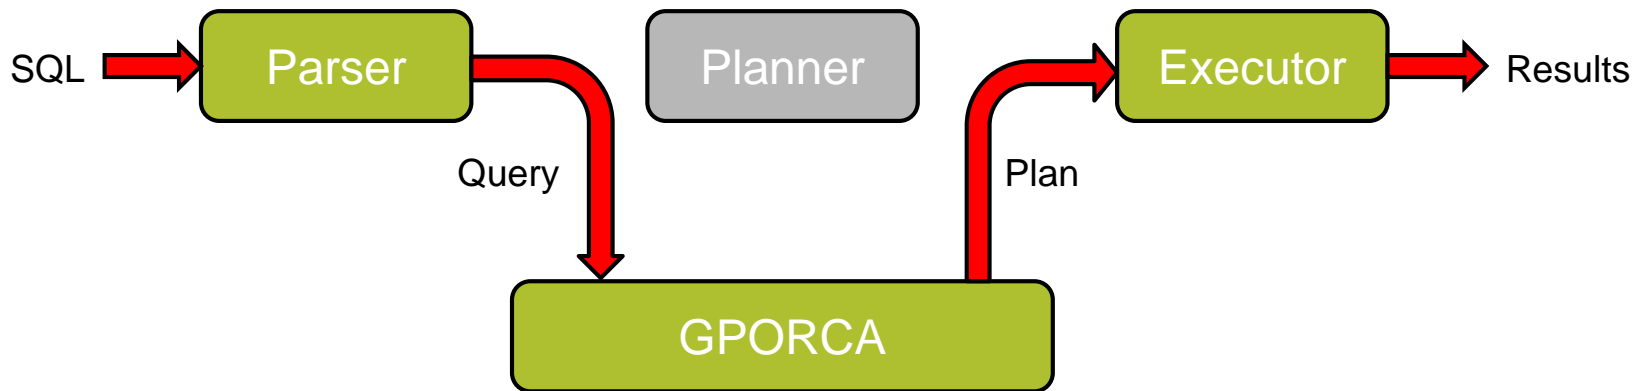
使用Planner:



Greenplum 分布式查询计划生成

Greenplum支持两种查询优化器：Planner 和 GPORCA

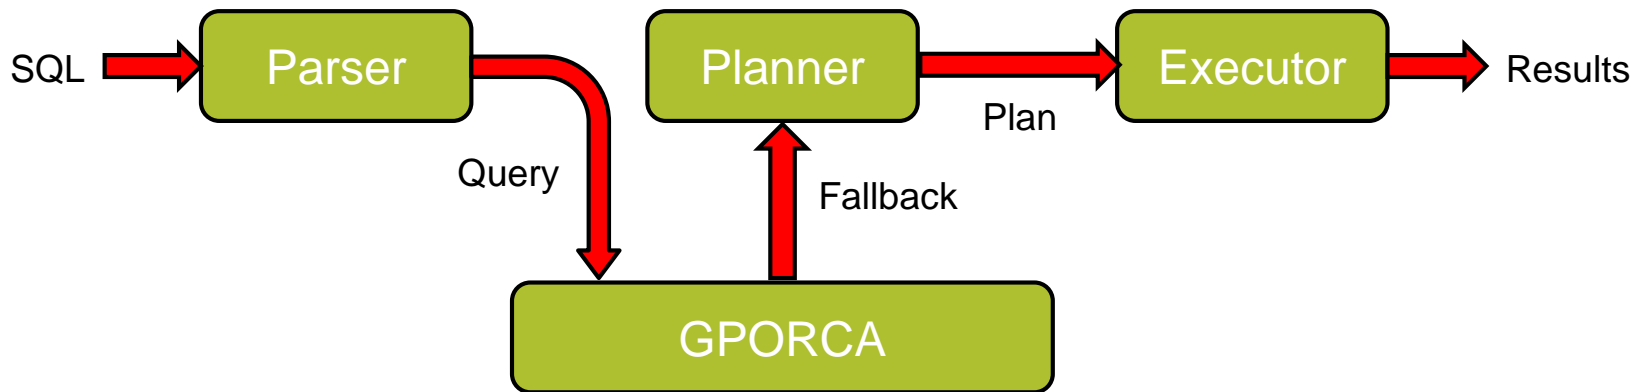
使用GPORCA:



Greenplum 分布式查询计划生成

Greenplum支持两种查询优化器：Planner 和 GPORCA

对于GPORCA不支持的特性，GPORCA会自动fall back到Planner:



GPORCA or Planner?

```
bootcamp=# explain select * from products;  
               QUERY PLAN
```

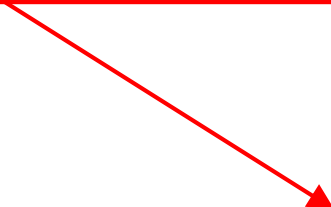
```
-----  
Gather Motion 2:1  (slicel; segments: 2)  (cost=0.00..431.00 rows=1 width=11)
```

```
-> Table Scan on products  (cost=0.00..431.00 rows=1 width=11)
```

```
Settings:  optimizer=on
```

```
Optimizer status: PQO version 1.628
```

```
(4 rows)
```



PQO: Pivotal Query Optimizer
+
Optimizer Version #

GPORCA or Planner?

```
bootcamp=# explain select * from products;  
              QUERY PLAN
```

```
-----  
Gather Motion 2:1  (slice1; segments: 2)  (cost=0.00..1.01 rows=1 width=11)
```

```
  -> Seq Scan on products  (cost=0.00..1.01 rows=1 width=11)
```

```
Optimizer status: legacy query optimizer  
(3 rows)
```




Planner Plan

Greenplum 分布式查询计划生成

```
create table t1(c1 int, c2 int, c3 int) distributed by (c1);  
create table t2(c1 int, c2 int, c3 int) distributed by (c1);
```

Greenplum 分布式查询计划生成: Scan

对于Scan节点，每个segment扫描其本地数据，最后master通过Gather Motion来收集结果。




```
SELECT * from t1 where t1.c2 = 1;
```

```
QUERY PLAN
```

```
-----  
Gather Motion 3:1  
  -> Seq Scan on t1  
      Filter: (c2 = 1)  
(3 rows)
```

Greenplum 分布式查询计划生成: Join

对于Join节点，如果我们执行的是基于分布键的等值连接，那么每个segment可以执行本地连接，最后master通过Gather Motion来收集结果。



```
SELECT * from t1, t2 where t1.c1 = t2.c1;  
QUERY PLAN
```

```
-----  
Gather Motion 3:1
```

```
-> Hash Join
```

```
Hash Cond: (t1.c1 = t2.c1)
```

```
-> Seq Scan on t1
```

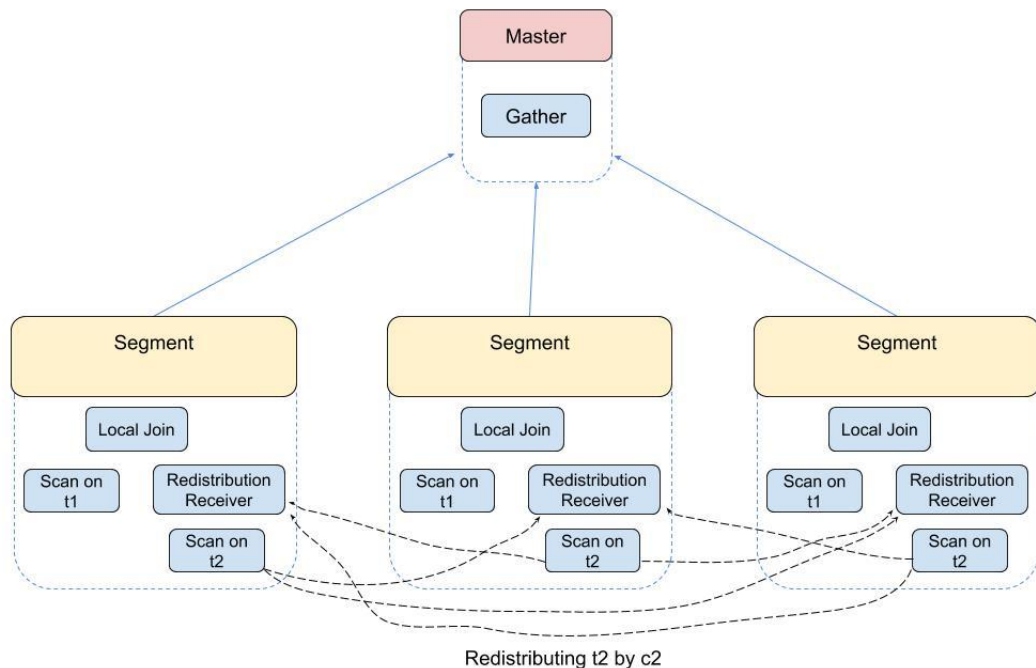
```
-> Hash
```

```
-> Seq Scan on t2
```

```
(6 rows)
```


Greenplum 分布式查询计划生成: Join

否则，我们可能会需要重分布其中一个表。



```
SELECT * from t1, t2 where t1.c1 = t2.c2;  
QUERY PLAN
```

Gather Motion 3:1

-> Hash Join

Hash Cond: (t1.c1 = t2.c2)

-> Seq Scan on t1

-> Hash

-> Redistribute Motion 3:3

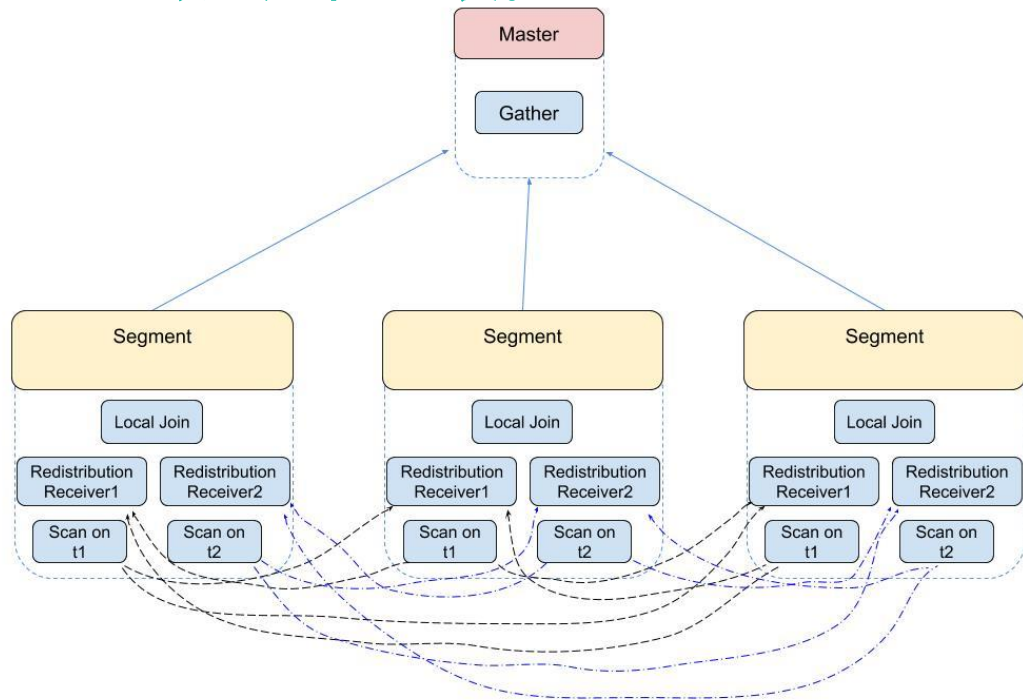
Hash Key: t2.c2

-> Seq Scan on t2

(8 rows)

Greenplum 分布式查询计划生成: Join

或重分布两个表。



Redistributing t1 and t2 by c2

```
SELECT * from t1, t2 where t1.c2 = t2.c2;  
QUERY PLAN
```

Gather Motion 3:1

-> Hash Join

Hash Cond: (t1.c2 = t2.c2)

-> Redistribute Motion 3:3

Hash Key: t1.c2

-> Seq Scan on t1

-> Hash

-> Redistribute Motion 3:3

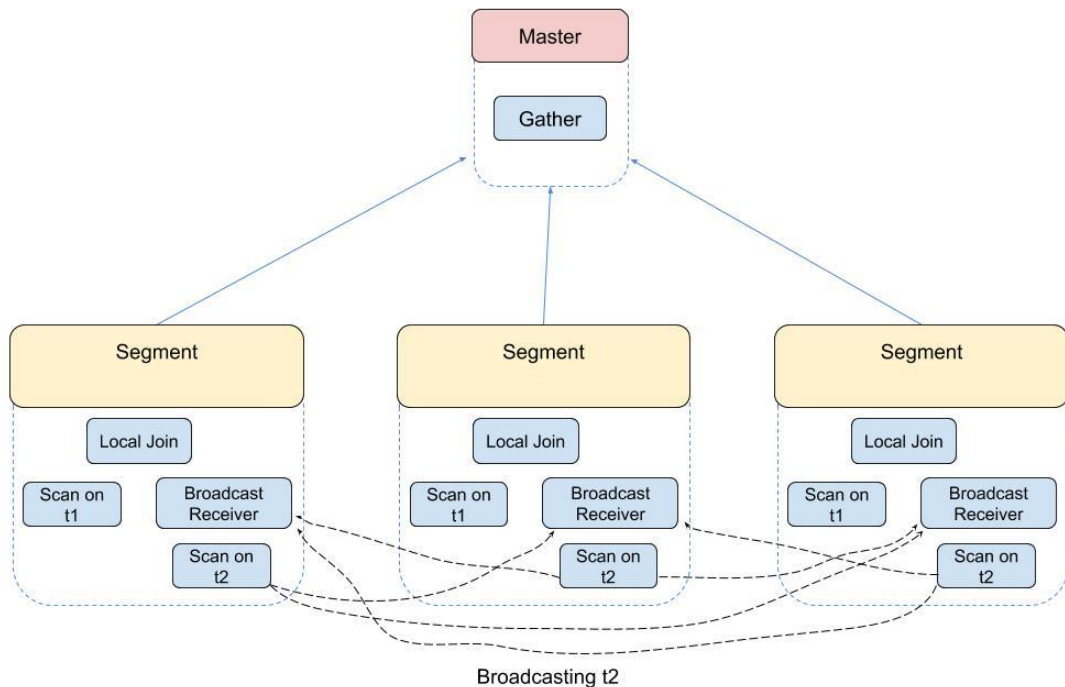
Hash Key: t2.c2

-> Seq Scan on t2

(10 rows)

Greenplum 分布式查询计划生成: Join

或者广播其中一个表。



```
SELECT * from t1, t2 where t1.c2 = t2.c2;  
QUERY PLAN
```

Gather Motion 3:1

-> Hash Join

Hash Cond: (t1.c2 = t2.c2)

-> Seq Scan on t1

-> Hash

-> Broadcast Motion 3:3

-> Seq Scan on t2

(7 rows)

Greenplum 分布式查询计划生成: Aggregate without Group


对于没有分组的聚集，**Greenplum**是通过两阶段聚集来完成的。

第一阶段：每个**segment**执行一次本地聚集。

第二阶段：**Master**通过**Gather Motion**收集第一阶段聚集的结果，然后执行第二阶段聚集。

Greenplum 分布式查询计划生成: Aggregate without Group

如果没有DISTINCT, 或者DISTINCT键是分布键, 那么可以直接执行两阶段聚集




```
SELECT avg(c1) from t1;  
QUERY PLAN
```

```
-----  
Aggregate  
-> Gather Motion 3:1  
    -> Aggregate  
        -> Seq Scan on t1  
(4 rows)
```

Greenplum 分布式查询计划生成: Aggregate without Group

否则，需要先重分布数据，然后再执行两阶段聚集。



```
SELECT avg(distinct c2) from t1;  
QUERY PLAN
```

```
-----  
Aggregate  
-> Gather Motion 3:1  
    -> Aggregate  
        -> Redistribute Motion 3:3  
            Hash Key: t1.c2  
                -> Seq Scan on t1
```

(6 rows)


Greenplum 分布式查询计划生成: Aggregate with Group

对于有分组的聚集，Greenplum的两个基本原则是：

1. 把属于同一组的数据重分布到同一segment上做聚集操作。
2. 尽量把不同的组分到不同的segment上从而提高并发度。

Greenplum 分布式查询计划生成: Aggregate with Group

如果是根据分布键分组，那么每个segment执行一阶段聚集。

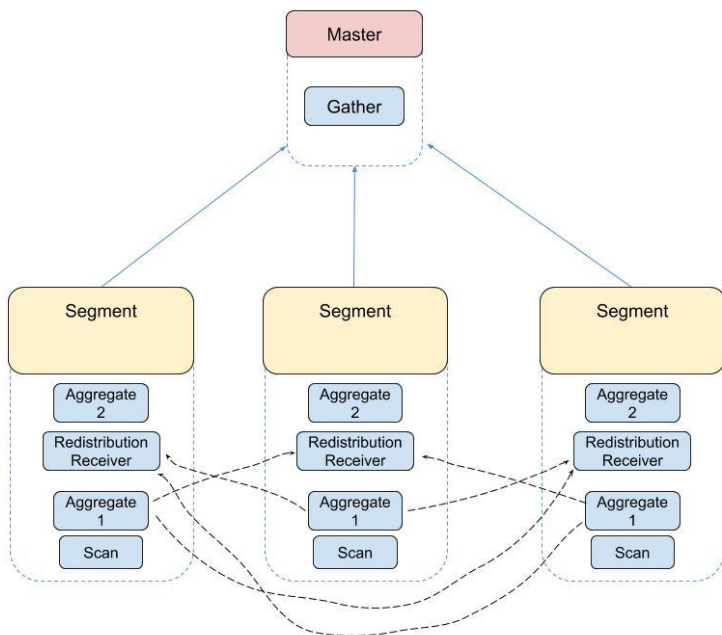


```
SELECT avg(c2) from t1 group by c1;  
QUERY PLAN
```

```
-----  
Gather Motion 3:1  
  -> HashAggregate  
        Group Key: c1  
        -> Seq Scan on t1  
(4 rows)
```


Greenplum 分布式查询计划生成: Aggregate with Group

如果是根据非分布键分组，同时没有DISTINCT操作，那么每个segment执行两阶段聚集。



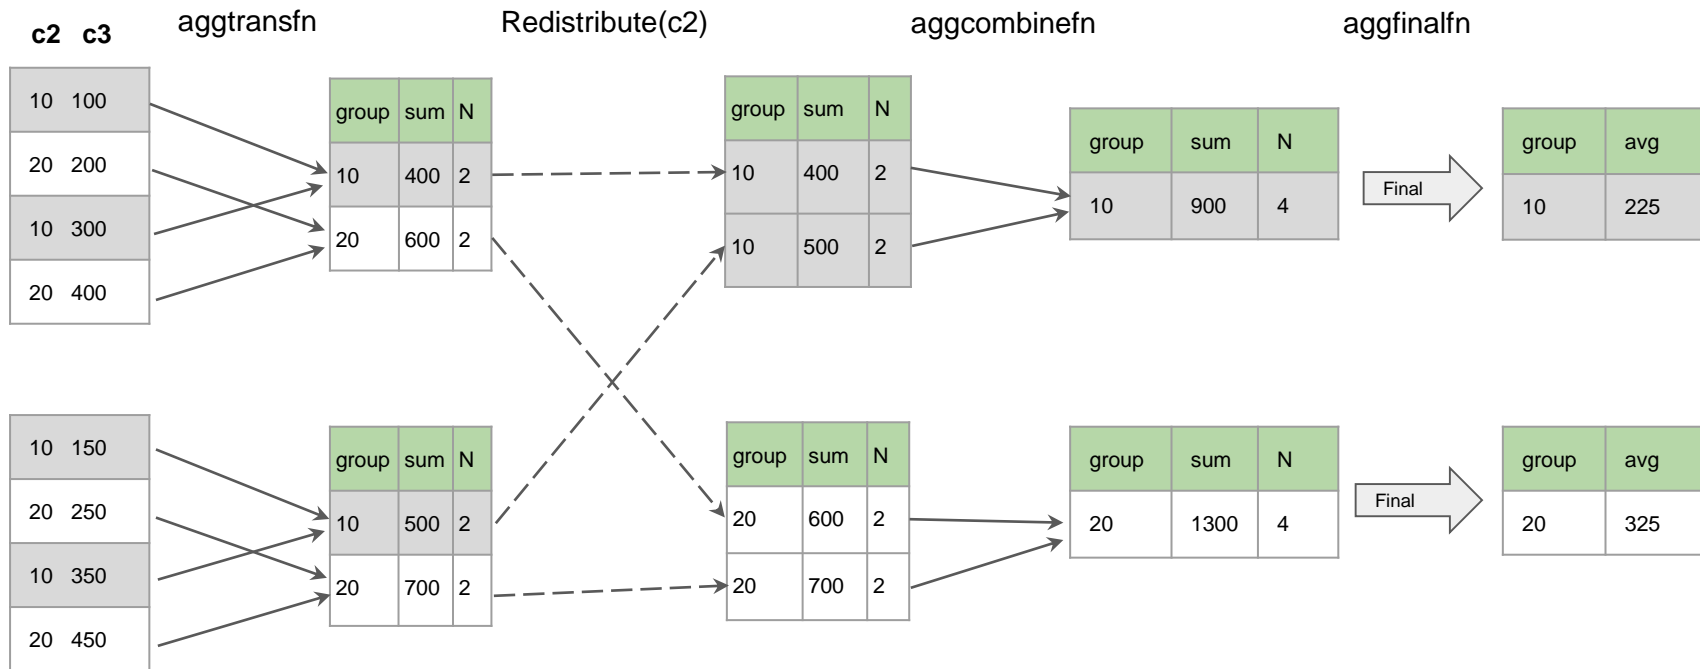
Redistributing (group key+transvalue) by group key

```
SELECT avg(c3) from t1 group by c2;  
QUERY PLAN
```

```
-----  
Gather Motion 3:1  
-> HashAggregate  
    Group Key: t1.c2  
-> Redistribute Motion 3:3  
    Hash Key: t1.c2  
-> HashAggregate  
    Group Key: t1.c2  
    -> Seq Scan on t1  
  
(8 rows)
```

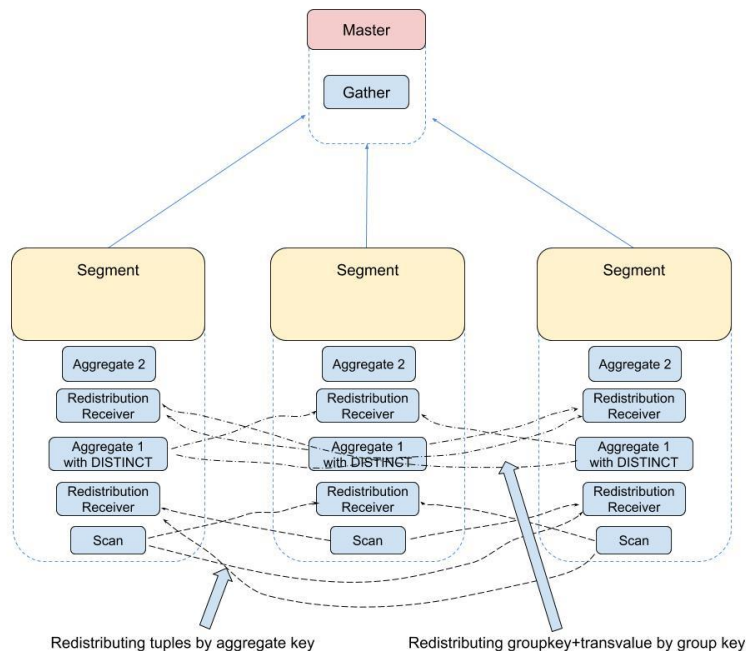
Greenplum 分布式查询计划生成: Aggregate with Group

```
SELECT avg(c3) from t1 group by c2;
```



Greenplum 分布式查询计划生成: Aggregate with Group

如果是根据非分布键分组，且有DISTINCT操作，那么每个segment先要重分布数据，然后再执行两阶段聚集。



```
SELECT avg(distinct c3) from t1 group by c2;  
QUERY PLAN
```

Gather Motion 3:1

-> GroupAggregate

Group Key: t1.c2

-> Sort

Sort Key: t1.c2

-> Redistribute Motion 3:3

Hash Key: t1.c2

-> GroupAggregate

Group Key: t1.c2

-> Sort

Sort Key: t1.c2

-> Redistribute Motion 3:3

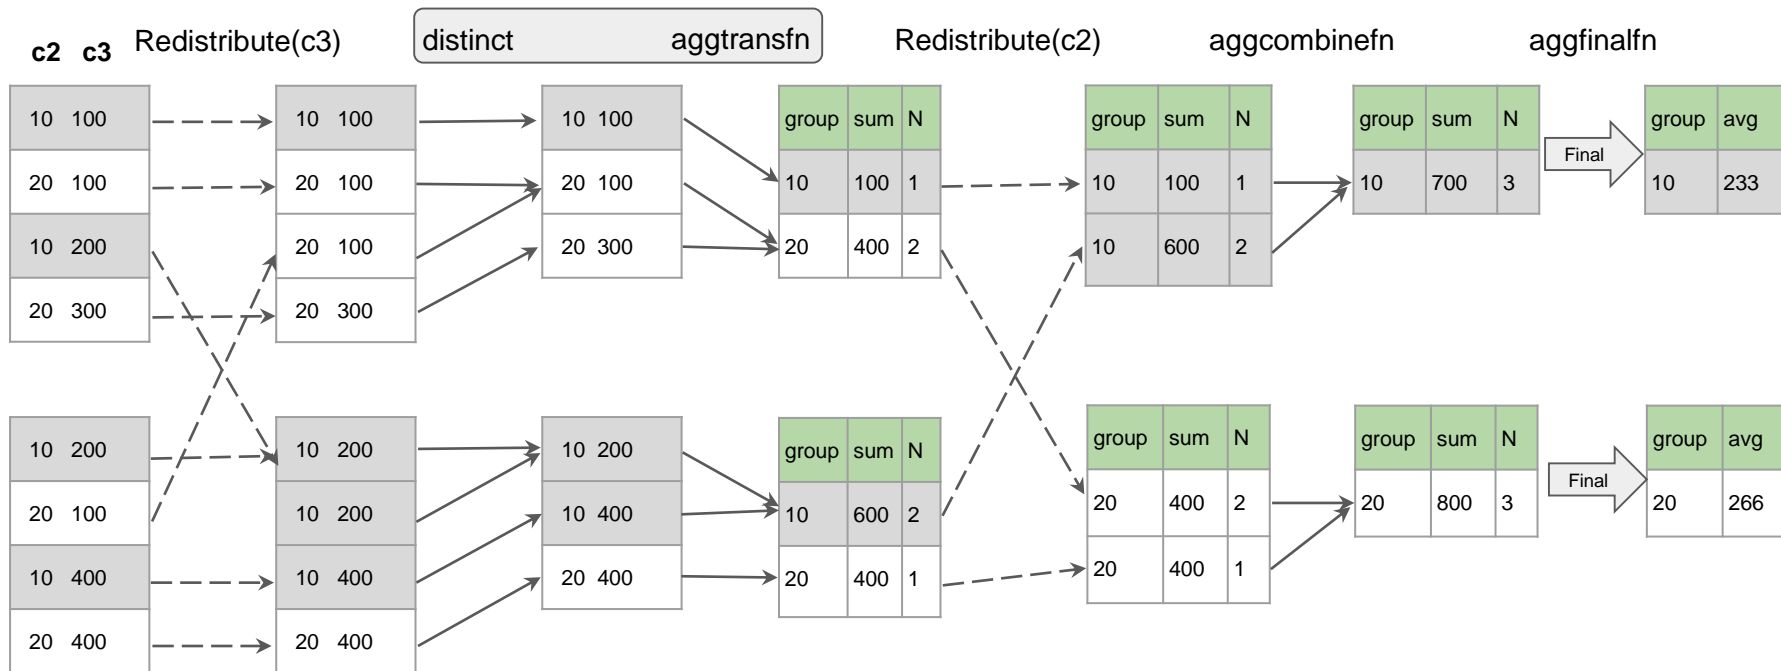
Hash Key: t1.c3

-> Seq Scan on t1

(14 rows)


Greenplum 分布式查询计划生成: Aggregate with Group

```
SELECT avg(distinct c3) from t1 group by c2;
```



Greenplum 分布式查询计划生成: Aggregate with Group

或者是每个segment执行三阶段聚集。



Redistributing (group key + aggregate key) by group key

```
SELECT avg(distinct c3) from t1 group by c2;
QUERY PLAN
```

Gather Motion 3:1

-> HashAggregate

Group Key: t1.c2

-> HashAggregate

Group Key: t1.c2, t1.c3

-> Redistribute Motion 3:3

Hash Key: t1.c2

-> HashAggregate

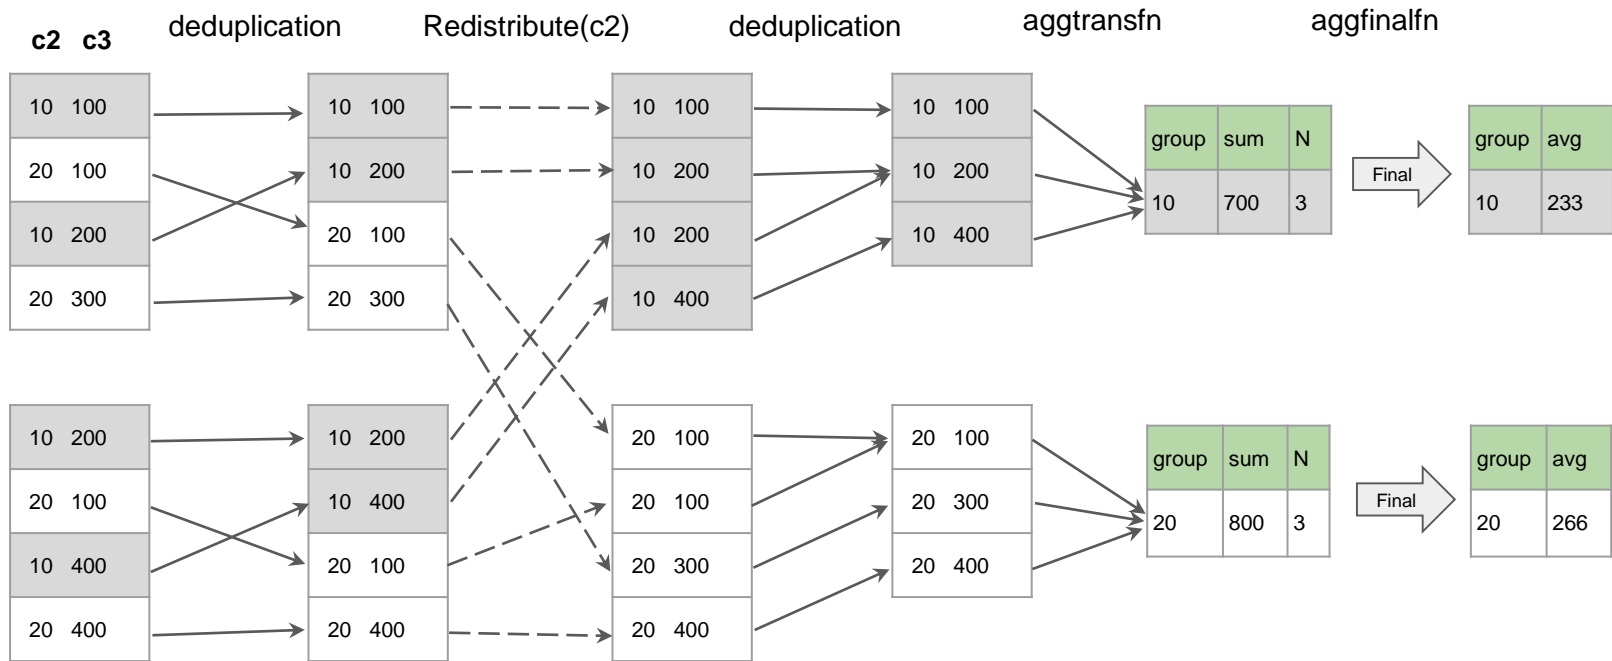
Group Key: t1.c2, t1.c3

-> Seq Scan on t1

(10 rows)

Greenplum 分布式查询计划生成: Aggregate with Group

```
SELECT avg(distinct c3) from t1 group by c2;
```



Greenplum 查询调优

Greenplum 查询调优: Example 1/2

```
# explain analyze select * from t1 join t2 on t1.c1 = t2.c2;
```

QUERY PLAN

```
-----
Gather Motion 3:1  (slice2; segments: 3)  (cost=1192.55..5169.33 rows=50380 width=16)
  Rows out:  51500 rows at destination with 88 ms to first row, 187 ms to end.
-> Hash Join  (cost=1192.55..5169.33 rows=16794 width=16)
   Hash Cond: t2.c2 = t1.c1
   Rows out:  Avg 17166.7 rows x 3 workers.  Max 50000 rows (seg0) with 132 ms to first row, 166 ms to end.
   Executor memory:  524K bytes avg, 1563K bytes max (seg0).
   Work_mem used:  524K bytes avg, 1563K bytes max (seg0). Workfile: (3 spilling)
   Work_mem wanted: 1563K bytes avg, 1563K bytes max (seg0) to lessen workfile I/O affecting 1 workers.
   (seg0)  Initial batch 0:
   (seg0)  Initial batches 1..7:
   (seg0)  Secondary Overflow batches 8..15:
   (seg0)  Hash chain length 50000.0 avg, 50000 max, using 1 of 1024 buckets.  Skipped 15 empty batches.
-> Redistribute Motion 3:3  (slice1; segments: 3)  (cost=0.00..3098.10 rows=33190 width=8)
   Hash Key: t2.c2
   Rows out:  Avg 33333.3 rows x 3 workers at destination.  Max 33348 rows (seg0) with 0.429 ms to first row, 38 ms to
end.

   -> Seq Scan on t2  (cost=0.00..1106.70 rows=33190 width=8)
       Rows out:  Avg 33333.3 rows x 3 workers.  Max 33348 rows (seg0) with 0.049 ms to first row, 4.865 ms to end.
-> Hash  (cost=562.80..562.80 rows=16794 width=8)
   Rows in:  Avg 99.5 rows x 2 workers.  Max 125 rows (seg1) with 1.178 ms to end, start offset by 0.929 ms.
   -> Seq Scan on t1  (cost=0.00..562.80 rows=16794 width=8)
       Rows out:  Avg 17166.7 rows x 3 workers.  Max 50000 rows (seg0) with 0.048 ms to first row, 7.870 ms to end.

Optimizer status: legacy query optimizer
(38 rows)
```


Greenplum 查询调优: Example 1/2

1. 数据倾斜

ALTER TABLE SET DISTRIBUTED BY

1. 数据溢出

SET statement_mem TO XXX

Greenplum 查询调优: Example 2/2

```
# EXPLAIN ANALYZE select avg(distinct c2) from t1 group by c2;
```

QUERY PLAN

```
-----
--
Gather Motion 3:1  (slice2; segments: 3)  (cost=144826.86..164922.45 rows=946623 width=36)
  Rows out:  1000000 rows at destination with 829 ms to first row, 6591 ms to end.
  -> GroupAggregate  (cost=144826.86..164922.45 rows=315541 width=36)
    Group By: c2
    Rows out:  Avg 333333.3 rows x 3 workers.  Max 333385 rows (seg2) with 828 ms to first row, 6623 ms to end.
    Executor memory:  10692717K bytes avg, 10694374K bytes max (seg2).
    Work_mem used:  33K bytes avg, 33K bytes max (seg0).
    -> Sort  (cost=144826.86..147581.13 rows=367236 width=4)
      Sort Key: c2
      Sort Method:  quicksort  Max Memory: 36857KB  Avg Memory: 36857KB (3 segments)
      Rows out:  Avg 366666.7 rows x 3 workers.  Max 366704 rows (seg0) with 823 ms to first row, 892 ms to end.
      Executor memory:  36857K bytes avg, 36857K bytes max (seg0).
      Work_mem used:  36857K bytes avg, 36857K bytes max (seg0). Workfile: (0 spilling)
      -> Redistribute Motion 3:3  (slice1; segments: 3)  (cost=0.00..34263.24 rows=367236 width=4)
        Hash Key: c2
        Rows out:  Avg 366666.7 rows x 3 workers at destination.  Max 366704 rows (seg0) with 1.061 ms to first row,
422 ms to end
.
      -> Seq Scan on t1  (cost=0.00..12229.08 rows=367236 width=4)
        Rows out:  Avg 366666.7 rows x 3 workers.  Max 366704 rows (seg0) with 0.052 ms to first row, 122 ms to
end.
Optimizer status: legacy query optimizer
(27 rows)
```

Greenplum 查询调优: Example 2/2

多阶段聚集

SET gp_enable_multiphase_agg TO on;

SET gp_enable_agg_distinct TO on;

2-phase

DQA

SET gp_enable_agg_distinct_pruning TO on;

3-phase DQA



Pivotal®

Transforming How The World Builds Software