

Clustering in PostgreSQL


Because one database server is never enough (and
neither is two)



PGConf Nepal
May 11-12, 2023

```
postgres=# select * from umair;
```

```
-[ RECORD 1 ]-----  
name          | Umair Shahid  
description    | 20 year veteran of the PostgreSQL community  
company        | Stormatics  
designation    | Founder  
location       | Islamabad, Pakistan  
family         | Mom, Wife & 2 kids  
kid1           | Son, 16 year old  
kid2           | Daughter, 13 year old
```






STORMATICS

PostgreSQL Solutions for the Enterprise

Clustering & Analytics, backed by 24/7 Support and Professional Services






10 years of popularity

As measured by



https://db-engines.com/en/ranking_trend



PostgreSQL




Oracle




MySQL




SQL Server



Most
Loved




Loved by Developers



Most Used
(Professional Developers)



stackoverflow
2022 Developer Survey



Most
Wanted

On to the topic now!



What is High Availability?

- Remain operational even in the face of hardware or software failure
- Minimize downtime
- Essential for mission-critical applications that require 24/7 availability
- Measured in 'Nines of Availability'



Nines of Availability

Availability	Downtime per year
90% (one nine)	36.53 days
99% (two nines)	3.65 days
99.9% (three nines)	8.77 hours
99.99% (four nines)	52.60 minutes
99.999% (five nines)	5.26 minutes

But my database resides
in the cloud, and the
cloud is always available

Right?



Wrong!



Amazon RDS Service Level Agreement

Multi-AZ configurations for MySQL, MariaDB, Oracle, and PostgreSQL are covered by the Amazon RDS Service Level Agreement ("SLA"). The RDS SLA affirms that AWS will use **commercially reasonable efforts** to make Multi-AZ instances of Amazon RDS available with a Monthly Uptime Percentage of at least **99.95%** during any monthly billing cycle. In the event Amazon RDS does not meet the Monthly Uptime Percentage commitment, affected customers will be eligible to receive a service credit.*

99.95% = three and a half nines = 4.38 hours of downtime per year!!!

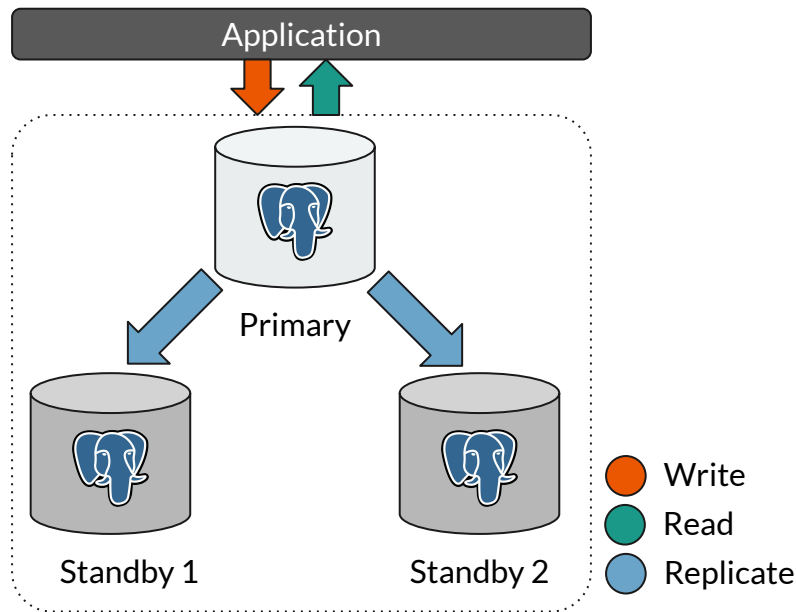
* <https://aws.amazon.com/rds/ha/>

So - what do I do if I want better reliability for my mission-critical data?

Clustering!

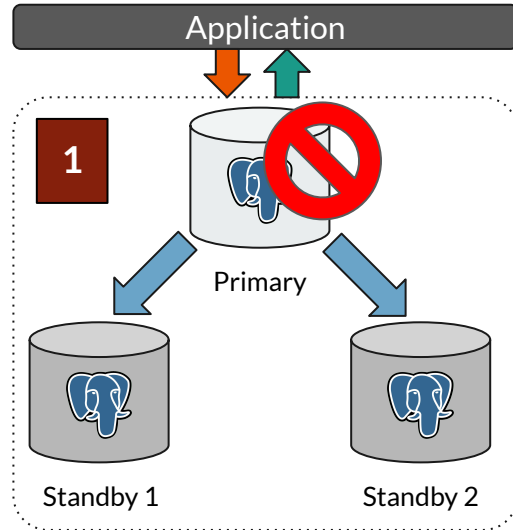


What is clustering?

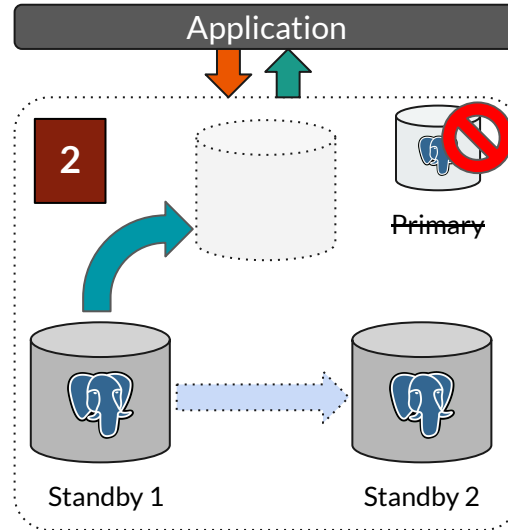


- Multiple database servers work together to provide redundancy
- Gives the appearance of a single database server
- Application communicates with the primary PostgreSQL instance
- Data is replicated to standby instances
- Auto failover in case the primary node goes down


What is auto failover?



- Primary node goes down

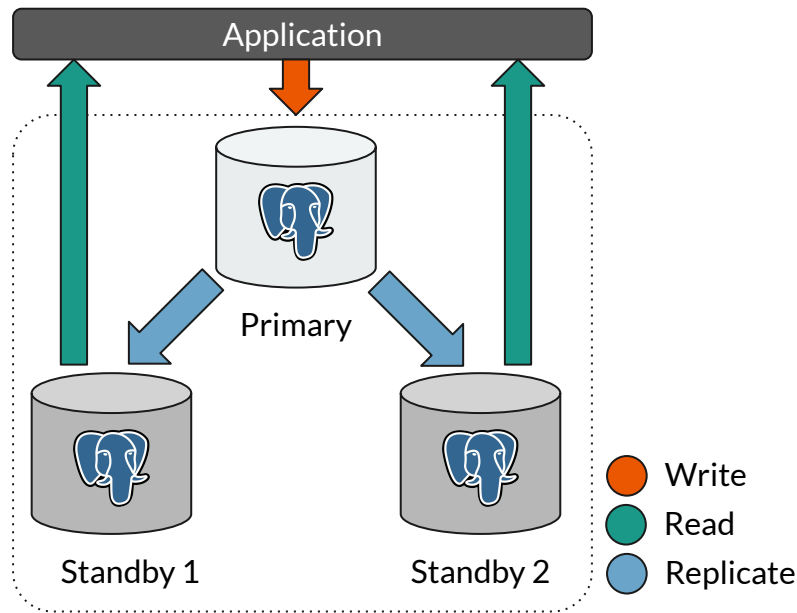


- Standby 1 gets promoted to Primary
- Standby 2 becomes subscriber to Standby 1



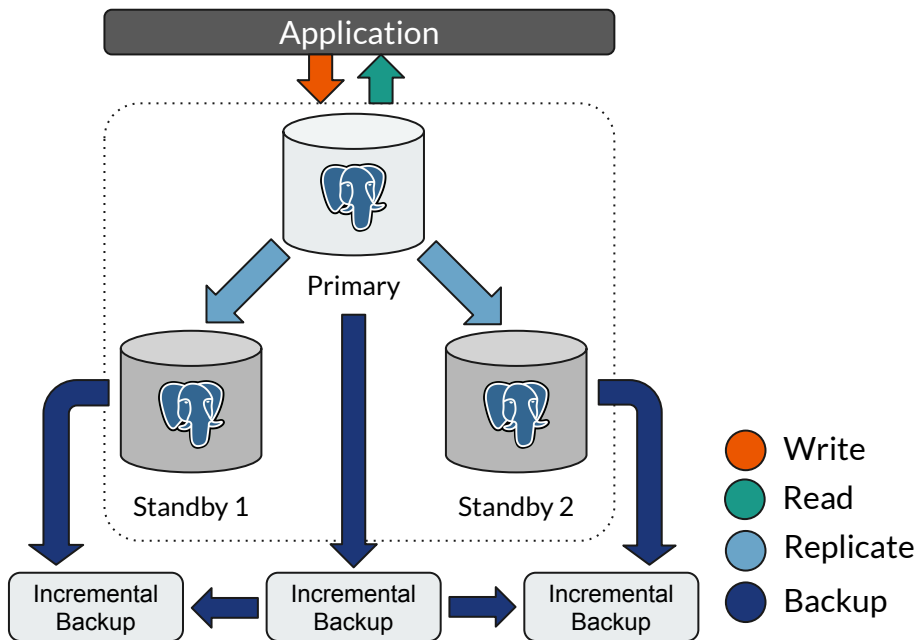
- New Standby is added to the cluster
- Application talks to the new Primary

Clusters with load balancing



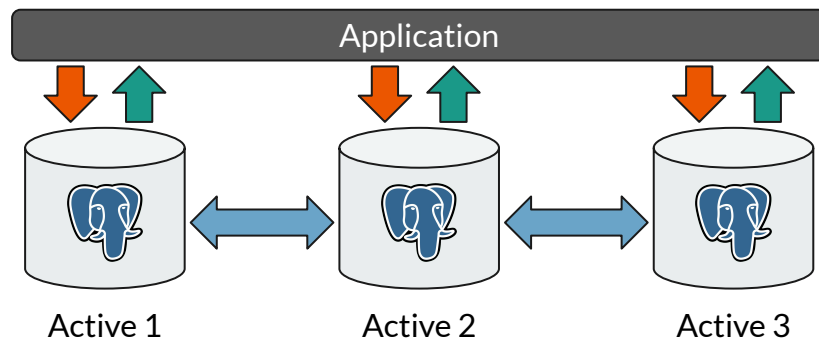
- Write to the primary PostgreSQL instance and read from standbys
- Data redundancy through replication to two standbys
- Auto failover in case the primary node goes down

Clusters with backups and disaster recovery



- Incremental backups
- Redundancy introduced from primary as well standbys
- RTO and RPO balance achieved per organizational requirements
- Point-in-time recovery

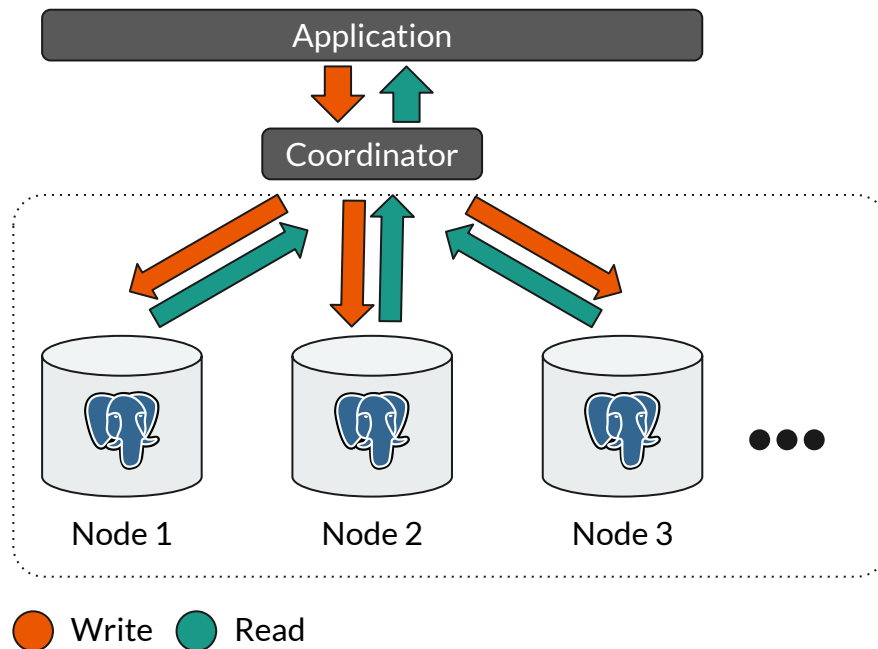
Multi-node clusters with Active-Active configuration*



- Shared-Everything architecture
- Load balancing for read as well as write operations
- Database redundancy to achieve high availability
- Asynchronous replication between nodes for better efficiency

**with conflict resolution at the application layer*

Multi-node clusters with data sharding and horizontal scaling



- Shared-Nothing architecture
- Automatic data sharding based on defined criteria
- Read and write operations are auto directed to the relevant node
- Each node can have its own standbys for high availability

Globally distributed clusters



- Spin up clusters on the public cloud, private cloud, on-prem, bare metal, VMs, or a hybrid of all the above
- Geo fencing for regulatory compliance
- High availability across data centers and geographies

Replication - synchronous vs asynchronous

Synchronous


- Data is transferred immediately
- Transaction waits for confirmation from replica before it commits
- Ensures data consistency across all nodes
- Performance overhead caused by latency
- Used where data accuracy is critical, even at the expense of performance

Asynchronous

- Data may not be transferred immediately
- Transaction commits without waiting for confirmation from replica
- Data may be inconsistent across nodes
- Faster and more scalable
- Used where performance matters more than data accuracy

Challenges in Clustering

- Split brain
- Network Latency
- False alarms
- Data inconsistency




*His decisions aren't any better than yours
— but they're WAY faster...*

Challenges in Clustering

- Split brain
 - Network Latency
 - False alarms
 - Data inconsistency

Split Brain



- **Defined:** Node in a highly available cluster lose connectivity with each other but continue to function independently
- **Challenge:** More than one node believes that it is the primary leading to inconsistencies and possible data loss

Split Brain - Prevention

- Quorum based voting
 - Majority nodes must agree on primary
 - Cluster stops if quorum is not achieved
- Redundancy and failover
 - Prevents a single point of failure
- Proper configuration
- Split brain resolver
 - Software based detection & resolution
 - Can shut down nodes in case of scenario
- Network segmentation
 - Physical network separation
 - Avoid congestion

Split Brain - Resolution

- Witness node
 - Third node that acts as a tie-breaker
 - The winner acts as primary, the loser is shut down
- Consensus algorithm
 - Paxos and Raft protocols are popular
- Manual resolution
 - DBA observes which nodes are competing to be the primary
 - Takes action based on best practices learnt from experience

Challenges in Clustering

- Split brain
- Network Latency
- False alarms
- Data inconsistency

Network Latency

- **Defined:** Time delay when data is transmitted from one point to another
- **Challenge:** Delayed replication can result in data loss. Delayed signals can trigger a false positive.

Network Latency - Preventing False Positives

- Deploy nodes in proximity
 - Reduces time delay and network hops
- Implement quorum-based system
 - If quorum isn't achieved failover won't occur
- High speed, low latency networking
 - High quality hardware and associated software
- Optimize database configuration
 - Parameter tuning based on workloads
 - max_connections, tcp_keealive_idle, ...
- Alerting system
 - Detect and alert admins of possible issues to preempt false positives

Challenges in Clustering

- Split brain
- Network Latency
- **False alarms**
- Data inconsistency

False Alarms

- **Defined:** A problem is reported, but in reality, there is no issue
- **Challenge:** Can trigger a failover when one isn't required, leading to unnecessary disruptions and impacting performance

False Alarms - Prevention

- Proper configuration
 - Best practices, past experience, and some hit & trial is required to ensure that the thresholds are configured appropriately
- Regular maintenance
 - Latest version of software and firmware to be used in order to avoid known bugs and exploits
- Regular testing
 - Testing of various use cases can help identify bottlenecks and possible misconfigurations
- Multiple monitoring tools
 - Multiple tools can help cross-reference alerts and confirm if failover is required

False Alarms - Resolution

In case a false alarm is triggered ...

- Check logs
 - Check the logs of cluster nodes, network devices, and other infrastructure components to identify any anomalies
- Notify stakeholders
 - Notify all stakeholders involved in the cluster's operations to prevent any unnecessary action
- Monitor cluster health
 - Monitor to cluster's health closely to ensure that it is functioning correctly and no further false alarms are triggered

Challenges in Clustering

- Split brain
- Network Latency
- False alarms
- Data inconsistency

Data Inconsistency

- **Defined:** Situations where data in different nodes of a cluster becomes out of sync, leading to inconsistent results and potential data corruption
- **Challenge:** Inaccurate query results that vary based on which node is queried. Such issues are very hard to debug.

Data Inconsistency - Prevention

- Synchronous replication
 - Ensures data is synchronized across all nodes before it is committed
 - Induces a performance overhead
- Load balancer
 - Ensure that all queries from a specific application are sent to the same node
 - Relies on eventual consistency of the cluster
- Monitoring tools
 - Help with early identification of possible issues so you can take preventive measures
- Maintenance windows
 - Minimize data disruption with planned downtime
- Regular testing
 - Test production use cases regularly to ensure the cluster is behaving as expected

Data Inconsistency - Resolution

- Resynchronization
 - Manual sync between nodes to correct data inconsistencies
- Rollback
 - Point-in-time recovery using tools like `pg_rewind` and `pg_backrest`
- Monitoring
 - Diligent monitoring to ensure that the issue doesn't recur

This all sounds really hard



Open source clustering tools for PostgreSQL

- repmgr
 - <https://repmgr.org/>
 - GPL v3
 - Provides automatic failover
 - Manage and monitor replication
- pgpool-II
 - <https://pgpool.net/>
 - Similar to BSD & MIT
 - Middleware between PostgreSQL and client applications
 - Connection pooling, load balancing, caching, and automatic failover
- Patroni
 - <https://patroni.readthedocs.io/en/latest/>
 - MIT
 - Template for PostgreSQL high availability clusters
 - Automatic failover, configuration management, & cluster management

Questions?



 pg_umair

