



Prague PostgreSQL
Developer Day

2023 • 02 • 01

Failsafe Patroni 3.0

Presented by

Alexander Kukushkin &
Polina Bungina



About us

Alexander Kukushkin

- Principal Software Engineer [@Microsoft](#)
- The Patroni guy
- akukushkin@microsoft.com
- Twitter: @cyberdemn

Polina Bungina

- Software Engineer [@ZalandoTech](#)
- polina.bungina@zalando.de
- Twitter: @hugh_capet



Agenda

Introduction to Patroni

Observer problem

Demo 1

DCS dailsafe feature

Demo 2

Conclusion

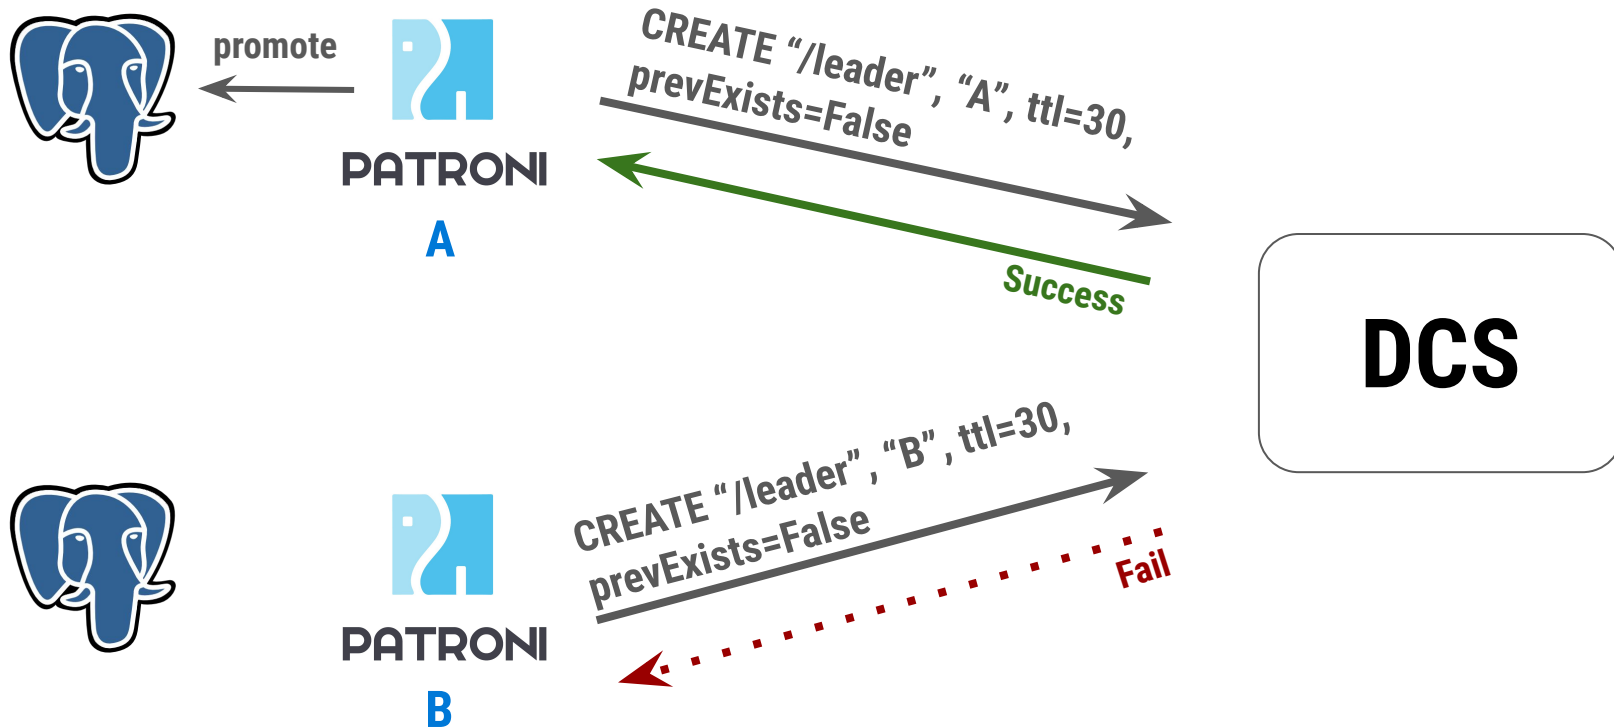
Do we need it at all?

- Service-Level Agreement (**SLA**)
- Recovery point objective (**RPO**)
- Recovery time objective (**RTO**)


Architecture overview

- Cluster state stored in Distributed Configuration Store (**DCS**)
 - ZooKeeper
 - Etcd
 - Consul
 - Kubernetes control-plane
- **Session/TTL** to expire data (i.e. leader key)
- **Atomic CAS** operations
- **Watches** for important keys

Leader race



Normal operational mode



primary



replica




UPDATE `"/leader", "A", ttl=30`
`prevValue="A"`

Success

DCS

WATCH(`"/leader"`)

Normal operational mode



replica



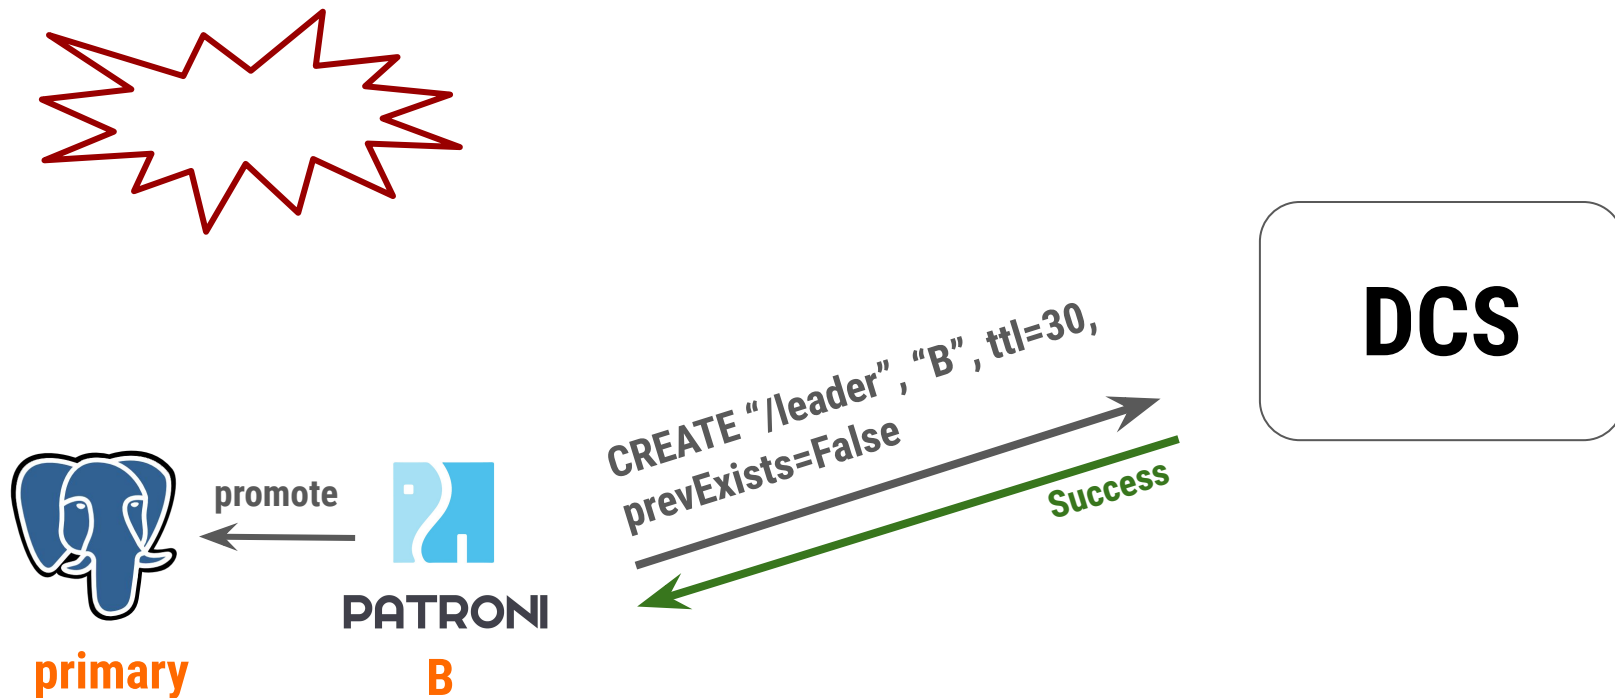
PATRONI

B


NOTIFY("/leader", expired=True)

DCS

Normal operational mode



DCS can't be accessed



primary



replica



A




B


UPDATE `"/leader"`, `"A"`, `ttl=30`
`prevValue="A"`




Fail



DCS can't be accessed



primary




replica




UPDATE `"/leader"`, `"A"`, `ttl=30`
`prevValue="A"`

Success

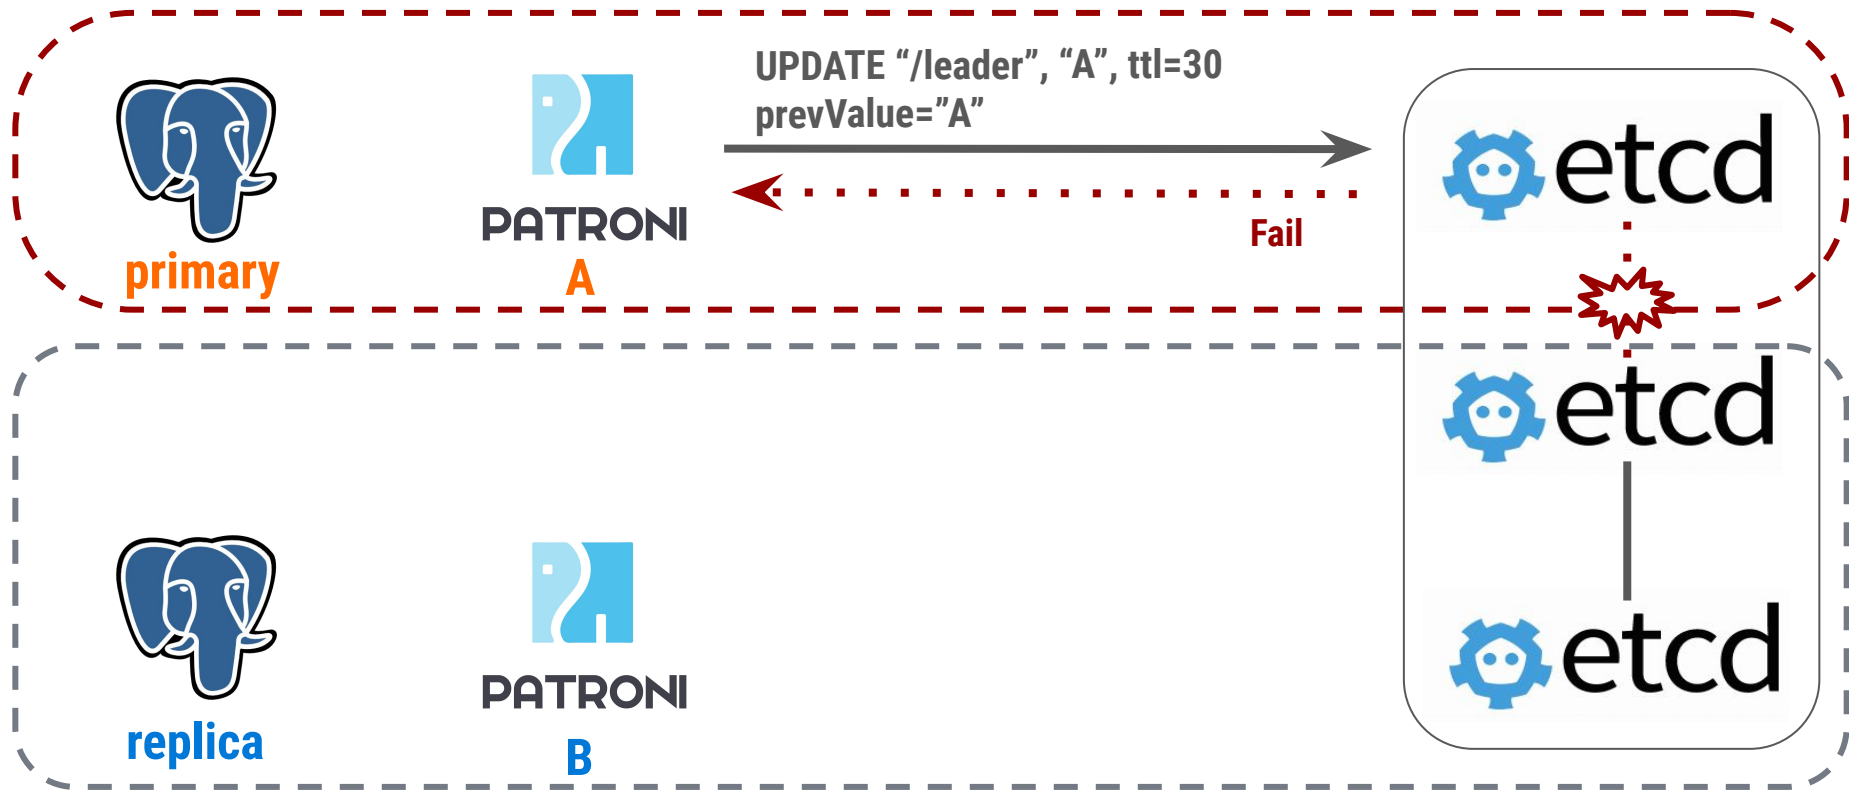


Why did update fail?

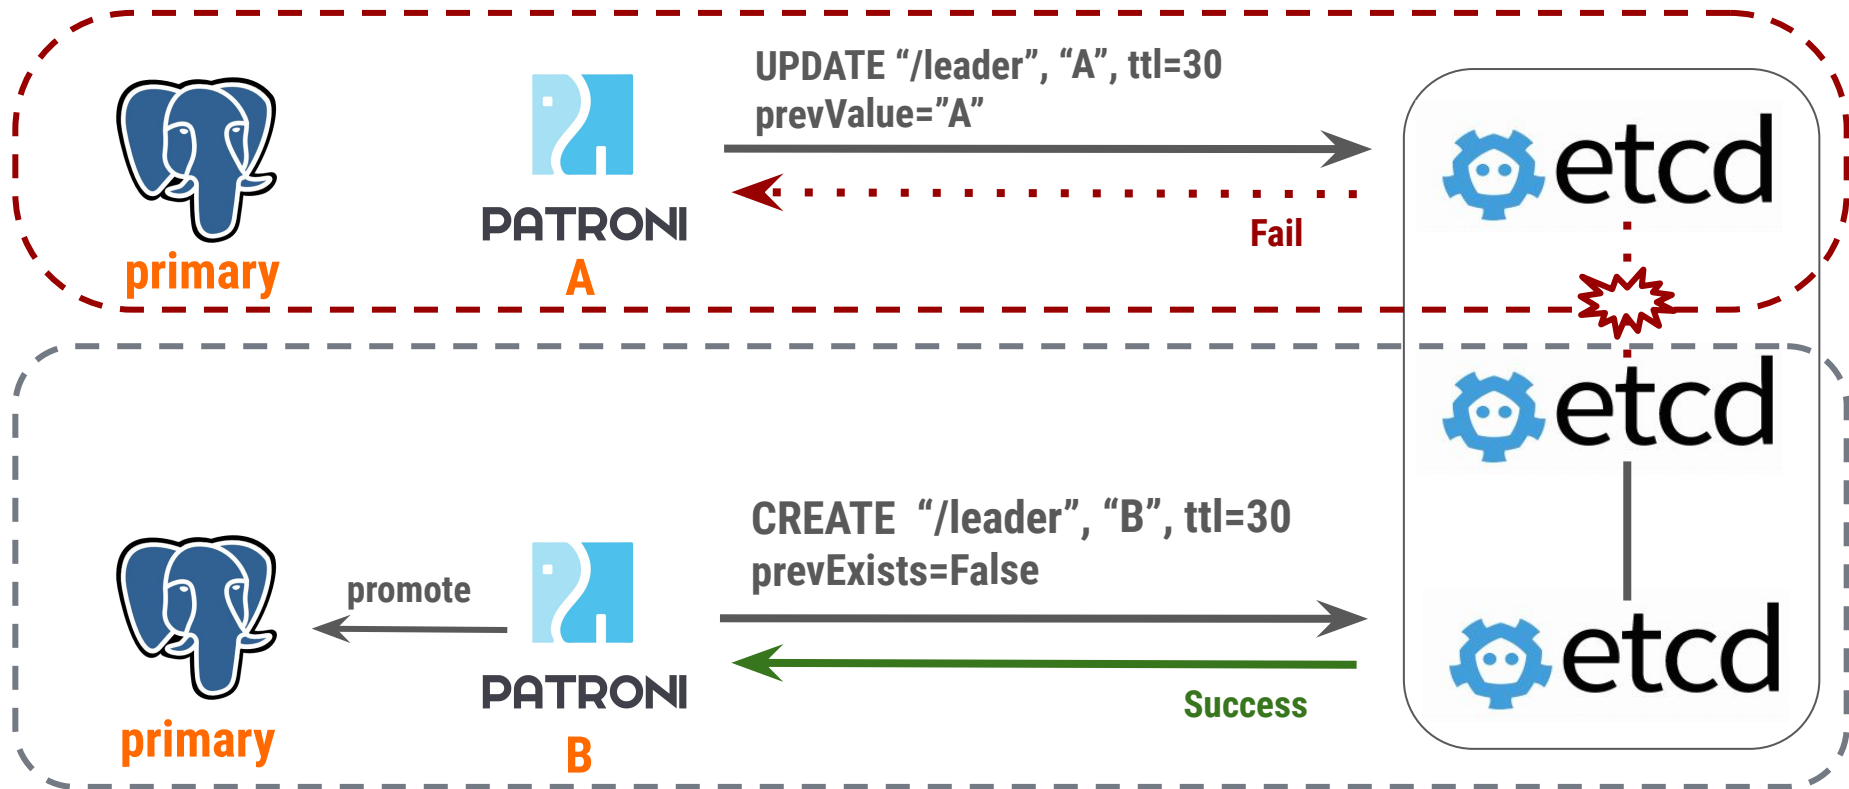
- DCS is down?
- Network issues?



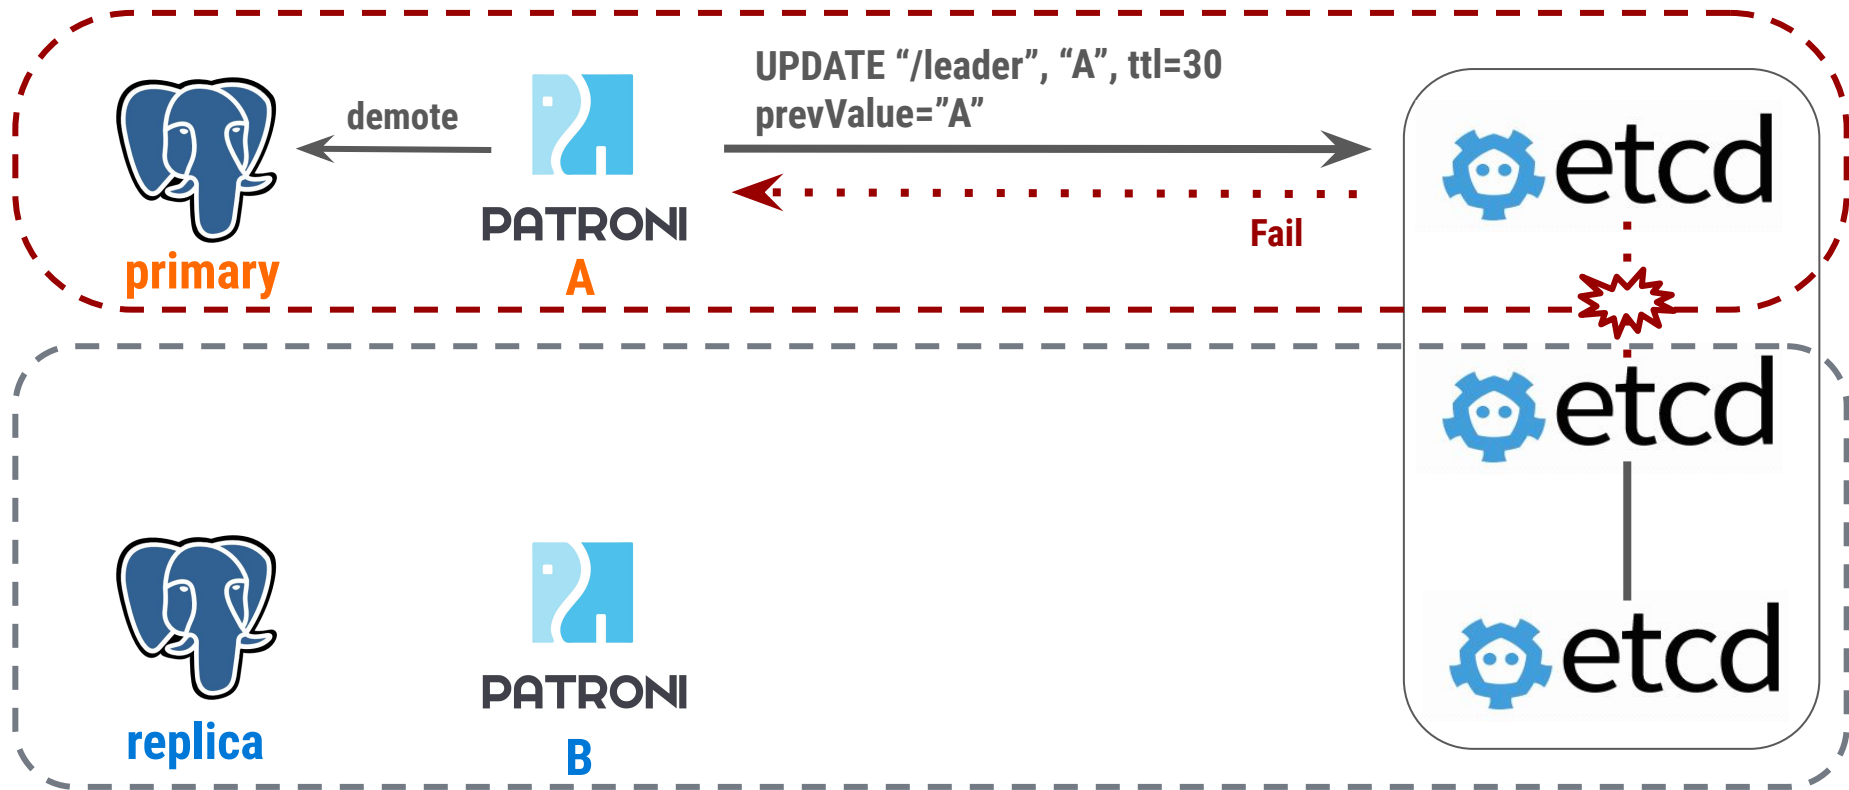
Network partition



Leader key expired




So, to be on the safe side...



So, to be on the safe side...



Still not perfect



CREATE `"/leader", "B", ttl=30`
`prevExists=False`

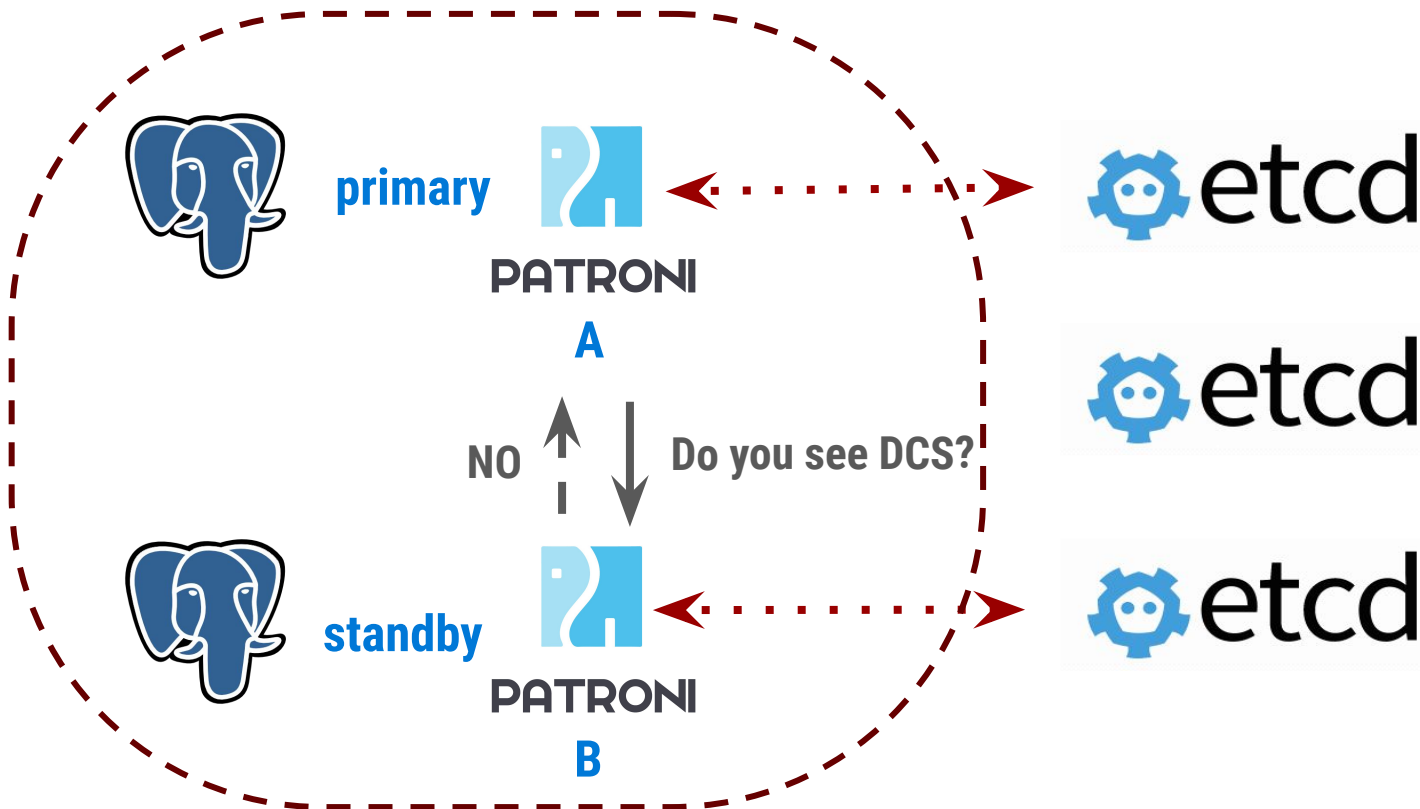




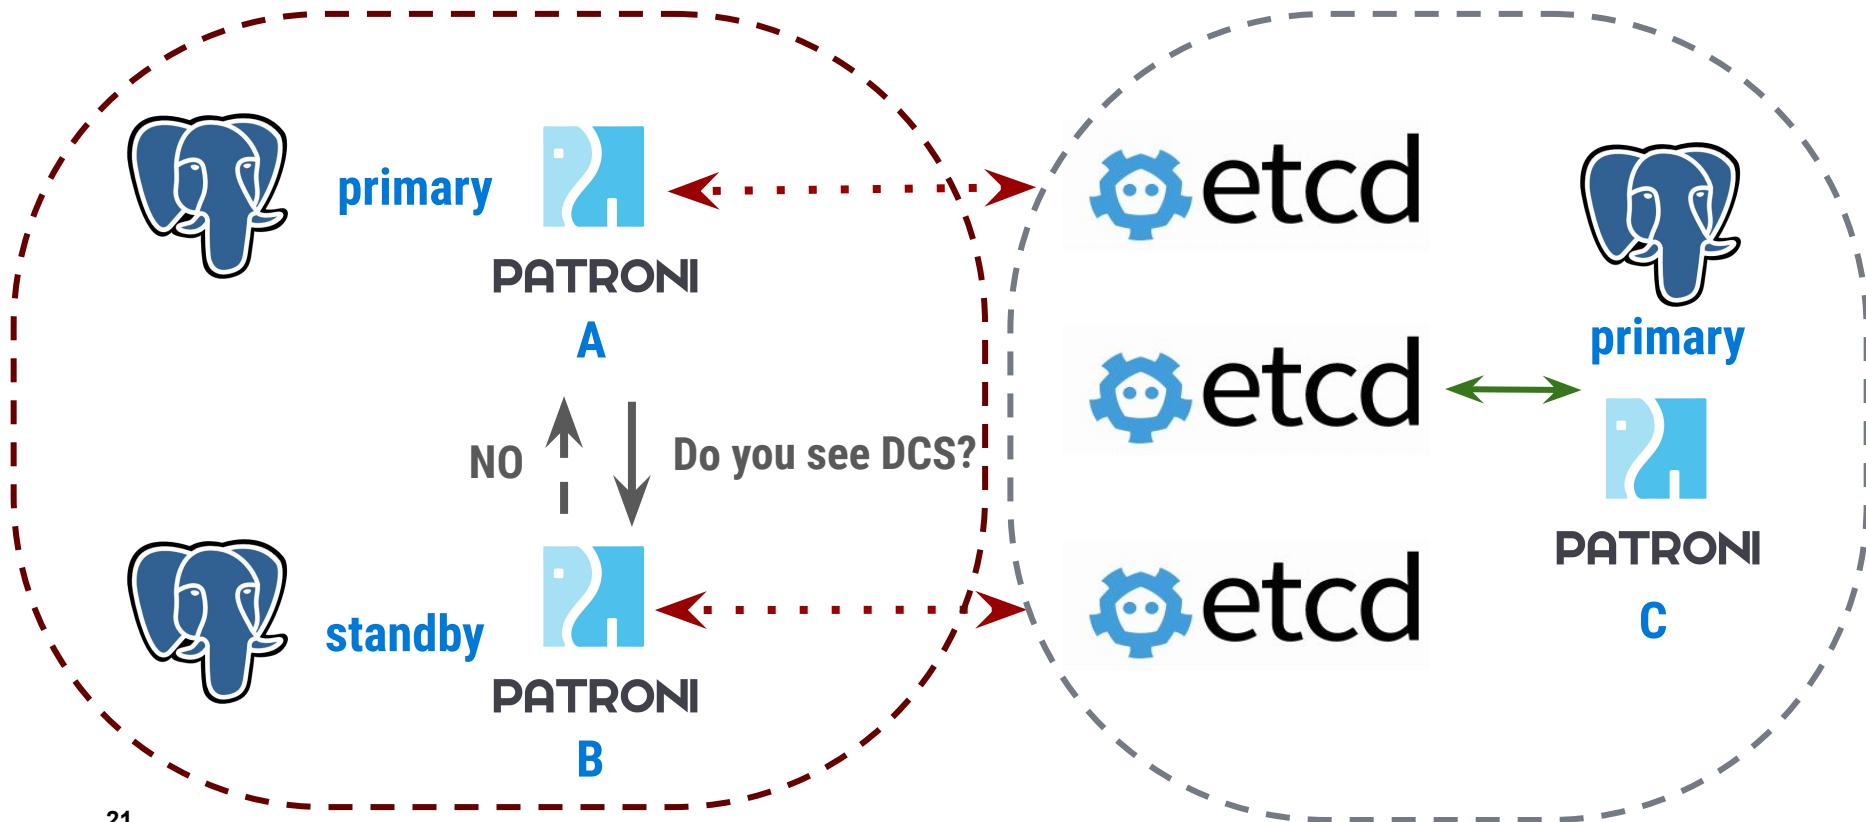
DCS down

- **Etcd, Zookeeper** - very unlikely (if configured correctly)
- **Consul** - local agent is a SPoF!
- **Kubernetes control-plane** - typical SLA for managed services is 99.95% (4h22m per year)

What if...



Split-brain!




Idea

- Continue to run as primary if can see ALL Patroni nodes
- Don't allow “unknown” nodes to become primary!

“Unknown” node?

- Patroni clusters are mainly “static”, but nodes can join and leave
- If topology changes - write list of Patroni nodes names to DCS
- Nodes outside of this list are “**unknown**” and not allowed to become primary


DCS failsafe mode




PATRONI
node1

UPDATE `"/leader", "A", ttl=30`
→
← Fail

1



4 200 OK ↑ POST /failsafe ↓ 2


PATRONI [cache primary data for ttl]
nodeN 3

`/failsafe: node1, node2, ..., nodeN`

Implementation details

- Introduce **/failsafe** key - list of **currently presented members** in the cluster
 - Maintained by the leader
 - Cache its value in Patroni (on all nodes)
- Introduce **POST /failsafe** REST API endpoint
 - Payload contains information about primary and permanent logical slots
 - Primary checks response code and demotes if not 200

Implementation details (continue)

- Replica disqualifies itself from the leader race if not listed in the DCS **/failsafe** key
- Primary executes the failsafe check only with nodes from the **failsafe** list
 - Continue as primary if ALL nodes are accessible
 - Otherwise demote
- Replicas call **pg_advance_replication_slot()** if necessary.

How to enable failsafe mode

```
$ patronictl edit-config
```

```
---
```

```
+++
```

```
@@ -4,3 +4,4 @@
```

```
    use_pg_rewind: true
```

```
retry_timeout: 10
```

```
ttl: 30
```

```
+failsafe_mode: on
```

```
Apply these changes? [y/N]: y
```

```
Configuration changed
```

```
$ etcdctl get /service/batman/failsafe
```

```
{  
  "postgresql0": "http://127.0.0.1:8008/patroni",  
  "postgresql1": "http://127.0.0.1:8009/patroni"  
}
```

```
$ curl http://127.0.0.1:8008/failsafe
```

```
{  
  "postgresql0": "http://127.0.0.1:8008/patroni",  
  "postgresql1": "http://127.0.0.1:8009/patroni"  
}
```

Monitoring

```
$ curl -s http://127.0.0.1:8008/patroni | jq .
```

```
{
  "state": "running",
  "postmaster_start_time": "2023-01-26 16:11:04.848424+00:00",
  "role": "master",
  "server_version": 150001,
  "xlog": {"location": 67419584},
  "timeline": 2,
  "replication": [
    {"username": "replicator", "application_name": "postgresql1", "client_addr": "127.0.0.1",
      "state": "streaming", "sync_state": "async", "sync_priority": 0}],
  "cluster_unlocked": true,
  "failsafe_mode_is_active": true,
  "dcs_last_seen": 1674749503,
  "database_system_identifier": "7192993973708324892",
  "patroni": {"version": "3.0.0", "scope": "demo"}
}
```



When not to use it

- When nodes could change their names after “restart” (with old storage)
 - If ALL nodes are restarted at the same time cluster will not recover automatically

Example:

- K8s deployment without StatefulSet
 - Crunchy Postgres Operator (PGO)



Thank you!
Questions?