



# 第九届PostgreSQL中国技术大会

## 2019 PostgreSQL Conference China

开源驱动 自主研发

主办:  PostgreSQL中文社区

协办: **ITPUB**

🕒 2019年11月29日-30日

📍 北京维景国际大酒店



第九届PostgreSQL中国技术大会  
2019 PostgreSQL Conference China

# Fault injection to test PostgreSQL code

Asim R P

[apraveen@pivotal.io](mailto:apraveen@pivotal.io)

Pune, India



# Agenda

- What is lacking in current testing frameworks?
- What is fault injection?
- Proposal to inject faults
- Tests to demonstrate proposed faultinjector patch



# Testing frameworks in PostgreSQL

- `pg_regress (src/test/regress)`
  - A test is written as a `.sql` file
  - Expected output is saved as an answer file `(.out)`
  - `pg_regress` runs a `.sql` file, compares the output with the answer
- `isolation (src/test/isolation)`
  - tests concurrent transactions
  - `.spec` file - SQL commands within each transaction and how to interleave them



# isolation: test that select waits for alter

session “s1”

```
setup { BEGIN ISOLATION LEVEL READ COMMITTED; }
```

```
step “alter1” { ALTER TABLE test ADD RENAME TO test2; }
```

```
step “commit1” { COMMIT; }
```

session “s2”

```
setup { SET default_transaction_isolation = read committed; }
```

```
step “select2” { SELECT * from test; }
```

```
permutation “alter1” “s2” “commit1”
```



# Testing frameworks in PostgreSQL

- TAP - tests written in perl (Test::More)
  - Orchestrate a cluster - master and one or more standby
  - initdb, start / stop cluster, etc.
  - src/test/perl/README
- src/test/modules (not run in CI)
  - modular testing of a specific component, e.g. planner
  - src/test/modules/README





# What is lacking in PostgreSQL tests?

- Test that a backend is killed after writing commit record but before updating the transaction status in pg\_clog
- Is synchronous replication really synchronous?
- Was a cached plan reused?
- Anything comes to your mind?



# Fault injection

- Unconference discussion at PGcon 2019 in Ottawa
- Already being used in Greenplum
- Patch proposed on pgsql-hackers:
  - <https://www.postgresql.org/message-id/flat/CANXE4TdxDESX1jKw48xet-5GvBFVSq%3D4cgNeioTQff372KO45A%40mail.gmail.com>





## Fault point (./configure CPPFLAGS=-DFAULT\_INJECTOR)

```
--- a/src/backend/executor/execIndexing.c
+++ b/src/backend/executor/execIndexing.c
@@ -289,6 +289,7 @@ ExecInsertIndexTuples(TupleTableSlot *slot,
                        bool
                        isnull[INDEX_MAX_KEYS];

    Assert(ItemPointerIsValid(tupleid));
+   SIMPLE_FAULT_INJECTOR("insert_index_tuples");

    /*
     * Get information from the result relation info structure.
```



# Fault point - macro definition

```
#ifdef FAULT_INJECTOR
#define SIMPLE_FAULT_INJECTOR(FaultName) \
    FaultInjector_TriggerFaultIfSet(FaultName, "", "")
#else
#define SIMPLE_FAULT_INJECTOR(FaultName)
#endif
```



```
--- a/src/backend/access/heap/heapam.c
+++ b/src/backend/access/heap/heapam.c
@@ -1875,6 +1875,12 @@ heap_insert(Relation relation, HeapTuple tup, CommandId cid,
         Buffer                vmbuffer = InvalidBuffer;
         bool                   all_visible_cleared = false;

+#ifdef FAULT_INJECTOR
+    FaultInjector_TriggerFaultIfSet(
+        "heap_insert",
+        "" /* database name */,
+        RelationGetRelationName(relation));
+#endif

/*
 * Fill in tuple header fields and toast the tuple if necessary.
 *
```



# Set a fault using its name

- `CREATE EXTENSION faultinjector;`
- Enable the “heap\_insert” fault when a tuple is inserted into “mytable”
  - `SELECT inject_fault('heap_insert', 'error', '', 'mytable', ...);`
- Wait until the fault is triggered (blocking call)
  - `SELECT wait_until_triggered_fault('heap_insert', 1);`
- Check the status
  - `SELECT inject_fault('heap_insert', 'status');`
- Reset the fault
  - `SELECT inject_fault('heap_insert', 'reset');`



# More on inject\_fault

inject\_fault( ... ,

<start\_occurrence> ,

trigger only after the fault point is reached as many  
number of times

<end\_occurrence> ,

stop after triggering as many number of times (fault  
state completed)

<sleep\_seconds>

seconds to sleep in a sleep fault)



# inject\_fault\_remote()

- Same as inject\_fault()
- Additional args: hostname and port number
- New libpq message to inject faults
- Useful for testing standby servers in a cluster





# Fault types

- error: elog(ERROR) leads to transaction abort
- skip: do nothing - indicates that the fault point was reached, also used for custom action
- suspend
- reset
- status



```
--- a/src/backend/access/transam/xlog.c
+++ b/src/backend/access/transam/xlog.c
@@ -8537,6 +8537,14 @@ CreateCheckpoint(int flags)
     VirtualTransactionId *vxids;
     int
         nvxids;

+#ifdef FAULT_INJECTOR
+    if (SIMPLE_FAULT_INJECTOR("checkpoint") == FaultInjector_FaultTypeSkip)
+    {
+        /* Custom logic here ... */
+        return;
+    }
+#endif
+
+    /*
+     * An end-of-recovery checkpoint is really a shutdown checkpoint, just
+     * issued at a different time.
```



# Fault states

- injected but not triggered - `SELECT inject_fault();`
- completed - triggered maximum number of times, will no longer trigger
- triggered - reached during execution at least once



# Faults offer fine grained control

- Enable complex testing scenarios
- Was a specific flag in shared memory set, and when?
- Was a branch in a function taken?
- Can be used in regress, isolation as well as TAP tests



# Demonstrative test: speculative insert

```
CREATE TABLE test(key TEXT, data TEXT);
```

```
CREATE UNIQUE INDEX ON test(key);
```

```
T1: INSERT INTO test values ('k1', 'inserted T1') ON CONFLICT  
DO UPDATE SET data = test.data || ' conflict update T1';
```

```
T2: INSERT INTO test values ('k1', 'inserted T2') ON CONFLICT  
DO UPDATE SET data = test.data || ' conflict update T2';
```



# Demonstrative test: speculative insert

Conflicts detected after a tuple is inserted into heap but before inserting into index are handled properly

- T1 goes first and inserts the tuple into heap
- Before T1 updates the index, T2 inserts into both, heap and the index.
- T2 sees no conflicts because T1 is still in progress.
- T1 moves ahead. T1 should detect a conflict and abort.





# Demonstrative test: speculative insert

- Find it in the faultinjector patch:
  - `src/test/isolation/specs/insert-conflict-with-faults.spec`



# Demonstrative test: synchronous replication

Ensure that commits on master block until a synchronous standby acknowledges the commit

- Master writes a commit record
- The backend process on master starts waiting for standby to flush WAL upto the commit LSN
- Standby goes down before the WAL receiver writes WAL upto the commit LSN



# Demonstrative test: synchronous replication

- Find it in the faultinjector patch, it's a TAP test:
  - `src/test/recovery/t/007_sync_rep.pl`

The background is a dark blue gradient. In the center, there is a faint, wireframe-style globe. Overlaid on this globe and the entire background is a complex network of thin, light blue lines that radiate outwards from various points, some of which are small dots. The overall aesthetic is technological and digital.

**THANK  
S**