

# { JSON }

## in MySQL and MariaDB Databases

Federico Razzoli




# € whoami

- Federico Razzoli
- Freelance consultant
- Working with databases since 2000

hello@[federico-razzoli.com](mailto:federico-razzoli.com)  
[federico-razzoli.com](http://federico-razzoli.com)

- I worked as a consultant for Percona and Ibuildings (mainly MySQL and MariaDB)
- I worked as a DBA for fast-growing companies like Catawiki, HumanState, TransferWise



# JSON Support

# JSON Support

- MySQL 7.0 has a JSON type and JSON functions
- MariaDB 10.2 has JSON alias (LONGTEXT) and JSON functions
  - From MariaDB 10.4 JSON columns must contain valid JSON data

# JSON Type Benefits

- Persons from both MySQL and MariaDB state that their approach is faster
  - Surprise, surprise!
- MySQL format is smaller
- When updating a property, MySQL does not rewrite the whole JSON document
- MySQL binary log only contains the modified properties
  - `binlog_row_value_options = 'PARTIAL_JSON'`
- However, in MySQL 8, make sure that:
  - `Default_tmp_storage_engine = 'TEMPTABLE'`

# JSON Features Comparison

- Let's compare JSON features, and more generic features that are useful working with JSON
- MySQL 8.0 / MariaDB 10.5
- MySQL only:
  - Index arrays (requires multi-valued indexes)
  - Schema validation
  - JSON\_TABLE()
  - Functional index / Index on VIRTUAL column
- MariaDB only:
  - CONNECT (use a JSON file as a table)
  - Stored aggregate functions

# Use Cases



# JSON in relational databases?

- Relational databases are **schemaful**
- All rows must have the same columns (name, type)
  - (though `NULL` is often used to relax the schema)
- `CHECK` constraints:
  - `email VARCHAR(100) CHECK (email LIKE '_%@_%._%')`
  - `birth_date DATE, death_date DATE, CHECK (birth_date <= death_date)`
- `UNIQUE` constraint
- Referential constraints (foreign keys):
  - Each row in a *child* table matches a row in a *parent* table



# JSON in relational databases?

- JSON is usually **schemaless**
- Two different objects can have different structures:

```
{  
  "product-type": "shirt",  
  "cost": 10.0,  
  "size": "M"  
}  
  
{  
  "product-type": "phone",  
  "cost": 69.99,  
  "size": [6.37, 2.9, 0.3],  
  "os": "Android"  
}
```

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?
- One table per product type?
  - Cannot enforce UNIQUE constraint
  - Try to find a product that could be of any type...
  - Try to get the min and max costs...

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?
- One table with one column for each property, NULL where it is not used?
  - Adding a product implies a slow, painful ALTER TABLE
  - If products type is deleted, obsolete columns tend to remain forever
  - Big table is bad for operations (backup, repair...)
  - INSERTs are slow
  - SELECT \*

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?
- One main table with common properties, and one separate table for each product type?
  - You can now see min/max costs, or find a product of any type
  - Many JOINS

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?
- Entity-Attribute-Value pattern (EAV)?
  - Store all data as texts
  - Crazy amount of JOINS

# JSON in relational databases?

- How would you store heterogeneous products in a relational database?
- Single table, where all non-common property are stored as JSON
  - Still not perfect
  - But it's good enough

# JSON in relational databases?

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50) NOT NULL,  
    name VARCHAR(50) NOT NULL,  
    cost DECIMAL(10, 2) NOT NULL,  
    quantity INT UNSIGNED NOT NULL,  
    attributes JSON NOT NULL  
);
```



# Indexing Properties

# Materialising Properties

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50),  
    name VARCHAR(50) NOT NULL,  
    cost DECIMAL(10, 2) NOT NULL,  
    quantity INT UNSIGNED NOT NULL,  
    attributes JSON NOT NULL,  
    colour VARCHAR(50) AS  
        (JSON_UNQUOTE(  
            JSON_EXTRACT(attributes, '$.colour')  
        )) STORED  
);
```

# MySQL shortcut

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50),  
    name VARCHAR(50) NOT NULL,  
    cost DECIMAL(10, 2) NOT NULL,  
    quantity INT UNSIGNED NOT NULL,  
    attributes JSON NOT NULL,  
    colour VARCHAR(50) AS  
        (attributes->>'$.colour') VIRTUAL  
);
```

# Index on a Property

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50),  
    name VARCHAR(50) NOT NULL,  
    cost DECIMAL(10, 2) NOT NULL,  
    quantity INT UNSIGNED NOT NULL,  
    attributes JSON NOT NULL,  
    colour VARCHAR(50) AS  
        (attributes->>'$.colour') VIRTUAL,  
    INDEX idx_colour (type, colour)  
);
```

# Index on a Property

```
SELECT id FROM product
  WHERE type = 'FURNITURE'
  AND colour = 'grey'
;
```

-- only uses the index in MySQL

```
SELECT id FROM product
  WHERE type = 'FURNITURE'
  AND attributes->>'$.colour' = 'grey'
;
```

# Indexing Properties

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50),  
    name VARCHAR(50) NOT NULL,  
    cost DECIMAL(10, 2) NOT NULL,  
    quantity INT UNSIGNED NOT NULL,  
    attributes JSON NOT NULL,  
    furniture_colour VARCHAR(50) AS  
        (IF(type = 'FURNITURE',  
            attributes->>'$.colour',  
            NULL  
        )) VIRTUAL,  
    INDEX idx_colour (furniture_colour)  
);
```

# Indexing Arrays

# Indexing ARRAYS

- MySQL and MariaDB historically don't support arrays
- Now they can use JSON arrays as “regular” arrays
- MySQL 8.0 supports Multi-Valued Indexes, that allow to index these arrays
  - Can be used to index JSON objects, *not covered here*



# Indexing Arrays

```
CREATE TABLE `order` (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    customer_id INT UNSIGNED NOT NULL,  
    product_ids JSON NOT NULL,  
    INDEX idx_products (  
        (CAST(product_ids AS UNSIGNED ARRAY))  
    )  
);  
  
INSERT INTO `order` (customer_id, product_ids) VALUES  
    (24, JSON_ARRAY(10, 20, 30));
```

# Indexing Arrays

```
SELECT DISTINCT customer_id  
  FROM `order`  
 WHERE 20 MEMBER OF (product_ids);
```

- JSON\_CONTAINS()
- JSON\_OVERLAPS()

# Relationships

# Foreign Keys on Properties

```
CREATE TABLE product (  
    id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,  
    type VARCHAR(50),  
    ...  
    furniture_colour INT UNSIGNED AS  
        (IF(type = 'FURNITURE',  
            attributes->>'$.colour',  
            NULL  
        )) STORED,  
    INDEX idx_furniture_colour (furniture_colour),  
    FOREIGN KEY fk_furniture_colour (furniture_colour)  
        REFERENCES furniture_colour (id)  
);
```

# Not suggesting to use Foreign Keys

- They come with several penalties
  - Slower writes
  - Propagated locks (for CASCADE FKs)
  - Limitations (only InnoDB, no partitions)
  - Cannot run **safe** online migrations
- Google for:  
Federico Razzoli "Foreign Key bugs in MySQL and MariaDB"

# But...

- But you can do the same without foreign keys
- You can easily add extra columns to `furniture_colour`
- You don't have to mind about colour name details (case, extra spaces...)
- You can easily change a colour name
- You keep the JSON documents smaller

Basically, you can use the relational databases philosophy  
While keeping some semi-structured data inside them

# Validating JSON documents



# Validating Against a JSON Schema

- Both MySQL and MariaDB only accept valid JSON documents in JSON columns
- The `CHECK` clause defines the rules that determine if a row is valid
- `JSON_SCHEMA_VALID()` in MySQL validates a JSON document against a JSON Schema
  - [json-schema.org](https://json-schema.org)
  - Draft 4
  - But from my tests we can use drafts 6 or 7



# Validating Schema

```
SELECT JSON_SCHEMA_VALID(@schema, @document);
```

```
SELECT JSON_SCHEMA_VALID((  
    SELECT schema  
    FROM schemas  
    WHERE type = 'furniture'  
) , document)  
FROM product  
WHERE id = 24  
;
```

# Validating Schema

Suppose we want to validate objects of this type:

```
{  
  "material": "iron",  
  "colour": "grey",  
  "size": [1.5, 1.5, 1.0]  
}
```

```
{  
  "$id": "https://federico-razzoli.com/schemas/furniture",  
  "$schema": "http://json-schema.org/draft-07/schema#",  
  "type": "object",  
  "properties": {  
    "material": {  
      "type": "string",  
      "enum": ["iron", "wood"]  
    },  
    "size": {  
      "type": "array",  
      "items": {  
        "type": "number",  
        "minItems": 3,  
        "maxItems": 3  
      }  
    },  
    "colour": {  
      "type": "string"  
    }  
  },  
  "required": ["material", "size"]  
}
```

# Debugging a JSON Schema

- `SELECT JSON_SCHEMA_VALID() + SHOW WARNINGS`
- `JSON_SCHEMA_VALIDATION_REPORT()`

# Default Properties

# Default Properties

- We can have a JSON document with the default properties
- And only store non-default values in each regular JSON document
- This is useful:
  - For large JSON objects (eg. software configuration)
  - For objects whose defaults may change (eg. software whose settings can be set per user and per team)

# Merging Documents

```
-- get user settings
SET @user_conf := (
    SELECT json_settings FROM user_configuration
    WHERE id = 24
);

-- get team settings
SET @team_conf := (
    SELECT json_settings FROM team_configuration
    WHERE id = 123
);

-- merge them. user settings overwrite team settings
SELECT JSON_MERGE_PATCH(@team_conf, @user_conf);
```

# JSON\_MERGE\_PATCH() example

```
mysql> SELECT JSON_MERGE_PATCH(  
    '{"items-per-page": 24, "suppress_warnings": true}',  
    '{"suppress_warnings": false}')
```

AS configuration;

```
+-----+  
| configuration |  
+-----+  
| {"items-per-page": 24, "suppress_warnings": false} |  
+-----+
```



JSON\_ARRAYAGG()  
JSON\_OBJECTAGG()

# JSON\_ARRAYAGG(), JSON\_OBJECTAGG()

- Available in:
  - MySQL 5.7
  - MariaDB 10.5, currently not GA
- However, they are:
  - Simple to emulate with `GROUP_CONCAT()`
  - Easy to code with MariaDB **Stored Aggregate Functions** (10.3)

# JSON\_ARRAYAGG()

```
mysql> SELECT type, JSON_ARRAYAGG(colour) AS colours
-> FROM product
-> GROUP BY type
-> ORDER BY type;
```

type	colours
chair	["green", "blue", "red", "black"]
shelf	["white", "brown"]
table	["white", "gray", "black"]

# JSON\_OBJECTAGG(): one document

```
mysql> SELECT
->   JSON_PRETTY(
->     JSON_OBJECTAGG(variable, value)
->   ) AS settings
-> FROM config \G
***** 1. row *****
settings: {
  "datadir": "/var/mysql/data",
  "log_bin": "1",
  "general_log": "0",
  "slow_query_log": "0",
  "innodb_log_file_size": "1G",
  "innodb_log_buffer_size": "16M",
  "innodb_buffer_pool_size": "16G"
}
```

# JSON\_OBJECTAGG() nested objects

```
mysql> SELECT
->     JSON_PRETTY(
->         JSON_OBJECTAGG(
->             variable,
->             JSON_OBJECT(
->                 'type', type,
->                 'value', value
->             )
->         )
->     ) AS settings
-> FROM config;
```

# JSON\_OBJECTAGG() nested objects

```
settings: {  
  "datadir": {  
    "type": "string",  
    "value": "/var/mysql/data"  
  },  
  "log_bin": {  
    "type": "bool",  
    "value": "1"  
  },  
  ...  
}
```

# JSON\_OBJECTAGG(): one document per row

```
mysql> SELECT JSON_OBJECTAGG(variable, value) AS settings  
-> FROM config  
-> GROUP BY variable;
```

```
+-----+  
| settings                                |  
+-----+  
| {"datadir": "/var/mysql/data"} |  
| {"general_log": "0"} |  
| {"innodb_buffer_pool_size": "16G"} |  
| {"innodb_log_buffer_size": "16M"} |  
| {"innodb_log_file_size": "1G"} |  
| {"log_bin": "1"} |  
| {"slow_query_log": "0"} |  
+-----+
```

# JSON\_TABLE()



# JSON\_TABLE()

- Only MySQL 8
- Transforms a JSON document into a table
- Each property, or some properties, become a table column
- Properties are identified by JSON Path expressions

# JSON\_TABLE() example

```
{
  "customers": [
    {
      "name": "John Doe",
      "emails": ["john@doe.com"],
      "phone": "123456"
    },
    {
      "name": "Jane Dee",
      "emails": ["Jane@dee.com"],
      "phone": "987654"
    },
    {
      "name": "Donald Duck",
      "emails": []
    }
  ],
  "products": [],
  "orders": []
}
```

# JSON\_TABLE() example

```
SELECT *  
  FROM JSON_TABLE(  
    @customers,  
    '$.customers[*]'  
  COLUMNS (  
    id FOR ORDINALITY,  
    name VARCHAR(50) PATH '$.name',  
    primary_email VARCHAR(50) PATH '$.emails[0]',  
    phone VARCHAR(50) PATH '$.phone'  
      DEFAULT "UNKNOWN" ON EMPTY  
      NULL ON ERROR  
  )  
) AS customers;
```

# JSON\_TABLE() example

id	name	primary_email	phone
1	John Doe	john@doe.com	123456
2	Jane Dee	Jane@dee.com	987654
3	Donald Duck	NULL	UNKNOWN

# MariaDB CONNECT

# What CONNECT is

- Connect is a storage engine for MariaDB
  - Won't compile on MySQL or Percona Server
- Reads data from different data sources:
  - Remote tables (MySQL protocol, ODBC, JDBC)
  - Files in various formats (JSON, XML, CSV, custom formats...)
  - Special data sources
  - Transformed data from other tables (pivot, ...)

# CONNECT and JSON

- JSON Table Type (10.1)
- On Windows, can read data from REST APIs (not easy to do)
- CONNECT comes with a library of JSON UDFs
  - Probably not battle-tested as the current built-in functions

# Thank you for listening!

Federico  
[hello@federico-razzoli.com](mailto:hello@federico-razzoli.com)