



Logical replication with pglogical

Moving the same old data around in new
and exciting ways



It's row-oriented replication

Really, that's pretty much it. The rest is details.

Done now?



pglogical

Open source – PostgreSQL license

Submitted to 9.6

Generic, re-usable, no custom PostgreSQL



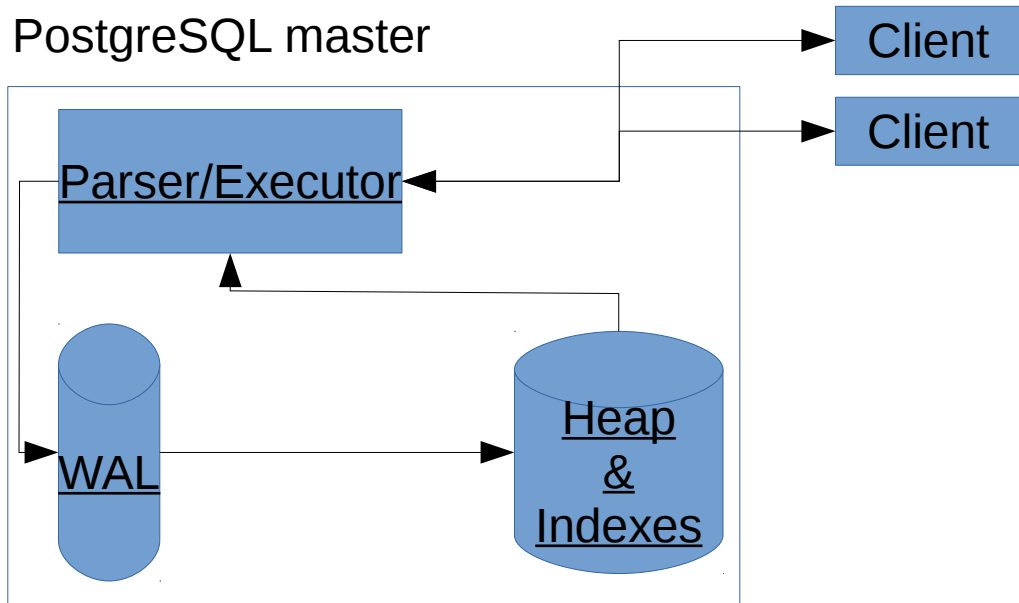
Architectures

- Standalone (no replication)
- Physical replication (block level)
- Logical replication (row level)



Standalone PostgreSQL


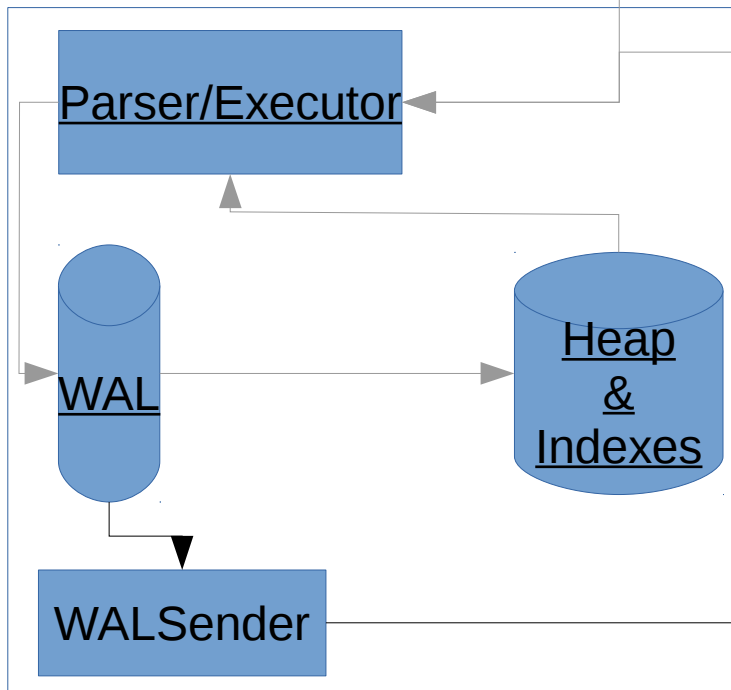
PostgreSQL master



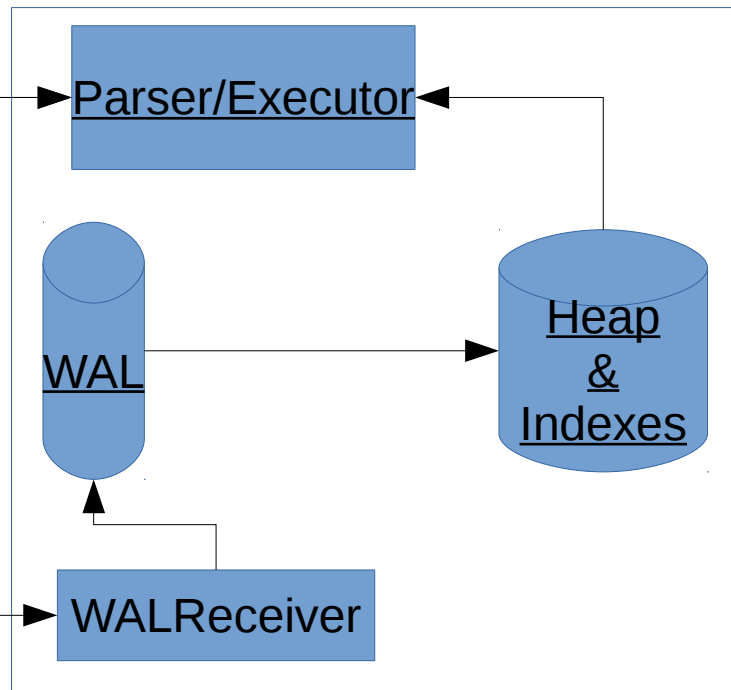


Physical replica & hot standby

PostgreSQL master



PostgreSQL replica / hot standby





Confused yet?



“Physical” replication

Copies *everything*:

- Every database
- VACUUM
- Index updates
-



“Physical” replication

Fast to apply changes to replicas.

Bandwidth-hungry.


All-or-nothing.

Hot standby limitations



Logical decoding

PostgreSQL master





Logical decoding

Collects just *row values*

No VACUUM traffic, index updates, etc

Can generate text-format values for replication to other Pg versions etc



Logical decoding

- Useful for logical replication
- ... but not just replication
 - Intrusion detection
 - Search
 - Messages buses
 -



Many logical decoding plugins

- **pglogical_output**
- **BDR** output plugin
- The demo **test_decoding** to stream SQL
- **decoder_raw** and **receiver_raw** to stream and apply SQL
- github.com/ildus/decoder_json and github.com/leptonix/decoding-json to stream JSON
- github.com/xstevens/decoderbufs to stream as protocol buffers
- github.com/confluentinc/bottledwater-pg to stream to Kafka
-?



pglogical_output

Make it *easy* and *generic*:

Both json & fast native proto

Selective replication, metadata, etc

Use it without writing a bunch of C



Logical *replication*

Selective – just the DBs/tables you want

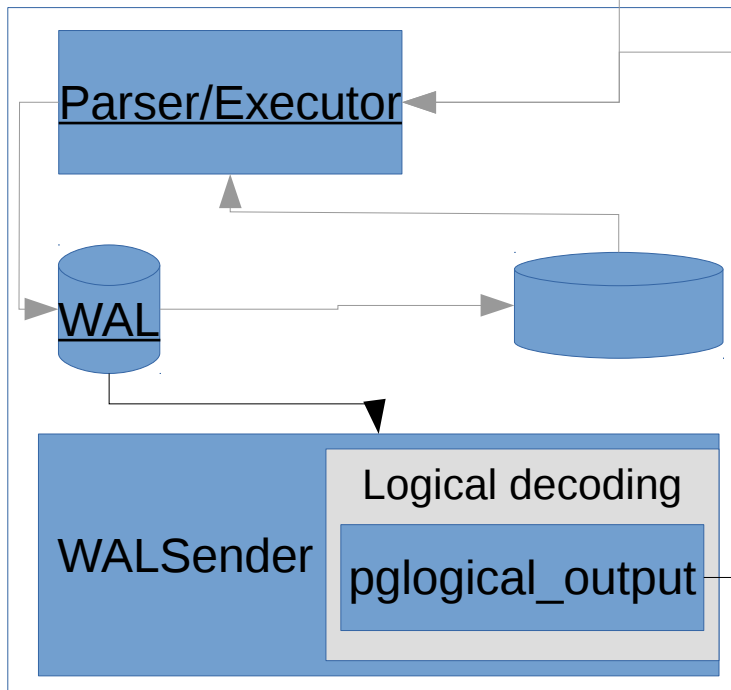
More flexible standby with read/write tables, no query cancels

Not just 1:1 – sharding, data gather, ...



Pglogical

PostgreSQL master




Client

Client

Client

Client

PostgreSQL logical replica





Pglogical: now

Selective replication

Crash safe

Downstream replica writeable

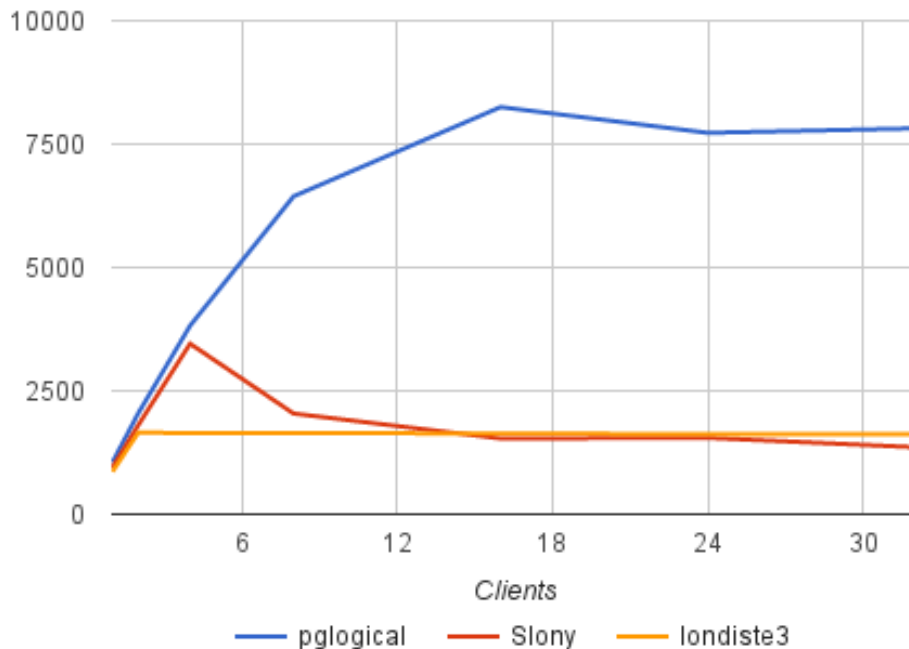
No query cancels on downstream

Efficient COPY-like apply process

No-downtime cross-version upgrade



Pglogical: performance now





Pglogical: future

Filter rows by WHERE clause

Logical and Physical Failover

Continuous ETL, transform

Continuous Data Warehouse ingestion

Automatic DDL replication

....



Sending data to apps

Take json or native output from
`pglogical_output`

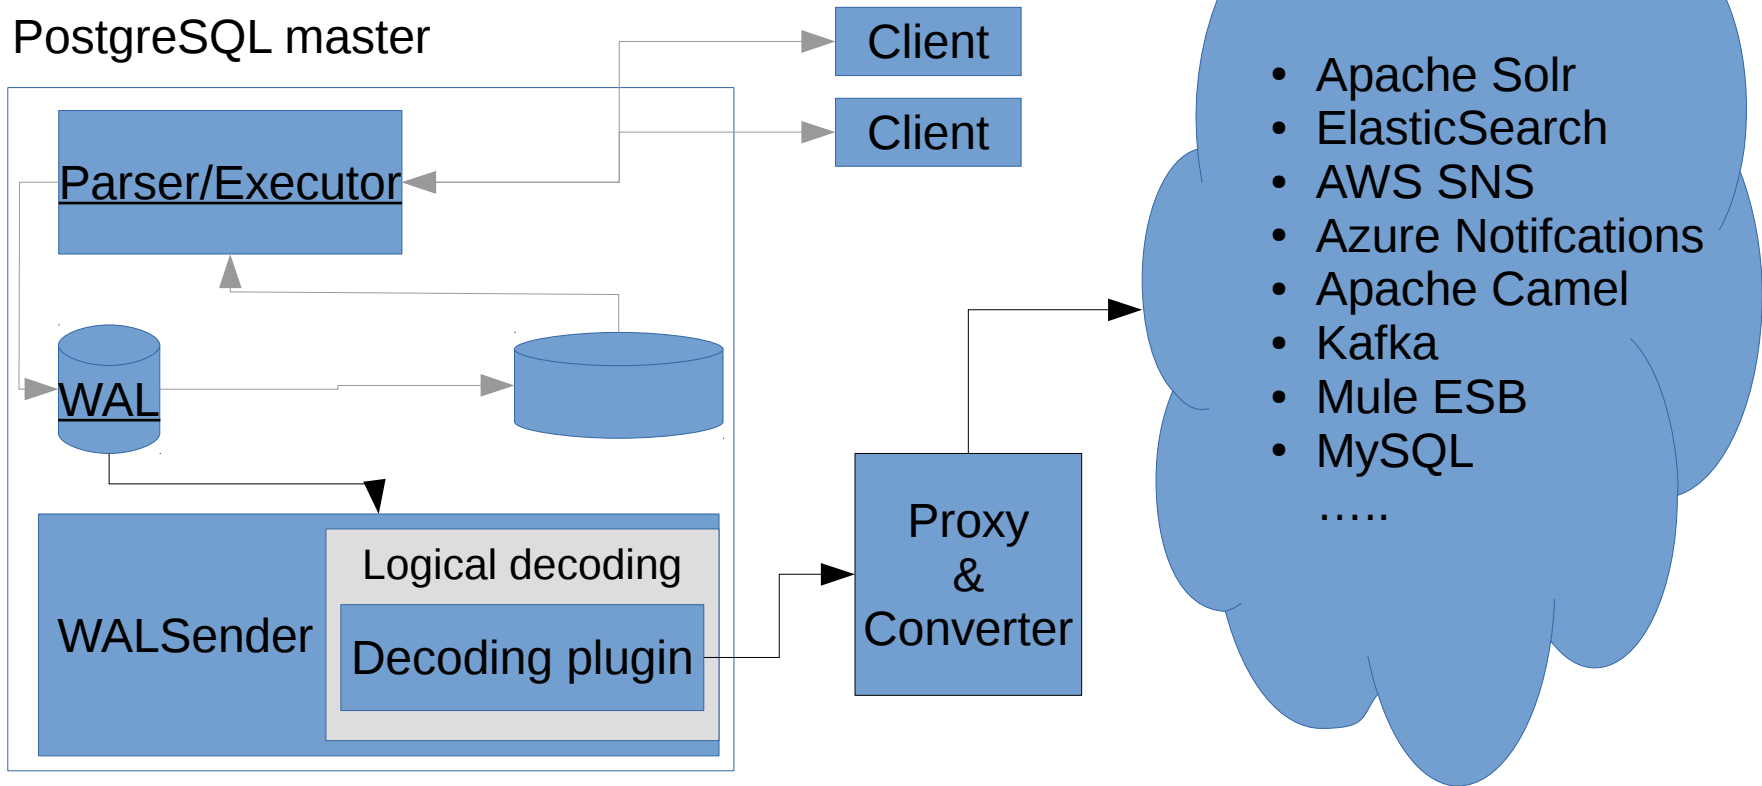
Proxy it to the app with a script

Ingest it into the app



Decoding to apps

PostgreSQL master





Demo: solr

Take json or native output

Proxy it

Ingest it



Demo: solr

The code is simple but not brief enough to list here.

<https://gist.github.com/ringerc/f74a12e430866ccd9227>



Demo: solr

The process:

- Make a normal psycopg2 connection
- Create a replication with `pg_create_logical_replication_slot slot` if it doesn't exist
- Loop over `pg_logical_slot_get_changes` to fetch the change stream
- Accumulate a whole transaction's worth of rows
- Transform the JSON from each call into something Solr will understand
- Send the transaction to Solr over http



Initial database state

No good way to send rows already in the database when we set up decoding.

Workaround:

```
COPY (SELECT row_to_json(x)
FROM my_table x) TO stdout
```



Initial state

Ugly? Very. Plenty of room for improvement.



pglogical pitfalls

- Unused slots can fill pg_xlog
- DDL isn't replicated yet
- Serial streaming of big xacts causes latency
- Big xacts need extra disk space
- Sequences not replicated yet



Using pglogical

I won't repeat the docs.

Yes, there are docs.

The `pglogical_output` protocol is documented too, for app devs.



Slots: a public service announcement

Replication slots *prevent the server from removing still-needed WAL from `pg_xlog`.*

An abandoned, unused slot can cause `pg_xlog` to fill up and your server to stop.

Unused logical slots also create bloat in the catalogs.



Slots: a public service announcement

Add `pg_replication_slots` replay lag to your monitoring and alerting system. Use `pg_xlog_location_diff(...)`

You'll already have alerts on `pg_xlog` disk space, of course. Right?

Just like you regularly test your backups.



Questions?



Questions?

Q: Can pglogical be used to receive data from non-PostgreSQL sources and stream it into PostgreSQL?

A: Not yet but there's room to support it in the design of the receiver. Good idea.

Q: Can I replicate to non-PostgreSQL databases?

A: No, and the pglogical downstream isn't designed to do that. You could use the `pglogical_output` plugin to provide the data extraction and streaming facilities you need to send data to your own downstream though.